

pestpp-glm

May 18, 2019

1 PESTPP-GLM

In this notebook, we will run PESTPP-GLM in standard parameter estimation mode and regularization mode. In both cases, we will use the baked-in bayes-linear posterior monte carlo analysis to get posterior forecast PDFs. We will use the prior monte carlo outputs as the prior forecast PDF.

```
In [1]: import os
import shutil
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib as mpl
plt.rcParams['font.size']=12
import flopy
import pyemu
```

flopy is installed in /Users/jeremyw/Dev/gw1876/activities_2day_mfm/notebooks/flopy

1.1 SUPER IMPORTANT: SET HOW MANY PARALLEL WORKERS TO USE

```
In [2]: num_workers = 20
```

```
In [3]: t_d = "template"
m_d = "master_glm"
```

```
In [4]: pst = pyemu.Pst(os.path.join(t_d, "freyberg.pst"))
pst.write_par_summary_table(filename="none")
```

```
Out[4]:
```

		type	transform	count	initial value	\
gr_prsity5	gr_prsity5		log	705	0	
pp_strt0	pp_strt0		log	32	0	
flow	flow		log	1	0	
gr_vka3	gr_vka3		log	705	0	
pp_vka2	pp_vka2		log	32	0	
cn_sy7	cn_sy7		log	1	0	
pp_strt1	pp_strt1		log	32	0	
gr_vka4	gr_vka4		log	705	0	

cn_prsity6	cn_prsity6	log	1	0
gr_ss5	gr_ss5	log	705	0
pp_hk0	pp_hk0	log	32	0
pp_prsity1	pp_prsity1	log	32	0
cn_strt8	cn_strt8	log	1	0
pp_rech1	pp_rech1	log	32	0
cn_ss6	cn_ss6	log	1	0
cn_vka8	cn_vka8	log	1	0
gr_strt5	gr_strt5	log	705	0
pp_prsity0	pp_prsity0	log	32	0
cn_vka6	cn_vka6	log	1	0
cn_hk7	cn_hk7	log	1	0
pp_strt2	pp_strt2	log	32	0
cn_hk8	cn_hk8	log	1	0
welflux	welflux	log	2	0 to 0.176091
cn_prsity8	cn_prsity8	log	1	0
cn_hk6	cn_hk6	log	1	0
gr_prsity4	gr_prsity4	log	705	0
cn_rech5	cn_rech5	log	1	-0.39794
gr_ss3	gr_ss3	log	705	0
cn_strt7	cn_strt7	log	1	0
gr_rech2	gr_rech2	log	705	0
...
strk	strk	log	40	0
pp_ss1	pp_ss1	log	32	0
welflux_k02	welflux_k02	log	6	0
gr_hk4	gr_hk4	log	705	0
gr_sy5	gr_sy5	log	705	0
pp_ss2	pp_ss2	log	32	0
cn_prsity7	cn_prsity7	log	1	0
gr_prsity3	gr_prsity3	log	705	0
gr_sy3	gr_sy3	log	705	0
cn_sy8	cn_sy8	log	1	0
cn_strt6	cn_strt6	log	1	0
gr_strt4	gr_strt4	log	705	0
pp_ss0	pp_ss0	log	32	0
pp_prsity2	pp_prsity2	log	32	0
pp_hk1	pp_hk1	log	32	0
pp_vka0	pp_vka0	log	32	0
pp_rech0	pp_rech0	log	32	0
gr_strt3	gr_strt3	log	705	0
gr_ss4	gr_ss4	log	705	0
gr_sy4	gr_sy4	log	705	0
pp_sy0	pp_sy0	log	32	0
drncond_k00	drncond_k00	log	10	0
gr_hk5	gr_hk5	log	705	0
cn_vka7	cn_vka7	log	1	0
pp_sy1	pp_sy1	log	32	0

cn_sy6	cn_sy6	log	1	0
pp_hk2	pp_hk2	log	32	0
gr_vka5	gr_vka5	log	705	0
cn_ss8	cn_ss8	log	1	0
pp_vka1	pp_vka1	log	32	0

	upper bound	lower bound	standard deviation
gr_prsity5	0.176091	-0.30103	0.11928
pp_strt0	0.0211893	-0.0222764	0.0108664
flow	0.09691	-0.124939	0.0554622
gr_vka3	1	-1	0.5
pp_vka2	1	-1	0.5
cn_sy7	0.243038	-0.60206	0.211275
pp_strt1	0.0211893	-0.0222764	0.0108664
gr_vka4	1	-1	0.5
cn_prsity6	0.176091	-0.30103	0.11928
gr_ss5	1	-1	0.5
pp_hk0	1	-1	0.5
pp_prsity1	0.176091	-0.30103	0.11928
cn_strt8	0.0211893	-0.0222764	0.0108664
pp_rech1	0.0413927	-0.0457575	0.0217875
cn_ss6	1	-1	0.5
cn_vka8	1	-1	0.5
gr_strt5	0.0211893	-0.0222764	0.0108664
pp_prsity0	0.176091	-0.30103	0.11928
cn_vka6	1	-1	0.5
cn_hk7	1	-1	0.5
pp_strt2	0.0211893	-0.0222764	0.0108664
cn_hk8	1	-1	0.5
welflux	0.176091 to 0.30103	-0.30103 to 0	0.0752575 to 0.11928
cn_prsity8	0.176091	-0.30103	0.11928
cn_hk6	1	-1	0.5
gr_prsity4	0.176091	-0.30103	0.11928
cn_rech5	-0.09691	-1	0.225772
gr_ss3	1	-1	0.5
cn_strt7	0.0211893	-0.0222764	0.0108664
gr_rech2	0.0413927	-0.0457575	0.0217875
...
strk	2	-2	1
pp_ss1	1	-1	0.5
welflux_k02	1	-1	0.5
gr_hk4	1	-1	0.5
gr_sy5	0.243038	-0.60206	0.211275
pp_ss2	1	-1	0.5
cn_prsity7	0.176091	-0.30103	0.11928
gr_prsity3	0.176091	-0.30103	0.11928
gr_sy3	0.243038	-0.60206	0.211275
cn_sy8	0.243038	-0.60206	0.211275

cn_strt6	0.0211893	-0.0222764	0.0108664
gr_strt4	0.0211893	-0.0222764	0.0108664
pp_ss0	1	-1	0.5
pp_prsity2	0.176091	-0.30103	0.11928
pp_hk1	1	-1	0.5
pp_vka0	1	-1	0.5
pp_rech0	0.0413927	-0.0457575	0.0217875
gr_strt3	0.0211893	-0.0222764	0.0108664
gr_ss4	1	-1	0.5
gr_sy4	0.243038	-0.60206	0.211275
pp_sy0	0.243038	-0.60206	0.211275
drncond_k00	1	-1	0.5
gr_hk5	1	-1	0.5
cn_vka7	1	-1	0.5
pp_sy1	0.243038	-0.60206	0.211275
cn_sy6	0.243038	-0.60206	0.211275
pp_hk2	1	-1	0.5
gr_vka5	1	-1	0.5
cn_ss8	1	-1	0.5
pp_vka1	1	-1	0.5

[65 rows x 7 columns]

1.1.1 reduce the number of adjustable parameters

This is the painful part: we cant use 10K+ pars because we cant wait around for that many runs and then the linear algebra of factoring a 10k+ by 10K+ matrix is also difficult. So that means we need to fix a lot a parameters #frownyface

```
In [5]: par = pst.parameter_data
```

```
In [6]: # grid-scale pars
```

```
gr_pars = par.loc[par.pargp.apply(lambda x: "gr" in x), "parnme"]
par.loc[gr_pars, "partrans"] = "fixed"
pst.npar_adj
```

```
Out[6]: 719
```

```
In [7]: # these are the sfr conductance parameters - Ive left all 40 adjustable
# but if you uncomment this, it will tie them into 1 parameter effectively
# strk_pars = par.loc[par.pargp=="strk", "parnme"]
# p1 = strk_pars.iloc[0]
# par.loc[strk_pars.iloc[1:], "partrans"] = "tied"
# par.loc[strk_pars.iloc[1:], "partied"] = p1
pst.npar_adj
```

```
Out[7]: 719
```

```
In [8]: par.loc[par.pargp.apply(lambda x: "pp" in x), "pargp"].unique()
```

```
Out[8]: array(['pp_hk0', 'pp_hk1', 'pp_hk2', 'pp_prsity0', 'pp_prsity1',
              'pp_prsity2', 'pp_rech0', 'pp_rech1', 'pp_ss0', 'pp_ss1', 'pp_ss2',
              'pp_strt0', 'pp_strt1', 'pp_strt2', 'pp_sy0', 'pp_sy1', 'pp_sy2',
              'pp_vka0', 'pp_vka1', 'pp_vka2'], dtype=object)
```

Fix the storage pilot points - we still have layer-scale storage pars adjustable

```
In [9]: #s_pars = par.loc[par.pargp.apply(lambda x: "pp" in x and ("ss" in x or "sy" in x)), "p
        #par.loc[s_pars, "partrans"] = "fixed"
        pst.npar_adj
```

```
Out[9]: 719
```

```
In [10]: adj_par = par.loc[par.partrans=="log", :]
         adj_par.pargp.value_counts().sort_values()
```

```
Out[10]: cn_strt7      1
         cn_prsity6    1
         cn_strt8      1
         cn_vka8        1
         flow          1
         cn_ss6         1
         cn_vka6        1
         cn_ss8         1
         cn_hk6         1
         cn_sy7         1
         cn_sy8         1
         cn_hk7         1
         cn_prsity8     1
         cn_sy6         1
         cn_hk8         1
         cn_rech4       1
         cn_ss7         1
         cn_prsity7     1
         cn_strt6       1
         cn_vka7        1
         cn_rech5       1
         welflux        2
         welflux_k02    6
         drncond_k00    10
         pp_strt1       32
         pp_prsity0     32
         pp_hk2         32
         pp_sy0         32
         pp_sy2         32
         pp_vka2        32
         pp_strt2       32
         pp_ss1         32
         pp_prsity1     32
```

```

pp_ss2          32
pp_hk0          32
pp_prsity2      32
pp_hk1          32
pp_vka0         32
pp_rech0        32
pp_strt0        32
pp_rech1        32
pp_vka1         32
pp_sy1          32
pp_ss0          32
strk            40
Name: pargp, dtype: int64

```

fix the future recharge pilot points, vka in layers 1 and 3 and the initial condition pilot points (we still have layer-scale pars for each of these types)

```

In [11]: fi_grps = ["pp_rech1", "pp_vka0", "pp_vka2", "pp_strt0", "pp_strt1", "pp_strt2"]
          par.loc[par.pargp.apply(lambda x: x in fi_grps), "partrans"] = "fixed"
          pst.npar_adj

```

```

Out[11]: 527

```

Ok, thats better...so lets run PESTPP-GLM. We will use a single “base parameter” jacobian matrix as the basis for 6 super parameter iterations. Then we will draw 100 realizations from the FOSM posterior parameter covariance matrix and run those 100 realizations to get the psoterior forecast PDFs

```

In [12]: pst.control_data.noptmax = 3
          pst.pestpp_options["n_iter_base"] = -1
          pst.pestpp_options["n_iter_super"] = 3
          pst.pestpp_options["num_reals"] = 50 # this is how many ies uses
          pst.pestpp_options["parcov"] = "prior_cov.jcb"
          pst.write(os.path.join(t_d, "freyberg_pp.pst"))

```

```

noptmax:3, npar_adj:527, nnz_obs:14

```

```

In [13]: pyemu.os_utils.start_slaves(t_d, "pestpp-glm", "freyberg_pp.pst", num_slaves=num_workers
          master_dir=m_d)

```

```

In [14]: df = df=pd.read_csv(os.path.join(m_d, "freyberg_pp.post.obsen.csv"), index_col=0)
          oe = pyemu.ObservationEnsemble.from_dataframe(pst=pst, df=df)

```

```

In [15]: ax = oe.phi_vector.hist()#bins=np.linspace(0,100,20))
          oe.phi_vector.sort_values().iloc[:20]

```

```


Out[15]: real_name
11      67.311123

```

```

19      73.629202
48      80.294490
37      85.556290
39     130.363406
46     142.101852
16     170.446318
34     192.605624
33     206.148352
24     220.418561
29     254.451689
1      255.306658
17     289.671776
26     398.593623
7      428.985089
45     516.951351
44     570.947286
38     664.273902
49     679.494764
23     681.123367
dtype: float64

```



Here we see the distribution of phi values across the 100 posterior realizations. Should we accept all of these??? The theoretical phi we should accept is number of nonzero obs (14).

To get a “posterior” ensemble, we need to throw out the realizations with large phi - lets just take the 20 best:

```
In [16]: oe_pt = oe.loc[oe.phi_vector.sort_values().index[:20],:] #just take the 20 lowest phi
```

We can also load and plot the FOSM forecast results along side of the ensemble results:

```
In [17]: f_df = pd.read_csv(os.path.join(m_d,"freyberg_pp.pred.usum.csv"),index_col=0)
f_df.index = f_df.index.map(str.lower)
f_df
```

```
Out [17]:
```

	prior_mean	prior_stddev	prior_lower_bound \
name			
fa_hw_19791230	-977.2390	295.32800	-1567.8900
fa_hw_19801229	-351.2160	409.77000	-1170.7600
fa_tw_19791230	-453.0330	409.35100	-1271.7400
fa_tw_19801229	108.9600	506.73200	-904.5040
hds_00_013_002_000	39.6102	3.96314	31.6840
hds_00_013_002_001	38.3838	4.05782	30.2681
part_status	2.0000	0.00000	2.0000
part_time	907.7020	570.98600	-234.2690

	prior_upper_bound	post_mean	post_stddev \
name			
fa_hw_19791230	-386.5840	-691.3690	248.764000
fa_hw_19801229	468.3240	38.8612	336.875000
fa_tw_19791230	365.6690	-315.6270	172.537000
fa_tw_19801229	1122.4200	335.7690	265.523000
hds_00_013_002_000	47.5365	37.6724	0.283577
hds_00_013_002_001	46.4994	36.4533	0.691648
part_status	2.0000	2.0000	0.000000
part_time	2049.6700	660.4670	437.861000

	post_lower_bound	post_upper_bound
name		
fa_hw_19791230	-1188.9000	-193.8410
fa_hw_19801229	-634.8880	712.6110
fa_tw_19791230	-660.7020	29.4469
fa_tw_19801229	-195.2760	866.8150
hds_00_013_002_000	37.1052	38.2395
hds_00_013_002_001	35.0700	37.8366
part_status	2.0000	2.0000
part_time	-215.2550	1536.1900


```
In [18]: obs = pst.observation_data
fnames = pst.pestpp_options["forecasts"].split(",")
for forecast in fnames:
    ax = plt.subplot(111)
    oe_pt.loc[:,forecast].hist(ax=ax,color="b",alpha=0.5,normed=True)
    ax.plot([obs.loc[forecast,"obsval"],obs.loc[forecast,"obsval"]],ax.get_ylim(),"r")
    axt = plt.twinx()
    x,y = pyemu.plot_utils.gaussian_distribution(f_df.loc[forecast,"prior_mean"],f_df
    axt.fill_between(x,0,y,facecolor="0.5",alpha=0.25)
```







```



x,y = pyemu.plot_utils.gaussian_distribution(f_df.loc[forecast,"post_mean"],f_df.
axt.fill_between(x,0,y,facecolor="b",alpha=0.25)
axt.set_ylim(0,axt.get_ylim()[1])
axt.set_yticks([])
ax.set_title(forecast)
plt.show()


```











1.1.2 Setup of Tikhonov regularization

Now lets setup and use some formal regularization to bring the final phi up to around 14. We will use first-order regularization based on the covariance matrix we build earlier:

```
In [19]: cov = pyemu.Cov.from_binary(os.path.join(t_d,"prior_cov.jcb"))
```

```
new binary format detected...
```

```
In [20]: pyemu.helpers.first_order_pearson_tikhonov(pst,cov)
```

```
getting CC matrix
processing
```

```
In [21]: pst.prior_information.head()
```

```
Out[21]:
```

	equation	obgnme	\
pilbl			
pcc_1	$1.0 * \log(\text{dc0000390005}) - 1.0 * \log(\text{dc0000390006}) = 0.0$	regul_cc	
pcc_2	$1.0 * \log(\text{dc0000390005}) - 1.0 * \log(\text{dc0000390007}) = 0.0$	regul_cc	
pcc_3	$1.0 * \log(\text{dc0000390005}) - 1.0 * \log(\text{dc0000390008}) = 0.0$	regul_cc	
pcc_4	$1.0 * \log(\text{dc0000390005}) - 1.0 * \log(\text{dc0000390009}) = 0.0$	regul_cc	
pcc_5	$1.0 * \log(\text{dc0000390005}) - 1.0 * \log(\text{dc0000390010}) = 0.0$	regul_cc	

	pilbl	weight
pilbl		
pcc_1	pcc_1	0.904837
pcc_2	pcc_2	0.818731
pcc_3	pcc_3	0.740818
pcc_4	pcc_4	0.670320
pcc_5	pcc_5	0.606531

```
In [22]: shutil.copy2(os.path.join(m_d, "freyberg_pp.jcb"), os.path.join(t_d, "restart_pp.jcb"))
```

```
Out[22]: 'template/restart_pp.jcb'
```

```
In [23]: pst.pestpp_options["base_jacobian"] = "restart_pp.jcb"
pst.reg_data.phimlim = pst.nnz_obs
pst.reg_data.phimaccept = pst.reg_data.phimlim * 1.1
pst.write(os.path.join(t_d, "freyberg_pp.pst"))
```

```
noptmax:3, npar_adj:527, nnz_obs:14
```

```
In [24]: pyemu.os_utils.start_slaves(t_d, "pestpp-glm", "freyberg_pp.pst", num_slaves=num_workers,
                                     master_dir=m_d)
```

```
In [25]: df = df=pd.read_csv(os.path.join(m_d, "freyberg_pp.post.obsen.csv"), index_col=0)
oe = pyemu.ObservationEnsemble.from_dataframe(pst=pst, df=df)
```


```
In [26]: ax = oe.phi_vector.hist()#bins=np.linspace(0,100,20)
oe.phi_vector.sort_values().iloc[:20]
```

```
Out[26]: real_name
48      36.238820
38      54.829565
11      63.978008
37     103.284082
46     109.420095
19     132.244209
26     140.618000
34     166.380438
17     174.431787
16     191.814022
24     201.539107
1      203.762385
39     233.843119
33     236.361174
5      270.972756
47     308.716219
29     395.603897
44     450.205197
```

```

6      469.426236
49     477.567389
dtype: float64

```



Same as before, to get a “posterior” ensemble, we need to throw out the realizations with large phi - lets just take the 20 best:

```

In [27]: oe_pt = oe.loc[oe.phi_vector.sort_values().index[:20],:]

In [28]: f_df = pd.read_csv(os.path.join(m_d,"freyberg_pp.pred.usum.csv"),index_col=0)
f_df.index = f_df.index.map(str.lower)
f_df

```

```

Out [28]:

```

	prior_mean	prior_stdev	prior_lower_bound \
name			
fa_hw_19791230	-977.2390	295.32800	-1567.8900
fa_hw_19801229	-351.2160	409.77000	-1170.7600
fa_tw_19791230	-453.0330	409.35100	-1271.7400
fa_tw_19801229	108.9600	506.73200	-904.5040
hds_00_013_002_000	39.6102	3.96314	31.6840
hds_00_013_002_001	38.3838	4.05782	30.2681
part_status	2.0000	0.00000	2.0000
part_time	907.7020	570.98600	-234.2690

	prior_upper_bound	post_mean	post_stdev \
name			
fa_hw_19791230	-386.5840	-696.3770	252.671000

fa_hw_19801229	468.3240	-128.8820	342.901000
fa_tw_19791230	365.6690	-292.7330	174.083000
fa_tw_19801229	1122.4200	206.5750	267.003000
hds_00_013_002_000	47.5365	37.8527	0.291081
hds_00_013_002_001	46.4994	36.9254	0.694710
part_status	2.0000	2.0000	0.000000
part_time	2049.6700	900.1010	438.233000

	post_lower_bound	post_upper_bound
name		
fa_hw_19791230	-1201.7200	-191.0340
fa_hw_19801229	-814.6840	556.9200
fa_tw_19791230	-640.8980	55.4319
fa_tw_19801229	-327.4310	740.5810
hds_00_013_002_000	37.2706	38.4349
hds_00_013_002_001	35.5360	38.3148
part_status	2.0000	2.0000
part_time	23.6357	1776.5700

```
In [29]: obs = pst.observation_data
fnames = pst.pestpp_options["forecasts"].split(",")
for forecast in fnames:
    ax = plt.subplot(111)
    oe_pt.loc[:,forecast].hist(ax=ax,color="b",alpha=0.5,normed=True)
    ax.plot([obs.loc[forecast,"obsval"],obs.loc[forecast,"obsval"]],ax.get_ylim(),"r")
    axt = plt.twinx()
    x,y = pyemu.plot_utils.gaussian_distribution(f_df.loc[forecast,"prior_mean"],f_df.loc[forecast,"prior_std"])
    axt.fill_between(x,0,y,facecolor="0.5",alpha=0.25)
    x,y = pyemu.plot_utils.gaussian_distribution(f_df.loc[forecast,"post_mean"],f_df.loc[forecast,"post_std"])
    axt.fill_between(x,0,y,facecolor="b",alpha=0.25)
    axt.set_ylim(0,axt.get_ylim()[1])
    axt.set_yticks([])
    ax.set_title(forecast)
plt.show()
```