

pestpp-glm

May 31, 2019

1 PESTPP-GLM

In this notebook, we will run PESTPP-GLM in standard parameter estimation mode and regularization mode. In both cases, we will use the baked-in bayes-linear posterior monte carlo analysis to get posterior forecast PDFs. We will use the prior monte carlo outputs as the prior forecast PDF.

```
In [1]: import os
import shutil
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib as mpl
plt.rcParams['font.size']=12
import flopy
import pyemu
```

flopy is installed in /Users/jeremyw/Dev/gw1876/activities_2day_mfm/notebooks/flopy

1.1 SUPER IMPORTANT: SET HOW MANY PARALLEL WORKERS TO USE

```
In [2]: num_workers = 20
```

```
In [3]: t_d = "template"
m_d = "master_glm"
```

```
In [4]: pst = pyemu.Pst(os.path.join(t_d, "freyberg.pst"))
pst.write_par_summary_table(filename="none")
```

```
Out[4]:
```

	type	transform	count	initial value	\
cn_strt8	cn_strt8	log	1	0	
gr_vka3	gr_vka3	log	705	0	
cn_hk7	cn_hk7	log	1	0	
pp_ss1	pp_ss1	log	32	0	
gr_hk3	gr_hk3	log	705	0	
cn_strt6	cn_strt6	log	1	0	
cn_vka7	cn_vka7	log	1	0	
cn_vka6	cn_vka6	log	1	0	

pp_ss0	pp_ss0	log	32	0
gr_rech2	gr_rech2	log	705	0
gr_prsity5	gr_prsity5	log	705	0
gr_vka4	gr_vka4	log	705	0
gr_prsity3	gr_prsity3	log	705	0
gr_sy4	gr_sy4	log	705	0
gr_strt3	gr_strt3	log	705	0
pp_rech1	pp_rech1	log	32	0
drncond_k00	drncond_k00	log	10	0
cn_rech5	cn_rech5	log	1	-0.39794
gr_ss3	gr_ss3	log	705	0
gr_strt5	gr_strt5	log	705	0
pp_vka0	pp_vka0	log	32	0
gr_hk4	gr_hk4	log	705	0
gr_prsity4	gr_prsity4	log	705	0
pp_vka2	pp_vka2	log	32	0
gr_ss5	gr_ss5	log	705	0
gr_sy5	gr_sy5	log	705	0
pp_ss2	pp_ss2	log	32	0
pp_sy0	pp_sy0	log	32	0
gr_vka5	gr_vka5	log	705	0
cn_ss6	cn_ss6	log	1	0
...
cn_rech4	cn_rech4	log	1	0
cn_vka8	cn_vka8	log	1	0
cn_sy8	cn_sy8	log	1	0
pp_strt1	pp_strt1	log	32	0
cn_prsity7	cn_prsity7	log	1	0
pp_prsity1	pp_prsity1	log	32	0
flow	flow	log	1	0
cn_sy7	cn_sy7	log	1	0
gr_strt4	gr_strt4	log	705	0
cn_sy6	cn_sy6	log	1	0
gr_hk5	gr_hk5	log	705	0
gr_ss4	gr_ss4	log	705	0
gr_rech3	gr_rech3	log	705	0
pp_prsity0	pp_prsity0	log	32	0
welflux_k02	welflux_k02	log	6	0
pp_strt2	pp_strt2	log	32	0
cn_hk8	cn_hk8	log	1	0
strk	strk	log	40	0
cn_ss8	cn_ss8	log	1	0
pp_hk1	pp_hk1	log	32	0
pp_rech0	pp_rech0	log	32	0
pp_prsity2	pp_prsity2	log	32	0
cn_strt7	cn_strt7	log	1	0
pp_vka1	pp_vka1	log	32	0
cn_prsity6	cn_prsity6	log	1	0

cn_ss7	cn_ss7	log	1	0
welflux	welflux	log	2	0 to 0.176091
cn_hk6	cn_hk6	log	1	0
pp_hk0	pp_hk0	log	32	0
pp_hk2	pp_hk2	log	32	0

	upper bound	lower bound	standard deviation
cn_strt8	0.0211893	-0.0222764	0.0108664
gr_vka3	1	-1	0.5
cn_hk7	1	-1	0.5
pp_ss1	1	-1	0.5
gr_hk3	1	-1	0.5
cn_strt6	0.0211893	-0.0222764	0.0108664
cn_vka7	1	-1	0.5
cn_vka6	1	-1	0.5
pp_ss0	1	-1	0.5
gr_rech2	0.0413927	-0.0457575	0.0217875
gr_prsity5	0.176091	-0.30103	0.11928
gr_vka4	1	-1	0.5
gr_prsity3	0.176091	-0.30103	0.11928
gr_sy4	0.243038	-0.60206	0.211275
gr_strt3	0.0211893	-0.0222764	0.0108664
pp_rech1	0.0413927	-0.0457575	0.0217875
drncond_k00	1	-1	0.5
cn_rech5	-0.09691	-1	0.225772
gr_ss3	1	-1	0.5
gr_strt5	0.0211893	-0.0222764	0.0108664
pp_vka0	1	-1	0.5
gr_hk4	1	-1	0.5
gr_prsity4	0.176091	-0.30103	0.11928
pp_vka2	1	-1	0.5
gr_ss5	1	-1	0.5
gr_sy5	0.243038	-0.60206	0.211275
pp_ss2	1	-1	0.5
pp_sy0	0.243038	-0.60206	0.211275
gr_vka5	1	-1	0.5
cn_ss6	1	-1	0.5
...
cn_rech4	0.0791812	-0.09691	0.0440228
cn_vka8	1	-1	0.5
cn_sy8	0.243038	-0.60206	0.211275
pp_strt1	0.0211893	-0.0222764	0.0108664
cn_prsity7	0.176091	-0.30103	0.11928
pp_prsity1	0.176091	-0.30103	0.11928
flow	0.09691	-0.124939	0.0554622
cn_sy7	0.243038	-0.60206	0.211275
gr_strt4	0.0211893	-0.0222764	0.0108664
cn_sy6	0.243038	-0.60206	0.211275

gr_hk5	1	-1	0.5
gr_ss4	1	-1	0.5
gr_rech3	0.0413927	-0.0457575	0.0217875
pp_prsity0	0.176091	-0.30103	0.11928
welflux_k02	1	-1	0.5
pp_strt2	0.0211893	-0.0222764	0.0108664
cn_hk8	1	-1	0.5
strk	2	-2	1
cn_ss8	1	-1	0.5
pp_hk1	1	-1	0.5
pp_rech0	0.0413927	-0.0457575	0.0217875
pp_prsity2	0.176091	-0.30103	0.11928
cn_strt7	0.0211893	-0.0222764	0.0108664
pp_vka1	1	-1	0.5
cn_prsity6	0.176091	-0.30103	0.11928
cn_ss7	1	-1	0.5
welflux	0.176091 to 0.30103	-0.30103 to 0	0.0752575 to 0.11928
cn_hk6	1	-1	0.5
pp_hk0	1	-1	0.5
pp_hk2	1	-1	0.5

[65 rows x 7 columns]

1.1.1 reduce the number of adjustable parameters

This is the painful part: we cant use 10K+ pars because we cant wait around for that many runs and then the linear algebra of factoring a 10k+ by 10K+ matrix is also difficult. So that means we need to fix a lot a parameters #frownyface

```
In [5]: par = pst.parameter_data
```

```
In [6]: # grid-scale pars
```

```
gr_pars = par.loc[par.pargp.apply(lambda x: "gr" in x), "parnme"]
par.loc[gr_pars, "partrans"] = "fixed"
pst.npar_adj
```

```
Out[6]: 719
```

```
In [7]: # these are the sfr conductance parameters - Ive left all 40 adjustable
# but if you uncomment this, it will tie them into 1 parameter effectively
# strk_pars = par.loc[par.pargp=="strk", "parnme"]
# p1 = strk_pars.iloc[0]
# par.loc[strk_pars.iloc[1:], "partrans"] = "tied"
# par.loc[strk_pars.iloc[1:], "partied"] = p1
pst.npar_adj
```

```
Out[7]: 719
```

```
In [8]: par.loc[par.pargp.apply(lambda x: "pp" in x), "pargp"].unique()
```

```
Out[8]: array(['pp_hk0', 'pp_hk1', 'pp_hk2', 'pp_prsity0', 'pp_prsity1',
              'pp_prsity2', 'pp_rech0', 'pp_rech1', 'pp_ss0', 'pp_ss1', 'pp_ss2',
              'pp_strt0', 'pp_strt1', 'pp_strt2', 'pp_sy0', 'pp_sy1', 'pp_sy2',
              'pp_vka0', 'pp_vka1', 'pp_vka2'], dtype=object)
```

Fix the storage pilot points - we still have layer-scale storage pars adjustable

```
In [9]: #s_pars = par.loc[par.pargp.apply(lambda x: "pp" in x and ("ss" in x or "sy" in x)), "p
        #par.loc[s_pars, "partrans"] = "fixed"
        pst.npar_adj
```

```
Out[9]: 719
```

```
In [10]: adj_par = par.loc[par.partrans=="log",:]
         adj_par.pargp.value_counts().sort_values()
```

```
Out[10]: cn_strt8      1
         cn_sy6        1
         cn_vka6        1
         cn_vka7        1
         cn_hk6         1
         flow          1
         cn_strt6       1
         cn_hk7         1
         cn_rech5       1
         cn_prsity6     1
         cn_ss7         1
         cn_sy7         1
         cn_strt7       1
         cn_ss6         1
         cn_ss8         1
         cn_prsity8     1
         cn_hk8         1
         cn_rech4       1
         cn_vka8        1
         cn_sy8         1
         cn_prsity7     1
         welflux        2
         welflux_k02    6
         drncond_k00    10
         pp_hk1         32
         pp_ss0         32
         pp_vka1        32
         pp_prsity1     32
         pp_rech1       32
         pp_vka0        32
         pp_ss2         32
         pp_sy0         32
         pp_ss1         32
```

```

pp_strt2      32
pp_sy1        32
pp_hk2        32
pp_strt1      32
pp_strt0      32
pp_rech0      32
pp_prsity0    32
pp_hk0        32
pp_prsity2    32
pp_vka2       32
pp_sy2        32
strk          40
Name: pargp, dtype: int64

```

fix the future recharge pilot points, vka in layers 1 and 3 and the initial condition pilot points (we still have layer-scale pars for each of these types)

```

In [11]: fi_grps = ["pp_rech1", "pp_vka0", "pp_vka2", "pp_strt0", "pp_strt1", "pp_strt2"]
          par.loc[par.pargp.apply(lambda x: x in fi_grps), "partrans"] = "fixed"
          pst.npar_adj

```

```

Out[11]: 527

```

Ok, thats better...so lets run PESTPP-GLM. We will use a single “base parameter” jacobian matrix as the basis for 6 super parameter iterations. Then we will draw 100 realizations from the FOSM posterior parameter covariance matrix and run those 100 realizations to get the psoterior forecast PDFs

```

In [12]: pst.control_data.noptmax = 3
          pst.pestpp_options["n_iter_base"] = -1
          pst.pestpp_options["n_iter_super"] = 3
          pst.pestpp_options["num_reals"] = 50 # this is how many ies uses
          pst.pestpp_options["parcov"] = "prior_cov.jcb"
          pst.write(os.path.join(t_d, "freyberg_pp.pst"))

```

```

noptmax:3, npar_adj:527, nnz_obs:14

```

```

In [13]: pyemu.os_utils.start_slaves(t_d, "pestpp-glm", "freyberg_pp.pst", num_slaves=num_workers
          master_dir=m_d)

```

```

In [14]: df = df=pd.read_csv(os.path.join(m_d, "freyberg_pp.post.obsen.csv"), index_col=0)
          oe = pyemu.ObservationEnsemble.from_dataframe(pst=pst, df=df)

```

```

In [15]: ax = oe.phi_vector.hist()#bins=np.linspace(0,100,20))
          oe.phi_vector.sort_values().iloc[:20]

```

```


Out[15]: real_name
          25      57.504700

```

```

49      59.965027
48      79.076502
22     121.743995
12     129.003257
21     235.929537
10     271.640362
26     303.531649
42     328.530748
29     370.831442
24     473.160536
47     474.512084
33     495.605323
44     505.914044
5      507.399642
1      537.591081
9      608.619145
19     616.414116
46     680.674779
37     805.598204
dtype: float64

```



Here we see the distribution of phi values across the 100 posterior realizations. Should we accept all of these??? The theoretical phi we should accept is number of nonzero obs (14).

To get a “posterior” ensemble, we need to throw out the realizations with large phi - lets just take the 20 best:

```
In [16]: oe_pt = oe.loc[oe.phi_vector.sort_values().index[:20],:] #just take the 20 lowest phi
```

We can also load and plot the FOSM forecast results along side of the ensemble results:

```
In [17]: f_df = pd.read_csv(os.path.join(m_d,"freyberg_pp.pred.usum.csv"),index_col=0)
f_df.index = f_df.index.map(str.lower)
f_df
```

```
Out [17]:
```

	prior_mean	prior_stddev	prior_lower_bound \
name			
fa_hw_19791230	-977.2390	295.32800	-1567.8900
fa_hw_19801229	-351.2160	409.77000	-1170.7600
fa_tw_19791230	-453.0330	409.35100	-1271.7400
fa_tw_19801229	108.9600	506.73200	-904.5040
hds_00_013_002_000	39.6102	3.96314	31.6840
hds_00_013_002_001	38.3838	4.05782	30.2681
part_status	2.0000	0.00000	2.0000
part_time	907.7020	570.98600	-234.2690

	prior_upper_bound	post_mean	post_stddev \
name			
fa_hw_19791230	-386.5840	-675.0430	247.837000
fa_hw_19801229	468.3240	37.6125	335.549000
fa_tw_19791230	365.6690	-299.4690	170.577000
fa_tw_19801229	1122.4200	333.7460	263.672000
hds_00_013_002_000	47.5365	37.6359	0.273990
hds_00_013_002_001	46.4994	36.4927	0.687833
part_status	2.0000	2.0000	0.000000
part_time	2049.6700	703.9470	436.945000

	post_lower_bound	post_upper_bound
name		
fa_hw_19791230	-1170.7200	-179.3700
fa_hw_19801229	-633.4850	708.7100
fa_tw_19791230	-640.6230	41.6858
fa_tw_19801229	-193.5970	861.0890
hds_00_013_002_000	37.0879	38.1838
hds_00_013_002_001	35.1171	37.8684
part_status	2.0000	2.0000
part_time	-169.9420	1577.8400


```
In [18]: obs = pst.observation_data
fnames = pst.pestpp_options["forecasts"].split(",")
for forecast in fnames:
    ax = plt.subplot(111)
    oe_pt.loc[:,forecast].hist(ax=ax,color="b",alpha=0.5,normed=True)
    ax.plot([obs.loc[forecast,"obsval"],obs.loc[forecast,"obsval"]],ax.get_ylim(),"r")
    axt = plt.twinx()
    x,y = pyemu.plot_utils.gaussian_distribution(f_df.loc[forecast,"prior_mean"],f_df
    axt.fill_between(x,0,y,facecolor="0.5",alpha=0.25)
```





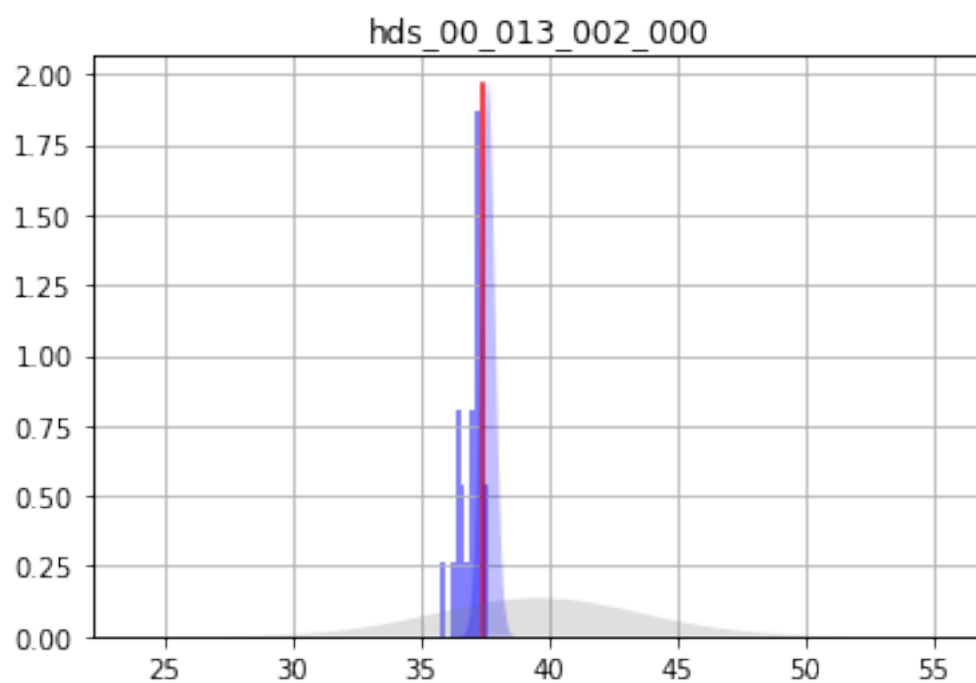

```



x,y = pyemu.plot_utils.gaussian_distribution(f_df.loc[forecast,"post_mean"],f_df.
axt.fill_between(x,0,y,facecolor="b",alpha=0.25)
axt.set_ylim(0,axt.get_ylim()[1])
axt.set_yticks([])
ax.set_title(forecast)
plt.show()


```











1.1.2 Setup of Tikhonov regularization

Now let's setup and use some formal regularization to bring the final phi up to around 14. We will use first-order regularization based on the covariance matrix we build earlier:

```
In [19]: cov = pyemu.Cov.from_binary(os.path.join(t_d,"prior_cov.jcb"))
```

```
new binary format detected...
```

```
In [20]: pyemu.helpers.first_order_pearson_tikhonov(pst,cov)
```

```
getting CC matrix
processing
```

```
In [21]: pst.prior_information.head()
```

```
Out[21]:
```

	equation	obgnme	\
pilbl			
pcc_1	$1.0 * \log(\text{dc0000390005}) - 1.0 * \log(\text{dc0000390006}) = 0.0$	regul_cc	
pcc_2	$1.0 * \log(\text{dc0000390005}) - 1.0 * \log(\text{dc0000390007}) = 0.0$	regul_cc	
pcc_3	$1.0 * \log(\text{dc0000390005}) - 1.0 * \log(\text{dc0000390008}) = 0.0$	regul_cc	
pcc_4	$1.0 * \log(\text{dc0000390005}) - 1.0 * \log(\text{dc0000390009}) = 0.0$	regul_cc	
pcc_5	$1.0 * \log(\text{dc0000390005}) - 1.0 * \log(\text{dc0000390010}) = 0.0$	regul_cc	

	pilbl	weight
pilbl		
pcc_1	pcc_1	0.904837
pcc_2	pcc_2	0.818731
pcc_3	pcc_3	0.740818
pcc_4	pcc_4	0.670320
pcc_5	pcc_5	0.606531

```
In [22]: shutil.copy2(os.path.join(m_d, "freyberg_pp.jcb"), os.path.join(t_d, "restart_pp.jcb"))
```

```
Out[22]: 'template/restart_pp.jcb'
```

```
In [23]: pst.pestpp_options["base_jacobian"] = "restart_pp.jcb"
pst.reg_data.phimlim = pst.nnz_obs
pst.reg_data.phimaccept = pst.reg_data.phimlim * 1.1
pst.write(os.path.join(t_d, "freyberg_pp.pst"))
```

```
noptmax:3, npar_adj:527, nnz_obs:14
```

```
In [24]: pyemu.os_utils.start_slaves(t_d, "pestpp-glm", "freyberg_pp.pst", num_slaves=num_workers,
                                     master_dir=m_d)
```

```
In [25]: df = df=pd.read_csv(os.path.join(m_d, "freyberg_pp.post.obsen.csv"), index_col=0)
oe = pyemu.ObservationEnsemble.from_dataframe(pst=pst, df=df)
```


```
In [26]: ax = oe.phi_vector.hist()#bins=np.linspace(0,100,20))
oe.phi_vector.sort_values().iloc[:20]
```

```
Out[26]: real_name
35      43.758830
10      66.299790
46      71.376560
12     118.219687
24     119.728221
8      140.461733
22     176.808587
21     200.494414
47     225.582200
30     255.737081
26     273.099728
44     290.829433
48     291.287982
5      321.916424
23     420.424853
38     430.891678
49     432.520905
29     522.375353
```

```

41      527.057867
33      566.113001
dtype: float64

```



Same as before, to get a “posterior” ensemble, we need to throw out the realizations with large phi - lets just take the 20 best:

```
In [27]: oe_pt = oe.loc[oe.phi_vector.sort_values().index[:20],:]
```

```
In [28]: f_df = pd.read_csv(os.path.join(m_d,"freyberg_pp.pred.usum.csv"),index_col=0)
f_df.index = f_df.index.map(str.lower)
f_df
```

```
Out [28]:
```

	prior_mean	prior_stdev	prior_lower_bound \
name			
fa_hw_19791230	-977.2390	295.32800	-1567.8900
fa_hw_19801229	-351.2160	409.77000	-1170.7600
fa_tw_19791230	-453.0330	409.35100	-1271.7400
fa_tw_19801229	108.9600	506.73200	-904.5040
hds_00_013_002_000	39.6102	3.96314	31.6840
hds_00_013_002_001	38.3838	4.05782	30.2681
part_status	2.0000	0.00000	2.0000
part_time	907.7020	570.98600	-234.2690

	prior_upper_bound	post_mean	post_stdev \
name			
fa_hw_19791230	-386.5840	-640.3370	250.802000

fa_hw_19801229	468.3240	-94.8227	339.971000
fa_tw_19791230	365.6690	-290.4730	170.616000
fa_tw_19801229	1122.4200	184.3570	263.999000
hds_00_013_002_000	47.5365	37.7893	0.274188
hds_00_013_002_001	46.4994	36.9350	0.687976
part_status	2.0000	2.0000	0.000000
part_time	2049.6700	966.2250	437.215000

	post_lower_bound	post_upper_bound
name		
fa_hw_19791230	-1141.9400	-138.7340
fa_hw_19801229	-774.7640	585.1190
fa_tw_19791230	-631.7040	50.7581
fa_tw_19801229	-343.6410	712.3540
hds_00_013_002_000	37.2409	38.3377
hds_00_013_002_001	35.5591	38.3110
part_status	2.0000	2.0000
part_time	91.7946	1840.6600

```
In [29]: obs = pst.observation_data
fnames = pst.pestpp_options["forecasts"].split(",")
for forecast in fnames:
    ax = plt.subplot(111)
    oe_pt.loc[:,forecast].hist(ax=ax,color="b",alpha=0.5,normed=True)
    ax.plot([obs.loc[forecast,"obsval"],obs.loc[forecast,"obsval"]],ax.get_ylim(),"r")
    axt = plt.twinx()
    x,y = pyemu.plot_utils.gaussian_distribution(f_df.loc[forecast,"prior_mean"],f_df.loc[forecast,"prior_std"])
    axt.fill_between(x,0,y,facecolor="0.5",alpha=0.25)
    x,y = pyemu.plot_utils.gaussian_distribution(f_df.loc[forecast,"post_mean"],f_df.loc[forecast,"post_std"])
    axt.fill_between(x,0,y,facecolor="b",alpha=0.25)
    axt.set_ylim(0,axt.get_ylim()[1])
    axt.set_yticks([])
    ax.set_title(forecast)
plt.show()
```