

pestpp-glm

May 3, 2019

1 PESTPP-GLM

In this notebook, we will run PESTPP-GLM in standard parameter estimation mode and regularization mode. In both cases, we will use the baked-in bayes-linear posterior monte carlo analysis to get posterior forecast PDFs. We will use the prior monte carlo outputs as the prior forecast PDF.

```
In [1]: import os
import shutil
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import flopy
import pyemu
```

flopy is installed in /Users/jeremyw/Dev/gw1876/activities_2day_mfm/notebooks/flopy

```
In [2]: t_d = "template"
m_d = "master_glm"
```

```
In [3]: pst = pyemu.Pst(os.path.join(t_d, "freyberg.pst"))
pst.write_par_summary_table(filename="none")
```

```
Out[3]:
```

	type	transform	count	initial value	\
gr_sy4	gr_sy4	log	705	0	
gr_prsity4	gr_prsity4	log	705	0	
cn_strt6	cn_strt6	log	1	0	
strk	strk	log	40	0	
cn_ss8	cn_ss8	log	1	0	
gr_rech3	gr_rech3	log	705	0	
gr_sy5	gr_sy5	log	705	0	
pp_rech1	pp_rech1	log	32	0	
gr_sy3	gr_sy3	log	705	0	
cn_sy8	cn_sy8	log	1	0	
gr_vka3	gr_vka3	log	705	0	
pp_hk1	pp_hk1	log	32	0	
cn_vka8	cn_vka8	log	1	0	

cn_hk7	cn_hk7	log	1	0
pp_prsity1	pp_prsity1	log	32	0
pp_ss0	pp_ss0	log	32	0
pp_strt0	pp_strt0	log	32	0
pp_hk0	pp_hk0	log	32	0
pp_ss1	pp_ss1	log	32	0
pp_hk2	pp_hk2	log	32	0
flow	flow	log	1	0
pp_ss2	pp_ss2	log	32	0
pp_rech0	pp_rech0	log	32	0
pp_prsity0	pp_prsity0	log	32	0
gr_strt4	gr_strt4	log	705	0
gr_strt5	gr_strt5	log	705	0
pp_prsity2	pp_prsity2	log	32	0
cn_rech5	cn_rech5	log	1	-0.39794
gr_rech2	gr_rech2	log	705	0
gr_hk3	gr_hk3	log	705	0
...
gr_ss5	gr_ss5	log	705	0
cn_strt8	cn_strt8	log	1	0
pp_strt2	pp_strt2	log	32	0
pp_vka0	pp_vka0	log	32	0
cn_ss6	cn_ss6	log	1	0
cn_hk8	cn_hk8	log	1	0
welflux_k02	welflux_k02	log	6	0
gr_strt3	gr_strt3	log	705	0
cn_vka7	cn_vka7	log	1	0
gr_vka4	gr_vka4	log	705	0
cn_prsity6	cn_prsity6	log	1	0
cn_prsity7	cn_prsity7	log	1	0
cn_sy6	cn_sy6	log	1	0
pp_sy0	pp_sy0	log	32	0
gr_ss3	gr_ss3	log	705	0
cn_hk6	cn_hk6	log	1	0
cn_sy7	cn_sy7	log	1	0
drncond_k00	drncond_k00	log	10	0
gr_hk5	gr_hk5	log	705	0
pp_sy1	pp_sy1	log	32	0
cn_prsity8	cn_prsity8	log	1	0
gr_prsity3	gr_prsity3	log	705	0
welflux	welflux	log	2	0 to 0.176091
gr_hk4	gr_hk4	log	705	0
gr_vka5	gr_vka5	log	705	0
cn_ss7	cn_ss7	log	1	0
pp_vka1	pp_vka1	log	32	0
pp_strt1	pp_strt1	log	32	0
cn_rech4	cn_rech4	log	1	0
gr_prsity5	gr_prsity5	log	705	0

	upper bound	lower bound	standard deviation
gr_sy4	0.243038	-0.60206	0.211275
gr_prsity4	0	-1	0.25
cn_strt6	0.0211893	-0.0222764	0.0108664
strk	2	-2	1
cn_ss8	1	-1	0.5
gr_rech3	0.0413927	-0.0457575	0.0217875
gr_sy5	0.243038	-0.60206	0.211275
pp_rech1	0.0413927	-0.0457575	0.0217875
gr_sy3	0.243038	-0.60206	0.211275
cn_sy8	0.243038	-0.60206	0.211275
gr_vka3	1	-1	0.5
pp_hk1	1	-1	0.5
cn_vka8	1	-1	0.5
cn_hk7	1	-1	0.5
pp_prsity1	0	-1	0.25
pp_ss0	1	-1	0.5
pp_strt0	0.0211893	-0.0222764	0.0108664
pp_hk0	1	-1	0.5
pp_ss1	1	-1	0.5
pp_hk2	1	-1	0.5
flow	0.09691	-0.124939	0.0554622
pp_ss2	1	-1	0.5
pp_rech0	0.0413927	-0.0457575	0.0217875
pp_prsity0	0	-1	0.25
gr_strt4	0.0211893	-0.0222764	0.0108664
gr_strt5	0.0211893	-0.0222764	0.0108664
pp_prsity2	0	-1	0.25
cn_rech5	-0.09691	-1	0.225772
gr_rech2	0.0413927	-0.0457575	0.0217875
gr_hk3	1	-1	0.5
...
gr_ss5	1	-1	0.5
cn_strt8	0.0211893	-0.0222764	0.0108664
pp_strt2	0.0211893	-0.0222764	0.0108664
pp_vka0	1	-1	0.5
cn_ss6	1	-1	0.5
cn_hk8	1	-1	0.5
welflux_k02	1	-1	0.5
gr_strt3	0.0211893	-0.0222764	0.0108664
cn_vka7	1	-1	0.5
gr_vka4	1	-1	0.5
cn_prsity6	0	-1	0.25
cn_prsity7	0	-1	0.25
cn_sy6	0.243038	-0.60206	0.211275
pp_sy0	0.243038	-0.60206	0.211275
gr_ss3	1	-1	0.5

cn_hk6	1	-1	0.5
cn_sy7	0.243038	-0.60206	0.211275
drncond_k00	1	-1	0.5
gr_hk5	1	-1	0.5
pp_sy1	0.243038	-0.60206	0.211275
cn_prsity8	0	-1	0.25
gr_prsity3	0	-1	0.25
welflux	0.176091 to 0.30103	-0.30103 to 0	0.0752575 to 0.11928
gr_hk4	1	-1	0.5
gr_vka5	1	-1	0.5
cn_ss7	1	-1	0.5
pp_vka1	1	-1	0.5
pp_strt1	0.0211893	-0.0222764	0.0108664
cn_rech4	0.0791812	-0.09691	0.0440228
gr_prsity5	0	-1	0.25

[65 rows x 7 columns]

1.0.1 reduce the number of adjustable parameters

This is the painful part: we cant use 10K+ pars because we cant wait around for that many runs and then the linear algebra of factoring a 10k+ by 10K+ matrix is also difficult. So that means we need to fix a lot a parameters #frownyface

```
In [4]: par = pst.parameter_data
```

```
In [5]: # grid-scale pars
gr_pars = par.loc[par.pargp.apply(lambda x: "gr" in x), "parnme"]
par.loc[gr_pars, "partrans"] = "fixed"
pst.npar_adj
```

```
Out[5]: 719
```

```
In [6]: # these are the sfr conductance parameters - Ive left all 40 adjustable
# but if you uncomment this, it will tie them into 1 parameter effectively
# strk_pars = par.loc[par.pargp=="strk", "parnme"]
# p1 = strk_pars.iloc[0]
# par.loc[strk_pars.iloc[1:], "partrans"] = "tied"
# par.loc[strk_pars.iloc[1:], "partied"] = p1
pst.npar_adj
```

```
Out[6]: 719
```

```
In [7]: par.loc[par.pargp.apply(lambda x: "pp" in x), "pargp"].unique()
```

```
Out[7]: array(['pp_hk0', 'pp_hk1', 'pp_hk2', 'pp_prsity0', 'pp_prsity1',
               'pp_prsity2', 'pp_rech0', 'pp_rech1', 'pp_ss0', 'pp_ss1', 'pp_ss2',
               'pp_strt0', 'pp_strt1', 'pp_strt2', 'pp_sy0', 'pp_sy1', 'pp_sy2',
               'pp_vka0', 'pp_vka1', 'pp_vka2'], dtype=object)
```

Fix the storage pilot points - we still have layer-scale storage pars adjustable

```
In [8]: #s_pars = par.loc[par.pargp.apply(lambda x: "pp" in x and ("ss" in x or "sy" in x)), "p"]  
        #par.loc[s_pars, "partrans"] = "fixed"  
        pst.npar_adj
```

```
Out[8]: 719
```

```
In [9]: adj_par = par.loc[par.partrans=="log", :]  
        adj_par.pargp.value_counts().sort_values()
```

```
Out[9]: cn_strt6      1  
        cn_rech4      1  
        cn_hk7        1  
        cn_sy8         1  
        cn_sy7         1  
        cn_ss8         1  
        cn_vka8         1  
        cn_vka6         1  
        flow           1  
        cn_rech5       1  
        cn_vka7         1  
        cn_strt7        1  
        cn_strt8        1  
        cn_hk8          1  
        cn_ss6          1  
        cn_prsity7      1  
        cn_sy6          1  
        cn_hk6          1  
        cn_prsity6      1  
        cn_prsity8      1  
        cn_ss7          1  
        welflux         2  
        welflux_k02     6  
        drncond_k00     10  
        pp_strt0        32  
        pp_rech0        32  
        pp_hk0          32  
        pp_sy2          32  
        pp_vka2         32  
        pp_ss0          32  
        pp_vka0         32  
        pp_sy0          32  
        pp_prsity2      32  
        pp_hk2          32  
        pp_strt2        32  
        pp_ss1          32  
        pp_ss2          32  
        pp_prsity1      32
```

```

pp_rech1      32
pp_sy1        32
pp_prsity0    32
pp_vka1       32
pp_strt1      32
pp_hk1        32
strk          40
Name: pargp, dtype: int64

```

fix the future recharge pilot points, vka in layers 1 and 3 and the initial condition pilot points (we still have layer-scale pars for each of these types)

```

In [10]: fi_grps = ["pp_rech1", "pp_vka0", "pp_vka2", "pp_strt0", "pp_strt1", "pp_strt2"]
          par.loc[par.pargp.apply(lambda x: x in fi_grps), "partrans"] = "fixed"
          pst.npar_adj

```

```

Out[10]: 527

```

Ok, thats better...so lets run PESTPP-GLM. We will use a single "base parameter" jacobian matrix as the basis for 6 super parameter iterations. Then we will draw 100 realizations from the FOSM posterior parameter covariance matrix and run those 100 realizations to get the psoterior forecast PDFs

```

In [11]: pst.control_data.noptmax = 3
          pst.pestpp_options["n_iter_base"] = -1
          pst.pestpp_options["n_iter_super"] = 3
          pst.pestpp_options["num_reals"] = 50 # this is how many ies uses
          pst.pestpp_options["parcov"] = "prior_cov.jcb"
          pst.write(os.path.join(t_d, "freyberg_pp.pst"))

In [12]: #pyemu.os_utils.start_slaves(t_d, "pestpp-glm", "freyberg_pp.pst", num_slaves=20, slave_r
          #                                     master_dir=m_d)

In [13]: df = df=pd.read_csv(os.path.join(m_d, "freyberg_pp.post.obsen.csv"), index_col=0)
          oe = pyemu.ObservationEnsemble.from_dataframe(pst=pst, df=df)

In [14]: ax = oe.phi_vector.hist()#bins=np.linspace(0, 100, 20))
          oe.phi_vector.sort_values().iloc[:20]

```

```


Out[14]: real_name
7      181.835972
23     307.875113
25     461.423208
21     508.389022
44     531.916245
19     710.543355
6      722.088590
45     882.682950
8      947.056165

```

```

9      1055.322092
24     1124.430746
33     1472.494248
11     1591.305572
17     1832.739231
34     2077.463682
30     2169.778154
27     2339.569383
46     2498.302887
49     2611.632003
41     2715.695479
dtype: float64

```



Here we see the distribution of phi values across the 100 posterior realizations. Should we accept all of these??? The theoretical phi we should accept is number of nonzero obs (14).

To get a “posterior” ensemble, we need to throw out the realizations with large phi - lets just take the 20 best:

```
In [15]: oe_pt = oe.loc[oe.phi_vector.sort_values().index[:20],:] #just take the 20 lowest phi
```

We can also load and plot the FOSM forecast results along side of the ensemble results:

```
In [16]: f_df = pd.read_csv(os.path.join(m_d,"freyberg_pp.pred.usum.csv"),index_col=0)
         f_df.index = f_df.index.map(str.lower)
         f_df
```

```
Out[16]:
```


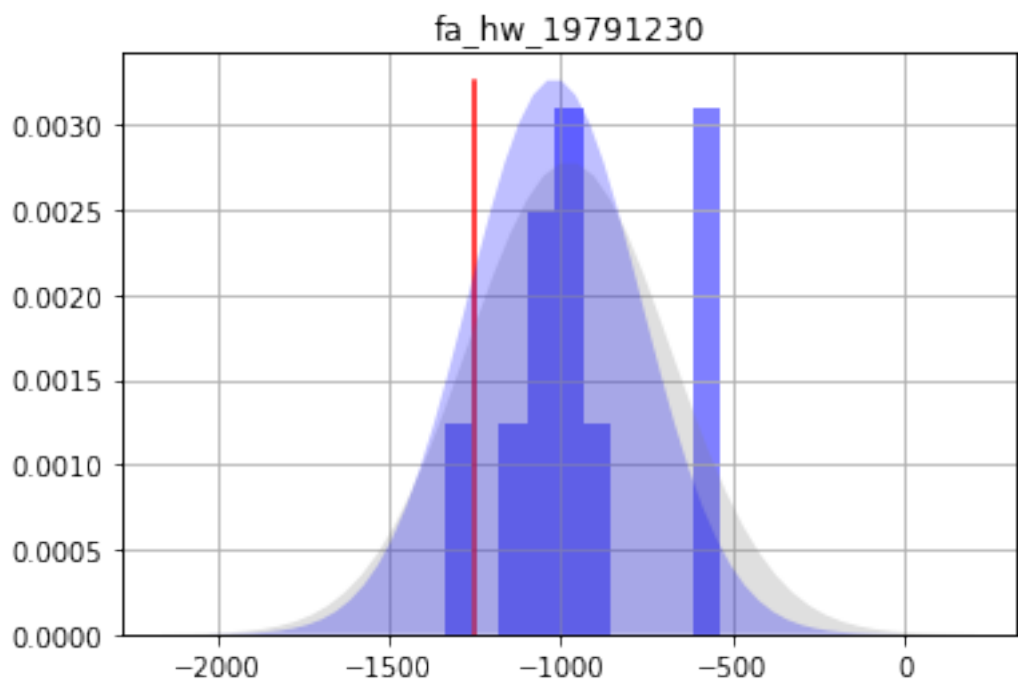
	prior_mean	prior_stdev	prior_lower_bound	\
name				



fa_hw_19791230	-977.2390	295.32800	-1567.8900
fa_hw_19801229	-351.2160	409.77000	-1170.7600
fa_tw_19791230	-453.0330	409.35100	-1271.7400
fa_tw_19801229	108.9600	506.73200	-904.5040
hds_00_013_002_000	39.6102	3.96314	31.6840
hds_00_013_002_001	38.3838	4.05782	30.2681
part_status	2.0000	0.00000	2.0000
part_time	907.7020	704.75100	-501.8010



	prior_upper_bound	post_mean	post_stdev \
name			
fa_hw_19791230	-386.5840	-1019.0700	251.100000
fa_hw_19801229	468.3240	-377.4560	340.542000
fa_tw_19791230	365.6690	-329.7400	175.584000
fa_tw_19801229	1122.4200	234.8450	268.440000
hds_00_013_002_000	47.5365	39.1429	0.295060
hds_00_013_002_001	46.4994	37.8161	0.696415
part_status	2.0000	2.0000	0.000000
part_time	2317.2000	1348.2500	603.218000


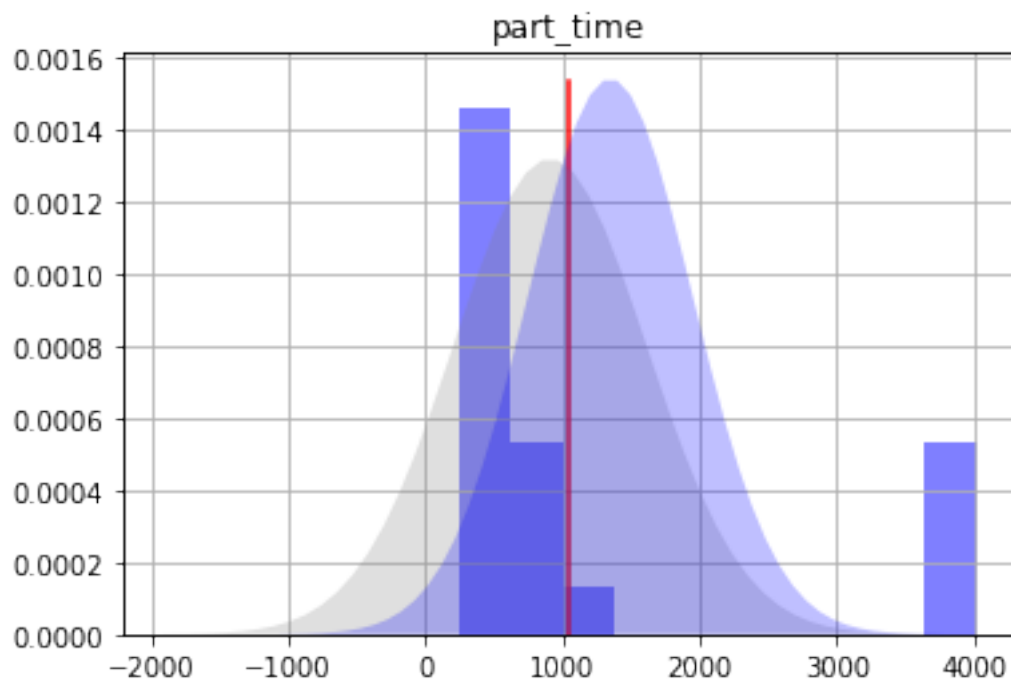
	post_lower_bound	post_upper_bound
name		
fa_hw_19791230	-1521.2700	-516.8720
fa_hw_19801229	-1058.5400	303.6270
fa_tw_19791230	-680.9080	21.4276
fa_tw_19801229	-302.0350	771.7250
hds_00_013_002_000	38.5528	39.7330
hds_00_013_002_001	36.4233	39.2090
part_status	2.0000	2.0000
part_time	141.8160	2554.6900

```
In [17]: obs = pst.observation_data
fnames = pst.pestpp_options["forecasts"].split(",")
for forecast in fnames:
    ax = plt.subplot(111)
    oe_pt.loc[:,forecast].hist(ax=ax,color="b",alpha=0.5,normed=True)
    ax.plot([obs.loc[forecast,"obsval"],obs.loc[forecast,"obsval"]],ax.get_ylim(),"r")
    axt = plt.twinx()
    x,y = pyemu.plot_utils.gaussian_distribution(f_df.loc[forecast,"prior_mean"],f_df.loc[forecast,"prior_stdev"])
    axt.fill_between(x,0,y,facecolor="0.5",alpha=0.25)
    x,y = pyemu.plot_utils.gaussian_distribution(f_df.loc[forecast,"post_mean"],f_df.loc[forecast,"post_stdev"])
    axt.fill_between(x,0,y,facecolor="b",alpha=0.25)
    axt.set_ylim(0,axt.get_ylim()[1])
    axt.set_yticks([])
    ax.set_title(forecast)
plt.show()
```







1.0.2 Setup of Tikhonov regularization

Now lets setup and use some formal regularization to bring the final phi up to around 14. We will use first-order regularization based on the covariance matrix we build earlier:

```
In [18]: cov = pyemu.Cov.from_binary(os.path.join(t_d,"prior_cov.jcb"))
```

```
new binary format detected...
```

```
In [19]: pyemu.helpers.first_order_pearson_tikhonov(pst,cov)
```

```
getting CC matrix  
processing
```

```
In [20]: pst.prior_information.head()
```

```
Out[20]:
```

	equation	obgnme	\
pilbl			
pcc_1	1.0 * log(dc0000390005) - 1.0 * log(dc0000390006) = 0.0	regul_cc	
pcc_2	1.0 * log(dc0000390005) - 1.0 * log(dc0000390007) = 0.0	regul_cc	
pcc_3	1.0 * log(dc0000390005) - 1.0 * log(dc0000390008) = 0.0	regul_cc	
pcc_4	1.0 * log(dc0000390005) - 1.0 * log(dc0000390009) = 0.0	regul_cc	
pcc_5	1.0 * log(dc0000390005) - 1.0 * log(dc0000390010) = 0.0	regul_cc	

	pilbl	weight
pilbl		
pcc_1	pcc_1	0.904837
pcc_2	pcc_2	0.818731
pcc_3	pcc_3	0.740818
pcc_4	pcc_4	0.670320
pcc_5	pcc_5	0.606531

```
In [21]: shutil.copy2(os.path.join(m_d,"freyberg_pp.jcb"),os.path.join(t_d,"restart_pp.jcb"))
```

```
Out[21]: 'template/restart_pp.jcb'
```

```
In [22]: pst.pestpp_options["base_jacobian"] = "restart_pp.jcb"  
pst.reg_data.phimlim = pst.nnz_obs  
pst.reg_data.phimaccept = pst.reg_data.phimlim * 1.1  
pst.write(os.path.join(t_d,"freyberg_pp.pst"))
```

```
In [23]: pyemu.os_utils.start_slaves(t_d,"pestpp-glm","freyberg_pp.pst",num_slaves=20,slave_ro  
master_dir=m_d)
```


```
In [24]: df = df=pd.read_csv(os.path.join(m_d,"freyberg_pp.post.obsen.csv"),index_col=0)  
oe = pyemu.ObservationEnsemble.from_dataframe(pst=pst,df=df)
```

```
In [25]: ax = oe.phi_vector.hist()  
oe.phi_vector.sort_values().iloc[:20]
```

```

Out [25]: real_name
26      93.513954
6      217.293403
23     226.100858
33     232.205306
22     308.050204
25     347.906426
45     429.314289
32     482.540726
44     557.405461
46     570.260765
47     595.014360
40     623.244712
4      638.025129
17     671.096086
15     676.631334
11     781.666502
5      834.788458
48    1223.998840
12    1249.309590
37    1631.682572
dtype: float64

```



Same as before, to get a “posterior” ensemble, we need to throw out the realizations with large phi - lets just take the 20 best:

```

In [26]: oe_pt = oe.loc[oe.phi_vector.sort_values().index[:20],:]

```

```
In [27]: f_df = pd.read_csv(os.path.join(m_d,"freyberg_pp.pred.usum.csv"),index_col=0)
f_df.index = f_df.index.map(str.lower)
f_df
```

```
Out [27]:
```

	prior_mean	prior_stdev	prior_lower_bound \
name			
fa_hw_19791230	-977.2390	295.32800	-1567.8900
fa_hw_19801229	-351.2160	409.77000	-1170.7600
fa_tw_19791230	-453.0330	409.35100	-1271.7400
fa_tw_19801229	108.9600	506.73200	-904.5040
hds_00_013_002_000	39.6102	3.96314	31.6840
hds_00_013_002_001	38.3838	4.05782	30.2681
part_status	2.0000	0.00000	2.0000
part_time	907.7020	704.75100	-501.8010

	prior_upper_bound	post_mean	post_stdev \
name			
fa_hw_19791230	-386.5840	-931.1650	247.835000
fa_hw_19801229	468.3240	-438.0700	335.549000
fa_tw_19791230	365.6690	-242.9430	170.577000
fa_tw_19801229	1122.4200	272.6840	263.671000
hds_00_013_002_000	47.5365	39.1357	0.273990
hds_00_013_002_001	46.4994	38.0574	0.687833
part_status	2.0000	2.0000	0.000000
part_time	2317.2000	1060.4200	601.307000

	post_lower_bound	post_upper_bound
name		
fa_hw_19791230	-1426.8300	-435.4950
fa_hw_19801229	-1109.1700	233.0270
fa_tw_19791230	-584.0980	98.2112
fa_tw_19801229	-254.6580	800.0260
hds_00_013_002_000	38.5877	39.6836
hds_00_013_002_001	36.6818	39.4331
part_status	2.0000	2.0000
part_time	-142.1900	2263.0400

```
In [28]: obs = pst.observation_data
fnames = pst.pestpp_options["forecasts"].split(",")
for forecast in fnames:
    ax = plt.subplot(111)
    oe_pt.loc[:,forecast].hist(ax=ax,color="b",alpha=0.5,normed=True)
    ax.plot([obs.loc[forecast,"obsval"],obs.loc[forecast,"obsval"]],ax.get_ylim(),"r")
    axt = plt.twinx()
    x,y = pyemu.plot_utils.gaussian_distribution(f_df.loc[forecast,"prior_mean"],f_df.
    axt.fill_between(x,0,y,facecolor="0.5",alpha=0.25)
    x,y = pyemu.plot_utils.gaussian_distribution(f_df.loc[forecast,"post_mean"],f_df.
    axt.fill_between(x,0,y,facecolor="b",alpha=0.25)
```

```

axt.set_ylim(0,axt.get_ylim()[1])
axt.set_yticks([])
ax.set_title(forecast)
plt.show()

```

