

pestpp-glm

May 7, 2019

1 PESTPP-GLM

In this notebook, we will run PESTPP-GLM in standard parameter estimation mode and regularization mode. In both cases, we will use the baked-in bayes-linear posterior monte carlo analysis to get posterior forecast PDFs. We will use the prior monte carlo outputs as the prior forecast PDF.

```
In [1]: import os
import shutil
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib as mpl
plt.rcParams['font.size']=12
import flopy
import pyemu
```

flopy is installed in /Users/jeremyw/Dev/gw1876/activities_2day_mfm/notebooks/flopy

```
In [2]: t_d = "template"
m_d = "master_glm"
```

```
In [3]: pst = pyemu.Pst(os.path.join(t_d,"freyberg.pst"))
pst.write_par_summary_table(filename="none")
```

```
Out[3]:
```

	type	transform	count	initial value	\
gr_prsity4	gr_prsity4	log	705	0	
cn_sy8	cn_sy8	log	1	0	
cn_ss8	cn_ss8	log	1	0	
gr_strt4	gr_strt4	log	705	0	
gr_rech3	gr_rech3	log	705	0	
pp_hk2	pp_hk2	log	32	0	
gr_ss3	gr_ss3	log	705	0	
gr_hk3	gr_hk3	log	705	0	
pp_ss2	pp_ss2	log	32	0	
gr_sy5	gr_sy5	log	705	0	
gr_ss5	gr_ss5	log	705	0	

cn_strt6	cn_strt6	log	1	0
gr_prsity5	gr_prsity5	log	705	0
pp_sy2	pp_sy2	log	32	0
pp_hk0	pp_hk0	log	32	0
strk	strk	log	40	0
cn_sy6	cn_sy6	log	1	0
gr_strt5	gr_strt5	log	705	0
gr_prsity3	gr_prsity3	log	705	0
drncond_k00	drncond_k00	log	10	0
cn_prsity7	cn_prsity7	log	1	0
cn_vka6	cn_vka6	log	1	0
gr_rech2	gr_rech2	log	705	0
gr_sy3	gr_sy3	log	705	0
pp_strt0	pp_strt0	log	32	0
pp_strt2	pp_strt2	log	32	0
pp_hk1	pp_hk1	log	32	0
gr_strt3	gr_strt3	log	705	0
gr_sy4	gr_sy4	log	705	0
gr_vka3	gr_vka3	log	705	0
...
pp_strt1	pp_strt1	log	32	0
cn_rech5	cn_rech5	log	1	-0.39794
cn_strt8	cn_strt8	log	1	0
flow	flow	log	1	0
gr_ss4	gr_ss4	log	705	0
cn_rech4	cn_rech4	log	1	0
cn_vka7	cn_vka7	log	1	0
gr_vka5	gr_vka5	log	705	0
gr_hk4	gr_hk4	log	705	0
cn_prsity8	cn_prsity8	log	1	0
pp_prsity2	pp_prsity2	log	32	0
pp_prsity0	pp_prsity0	log	32	0
pp_vka2	pp_vka2	log	32	0
cn_hk6	cn_hk6	log	1	0
welflux_k02	welflux_k02	log	6	0
cn_hk7	cn_hk7	log	1	0
cn_vka8	cn_vka8	log	1	0
gr_hk5	gr_hk5	log	705	0
cn_ss7	cn_ss7	log	1	0
pp_sy0	pp_sy0	log	32	0
cn_hk8	cn_hk8	log	1	0
pp_vka1	pp_vka1	log	32	0
pp_prsity1	pp_prsity1	log	32	0
cn_strt7	cn_strt7	log	1	0
pp_rech1	pp_rech1	log	32	0
cn_sy7	cn_sy7	log	1	0
cn_ss6	cn_ss6	log	1	0
pp_ss0	pp_ss0	log	32	0

pp_vka0	pp_vka0	log	32	0
welflux	welflux	log	2	0 to 0.176091

	upper bound	lower bound	standard deviation
gr_prsity4	0	-1	0.25
cn_sy8	0.243038	-0.60206	0.211275
cn_ss8	1	-1	0.5
gr_strt4	0.0211893	-0.0222764	0.0108664
gr_rech3	0.0413927	-0.0457575	0.0217875
pp_hk2	1	-1	0.5
gr_ss3	1	-1	0.5
gr_hk3	1	-1	0.5
pp_ss2	1	-1	0.5
gr_sy5	0.243038	-0.60206	0.211275
gr_ss5	1	-1	0.5
cn_strt6	0.0211893	-0.0222764	0.0108664
gr_prsity5	0	-1	0.25
pp_sy2	0.243038	-0.60206	0.211275
pp_hk0	1	-1	0.5
strk	2	-2	1
cn_sy6	0.243038	-0.60206	0.211275
gr_strt5	0.0211893	-0.0222764	0.0108664
gr_prsity3	0	-1	0.25
drncond_k00	1	-1	0.5
cn_prsity7	0	-1	0.25
cn_vka6	1	-1	0.5
gr_rech2	0.0413927	-0.0457575	0.0217875
gr_sy3	0.243038	-0.60206	0.211275
pp_strt0	0.0211893	-0.0222764	0.0108664
pp_strt2	0.0211893	-0.0222764	0.0108664
pp_hk1	1	-1	0.5
gr_strt3	0.0211893	-0.0222764	0.0108664
gr_sy4	0.243038	-0.60206	0.211275
gr_vka3	1	-1	0.5
...
pp_strt1	0.0211893	-0.0222764	0.0108664
cn_rech5	-0.09691	-1	0.225772
cn_strt8	0.0211893	-0.0222764	0.0108664
flow	0.09691	-0.124939	0.0554622
gr_ss4	1	-1	0.5
cn_rech4	0.0791812	-0.09691	0.0440228
cn_vka7	1	-1	0.5
gr_vka5	1	-1	0.5
gr_hk4	1	-1	0.5
cn_prsity8	0	-1	0.25
pp_prsity2	0	-1	0.25
pp_prsity0	0	-1	0.25
pp_vka2	1	-1	0.5

cn_hk6	1	-1	0.5
welflux_k02	1	-1	0.5
cn_hk7	1	-1	0.5
cn_vka8	1	-1	0.5
gr_hk5	1	-1	0.5
cn_ss7	1	-1	0.5
pp_sy0	0.243038	-0.60206	0.211275
cn_hk8	1	-1	0.5
pp_vka1	1	-1	0.5
pp_prsity1	0	-1	0.25
cn_strt7	0.0211893	-0.0222764	0.0108664
pp_rech1	0.0413927	-0.0457575	0.0217875
cn_sy7	0.243038	-0.60206	0.211275
cn_ss6	1	-1	0.5
pp_ss0	1	-1	0.5
pp_vka0	1	-1	0.5
welflux	0.176091 to 0.30103	-0.30103 to 0	0.0752575 to 0.11928

[65 rows x 7 columns]

1.0.1 reduce the number of adjustable parameters

This is the painful part: we cant use 10K+ pars because we cant wait around for that many runs and then the linear algebra of factoring a 10k+ by 10K+ matrix is also difficult. So that means we need to fix a lot a parameters #frownyface

```
In [4]: par = pst.parameter_data
```

```
In [5]: # grid-scale pars
```

```
gr_pars = par.loc[par.pargp.apply(lambda x: "gr" in x), "parnme"]
par.loc[gr_pars, "partrans"] = "fixed"
pst.npar_adj
```

```
Out[5]: 719
```

```
In [6]: # these are the sfr conductance parameters - Ive left all 40 adjustable
# but if you uncomment this, it will tie them into 1 parameter effectively
# strk_pars = par.loc[par.pargp=="strk", "parnme"]
# p1 = strk_pars.iloc[0]
# par.loc[strk_pars.iloc[1:], "partrans"] = "tied"
# par.loc[strk_pars.iloc[1:], "partied"] = p1
pst.npar_adj
```

```
Out[6]: 719
```

```
In [7]: par.loc[par.pargp.apply(lambda x: "pp" in x), "pargp"].unique()
```

```
Out[7]: array(['pp_hk0', 'pp_hk1', 'pp_hk2', 'pp_prsity0', 'pp_prsity1',
               'pp_prsity2', 'pp_rech0', 'pp_rech1', 'pp_ss0', 'pp_ss1', 'pp_ss2',
               'pp_strt0', 'pp_strt1', 'pp_strt2', 'pp_sy0', 'pp_sy1', 'pp_sy2',
               'pp_vka0', 'pp_vka1', 'pp_vka2'], dtype=object)
```

Fix the storage pilot points - we still have layer-scale storage pars adjustable

```
In [8]: #s_pars = par.loc[par.pargp.apply(lambda x: "pp" in x and ("ss" in x or "sy" in x)), "p"]  
        #par.loc[s_pars, "partrans"] = "fixed"  
        pst.npar_adj
```

```
Out[8]: 719
```

```
In [9]: adj_par = par.loc[par.partrans=="log", :]  
        adj_par.pargp.value_counts().sort_values()
```

```
Out[9]: cn_sy8          1  
        cn_ss8          1  
        cn_strt6         1  
        cn_strt8         1  
        cn_rech4         1  
        cn_prsity7        1  
        cn_prsity6        1  
        cn_sy6           1  
        cn_vka6           1  
        cn_vka7           1  
        cn_prsity8        1  
        flow             1  
        cn_hk6           1  
        cn_hk7           1  
        cn_vka8           1  
        cn_rech5         1  
        cn_hk8           1  
        cn_strt7         1  
        cn_sy7           1  
        cn_ss7           1  
        cn_ss6           1  
        welflux          2  
        welflux_k02       6  
        drncond_k00      10  
        pp_hk2           32  
        pp_strt2          32  
        pp_hk0           32  
        pp_hk1           32  
        pp_vka2          32  
        pp_rech0          32  
        pp_strt1          32  
        pp_vka0          32  
        pp_prsity0        32  
        pp_ss0           32  
        pp_rech1          32  
        pp_vka1          32  
        pp_sy2           32  
        pp_strt0          32
```

```

pp_sy1      32
pp_sy0      32
pp_ss2      32
pp_prsity1   32
pp_ss1      32
pp_prsity2   32
strk        40
Name: pargp, dtype: int64

```

fix the future recharge pilot points, vka in layers 1 and 3 and the initial condition pilot points (we still have layer-scale pars for each of these types)

```

In [10]: fi_grps = ["pp_rech1", "pp_vka0", "pp_vka2", "pp_strt0", "pp_strt1", "pp_strt2"]
         par.loc[par.pargp.apply(lambda x: x in fi_grps), "partrans"] = "fixed"
         pst.npar_adj

```

```

Out[10]: 527

```

Ok, thats better...so lets run PESTPP-GLM. We will use a single "base parameter" jacobian matrix as the basis for 6 super parameter iterations. Then we will draw 100 realizations from the FOSM posterior parameter covariance matrix and run those 100 realizations to get the psoterior forecast PDFs

```

In [11]: pst.control_data.noptmax = 3
         pst.pestpp_options["n_iter_base"] = -1
         pst.pestpp_options["n_iter_super"] = 3
         pst.pestpp_options["num_reals"] = 50 # this is how many ies uses
         pst.pestpp_options["parcov"] = "prior_cov.jcb"
         pst.write(os.path.join(t_d, "freyberg_pp.pst"))

```

```

In [12]: pyemu.os_utils.start_slaves(t_d, "pestpp-glm", "freyberg_pp.pst", num_slaves=20, slave_ro
         master_dir=m_d)

```

```

In [13]: df = df=pd.read_csv(os.path.join(m_d, "freyberg_pp.post.obsen.csv"), index_col=0)
         oe = pyemu.ObservationEnsemble.from_dataframe(pst=pst, df=df)

```

```

In [14]: ax = oe.phi_vector.hist() #bins=np.linspace(0, 100, 20))
         oe.phi_vector.sort_values().iloc[:20]

```

```


Out[14]: real_name
7      55.468202
24     84.260295
4     124.191279
1     168.977488
25    174.596536
21    199.858942
39    205.571352
20    373.756778
42    385.789083

```

```

12    428.295418
6     429.258788
16    430.476563
45    462.553423
19    495.599815
41    510.791903
22    512.655947
48    552.963192
33    573.280556
3     615.596134
37    772.791092
dtype: float64

```



Here we see the distribution of phi values across the 100 posterior realizations. Should we accept all of these??? The theoretical phi we should accept is number of nonzero obs (14).

To get a “posterior” ensemble, we need to throw out the realizations with large phi - lets just take the 20 best:

```
In [15]: oe_pt = oe.loc[oe.phi_vector.sort_values().index[:20],:] #just take the 20 lowest phi
```

We can also load and plot the FOSM forecast results along side of the ensemble results:

```
In [16]: f_df = pd.read_csv(os.path.join(m_d,"freyberg_pp.pred.usum.csv"),index_col=0)
         f_df.index = f_df.index.map(str.lower)
         f_df
```

```
Out[16]:
```



	prior_mean	prior_stdev	prior_lower_bound	\
name				



fa_hw_19791230	-977.2390	295.32800	-1567.8900
fa_hw_19801229	-351.2160	409.77000	-1170.7600
fa_tw_19791230	-453.0330	409.35100	-1271.7400
fa_tw_19801229	108.9600	506.73200	-904.5040
hds_00_013_002_000	39.6102	3.96314	31.6840
hds_00_013_002_001	38.3838	4.05782	30.2681
part_status	2.0000	0.00000	2.0000
part_time	907.7020	704.75100	-501.8010



	prior_upper_bound	post_mean	post_stdev \
name			
fa_hw_19791230	-386.5840	-811.3660	247.865000
fa_hw_19801229	468.3240	-300.1030	335.551000
fa_tw_19791230	365.6690	-103.5390	170.577000
fa_tw_19801229	1122.4200	366.6750	263.681000
hds_00_013_002_000	47.5365	39.4199	0.273990
hds_00_013_002_001	46.4994	38.1797	0.687836
part_status	2.0000	2.0000	0.000000
part_time	2317.2000	730.7160	601.308000



	post_lower_bound	post_upper_bound
name		
fa_hw_19791230	-1307.1000	-315.6370
fa_hw_19801229	-971.2060	371.0000
fa_tw_19791230	-444.6930	237.6160
fa_tw_19801229	-160.6870	894.0380
hds_00_013_002_000	38.8719	39.9679
hds_00_013_002_001	36.8040	39.5553
part_status	2.0000	2.0000
part_time	-471.9010	1933.3300

```
In [17]: obs = pst.observation_data
fnames = pst.pestpp_options["forecasts"].split(",")
for forecast in fnames:
    ax = plt.subplot(111)
    oe_pt.loc[:,forecast].hist(ax=ax,color="b",alpha=0.5,normed=True)
    ax.plot([obs.loc[forecast,"obsval"],obs.loc[forecast,"obsval"]],ax.get_ylim(),"r")
    axt = plt.twinx()
    x,y = pyemu.plot_utils.gaussian_distribution(f_df.loc[forecast,"prior_mean"],f_df.loc[forecast,"prior_stdev"])
    axt.fill_between(x,0,y,facecolor="0.5",alpha=0.25)
    x,y = pyemu.plot_utils.gaussian_distribution(f_df.loc[forecast,"post_mean"],f_df.loc[forecast,"post_stdev"])
    axt.fill_between(x,0,y,facecolor="b",alpha=0.25)
    axt.set_ylim(0,axt.get_ylim()[1])
    axt.set_yticks([])
    ax.set_title(forecast)
plt.show()
```







1.0.2 Setup of Tikhonov regularization

Now lets setup and use some formal regularization to bring the final phi up to around 14. We will use first-order regularization based on the covariance matrix we build earlier:

```
In [18]: cov = pyemu.Cov.from_binary(os.path.join(t_d,"prior_cov.jcb"))
```

```
new binary format detected...
```

```
In [19]: pyemu.helpers.first_order_pearson_tikhonov(pst,cov)
```

```
getting CC matrix  
processing
```

```
In [20]: pst.prior_information.head()
```

```
Out[20]:
```

	equation	obgnme	\
pilbl			
pcc_1	1.0 * log(dc0000390005) - 1.0 * log(dc0000390006) = 0.0	regul_cc	
pcc_2	1.0 * log(dc0000390005) - 1.0 * log(dc0000390007) = 0.0	regul_cc	
pcc_3	1.0 * log(dc0000390005) - 1.0 * log(dc0000390008) = 0.0	regul_cc	
pcc_4	1.0 * log(dc0000390005) - 1.0 * log(dc0000390009) = 0.0	regul_cc	
pcc_5	1.0 * log(dc0000390005) - 1.0 * log(dc0000390010) = 0.0	regul_cc	

	pilbl	weight
pilbl		
pcc_1	pcc_1	0.904837
pcc_2	pcc_2	0.818731
pcc_3	pcc_3	0.740818
pcc_4	pcc_4	0.670320
pcc_5	pcc_5	0.606531

```
In [21]: shutil.copy2(os.path.join(m_d,"freyberg_pp.jcb"),os.path.join(t_d,"restart_pp.jcb"))
```

```
Out[21]: 'template/restart_pp.jcb'
```


```
In [22]: pst.pestpp_options["base_jacobian"] = "restart_pp.jcb"  
pst.reg_data.phimlim = pst.nnz_obs  
pst.reg_data.phimaccept = pst.reg_data.phimlim * 1.1  
pst.write(os.path.join(t_d,"freyberg_pp.pst"))
```

```
In [23]: pyemu.os_utils.start_slaves(t_d,"pestpp-glm","freyberg_pp.pst",num_slaves=20,slave_ro  
master_dir=m_d)
```

```
In [24]: df = df=pd.read_csv(os.path.join(m_d,"freyberg_pp.post.obsen.csv"),index_col=0)  
oe = pyemu.ObservationEnsemble.from_dataframe(pst=pst,df=df)
```

```
In [25]: ax = oe.phi_vector.hist()#bins=np.linspace(0,100,20))  
oe.phi_vector.sort_values().iloc[:20]
```

```
Out [25]: real_name
7      21.906757
24     94.064807
25    130.642360
44    144.623771
1     171.822277
39    173.064244
19    189.988899
12    332.255646
21    346.544001
42    377.272242
4     408.480862
15    466.355447
48    556.091368
6     641.565693
26    683.294694
16    687.194301
22    696.663817
45    761.306957
33    792.380236
20    844.633515
dtype: float64
```



Same as before, to get a “posterior” ensemble, we need to throw out the realizations with large phi - lets just take the 20 best:

```
In [26]: oe_pt = oe.loc[oe.phi_vector.sort_values().index[:20],:]
```

```
In [27]: f_df = pd.read_csv(os.path.join(m_d,"freyberg_pp.pred.usum.csv"),index_col=0)
f_df.index = f_df.index.map(str.lower)
f_df
```

```
Out [27]:
```

	prior_mean	prior_stdev	prior_lower_bound \
name			
fa_hw_19791230	-977.2390	295.32800	-1567.8900
fa_hw_19801229	-351.2160	409.77000	-1170.7600
fa_tw_19791230	-453.0330	409.35100	-1271.7400
fa_tw_19801229	108.9600	506.73200	-904.5040
hds_00_013_002_000	39.6102	3.96314	31.6840
hds_00_013_002_001	38.3838	4.05782	30.2681
part_status	2.0000	0.00000	2.0000
part_time	907.7020	704.75100	-501.8010

	prior_upper_bound	post_mean	post_stdev \
name			
fa_hw_19791230	-386.5840	-650.4570	248.019000
fa_hw_19801229	468.3240	-254.1730	335.584000
fa_tw_19791230	365.6690	-57.8134	170.817000
fa_tw_19801229	1122.4200	279.3360	263.851000
hds_00_013_002_000	47.5365	39.4960	0.277776
hds_00_013_002_001	46.4994	38.4735	0.689849
part_status	2.0000	2.0000	0.000000
part_time	2317.2000	1003.5700	603.650000

	post_lower_bound	post_upper_bound
name		
fa_hw_19791230	-1146.5000	-154.4190
fa_hw_19801229	-925.3410	416.9940
fa_tw_19791230	-399.4470	283.8200
fa_tw_19801229	-248.3660	807.0390
hds_00_013_002_000	38.9405	40.0516
hds_00_013_002_001	37.0938	39.8532
part_status	2.0000	2.0000
part_time	-203.7280	2210.8700

```
In [28]: obs = pst.observation_data
fnames = pst.pestpp_options["forecasts"].split(",")
for forecast in fnames:
    ax = plt.subplot(111)
    oe_pt.loc[:,forecast].hist(ax=ax,color="b",alpha=0.5,normed=True)
    ax.plot([obs.loc[forecast,"obsval"],obs.loc[forecast,"obsval"]],ax.get_ylim(),"r")
    axt = plt.twinx()
    x,y = pyemu.plot_utils.gaussian_distribution(f_df.loc[forecast,"prior_mean"],f_df.
    axt.fill_between(x,0,y,facecolor="0.5",alpha=0.25)
    x,y = pyemu.plot_utils.gaussian_distribution(f_df.loc[forecast,"post_mean"],f_df.
    axt.fill_between(x,0,y,facecolor="b",alpha=0.25)
```

```
axt.set_ylim(0,axt.get_ylim()[1])
axt.set_yticks([])
ax.set_title(forecast)
plt.show()
```

