

pestpp-glm

May 2, 2019

1 PESTPP-GLM

In this notebook, we will run PESTPP-GLM in standard parameter estimation mode and regularization mode. In both cases, we will use the baked-in bayes-linear posterior monte carlo analysis to get posterior forecast PDFs. We will use the prior monte carlo outputs as the prior forecast PDF.

```
In [1]: import os
import shutil
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import flopy
import pyemu
```

flopy is installed in /Users/jeremyw/Dev/gw1876/activities_2day_mfm/notebooks/flopy

```
In [2]: t_d = "template"
m_d = "master_glm"
```

```
In [3]: pst = pyemu.Pst(os.path.join(t_d, "freyberg.pst"))
pst.write_par_summary_table(filename="none")
```

```
Out[3]:
```

| | type | transform | count | initial value \ |
|-------------|-------------|-----------|-------|-----------------|
| gr_vka4 | gr_vka4 | log | 705 | 0 |
| cn_vka7 | cn_vka7 | log | 1 | 0 |
| welflux | welflux | log | 2 | 0 to 0.176091 |
| gr_rech2 | gr_rech2 | log | 705 | 0 |
| pp_vka0 | pp_vka0 | log | 32 | 0 |
| gr_sy5 | gr_sy5 | log | 705 | 0 |
| gr_ss3 | gr_ss3 | log | 705 | 0 |
| gr_vka5 | gr_vka5 | log | 705 | 0 |
| cn_vka8 | cn_vka8 | log | 1 | 0 |
| cn_hk6 | cn_hk6 | log | 1 | 0 |
| cn_rech4 | cn_rech4 | log | 1 | 0 |
| drncond_k00 | drncond_k00 | log | 10 | 0 |
| pp_hk1 | pp_hk1 | log | 32 | 0 |

| | | | | |
|-------------|-------------|-----|-----|----------|
| gr_sy4 | gr_sy4 | log | 705 | 0 |
| flow | flow | log | 1 | 0 |
| cn_ss6 | cn_ss6 | log | 1 | 0 |
| strk | strk | log | 40 | 0 |
| cn_strt6 | cn_strt6 | log | 1 | 0 |
| pp_prsity2 | pp_prsity2 | log | 32 | 0 |
| pp_sy1 | pp_sy1 | log | 32 | 0 |
| gr_strt4 | gr_strt4 | log | 705 | 0 |
| gr_hk5 | gr_hk5 | log | 705 | 0 |
| gr_hk4 | gr_hk4 | log | 705 | 0 |
| cn_sy7 | cn_sy7 | log | 1 | 0 |
| cn_ss7 | cn_ss7 | log | 1 | 0 |
| cn_ss8 | cn_ss8 | log | 1 | 0 |
| pp_strt0 | pp_strt0 | log | 32 | 0 |
| gr_prsity4 | gr_prsity4 | log | 705 | 0 |
| gr_rech3 | gr_rech3 | log | 705 | 0 |
| gr_hk3 | gr_hk3 | log | 705 | 0 |
| ... | ... | ... | ... | ... |
| cn_hk8 | cn_hk8 | log | 1 | 0 |
| cn_prsity6 | cn_prsity6 | log | 1 | 0 |
| pp_ss0 | pp_ss0 | log | 32 | 0 |
| cn_prsity7 | cn_prsity7 | log | 1 | 0 |
| cn_rech5 | cn_rech5 | log | 1 | -0.39794 |
| welflux_k02 | welflux_k02 | log | 6 | 0 |
| gr_vka3 | gr_vka3 | log | 705 | 0 |
| pp_prsity1 | pp_prsity1 | log | 32 | 0 |
| gr_strt3 | gr_strt3 | log | 705 | 0 |
| pp_hk0 | pp_hk0 | log | 32 | 0 |
| gr_ss5 | gr_ss5 | log | 705 | 0 |
| pp_ss2 | pp_ss2 | log | 32 | 0 |
| cn_sy6 | cn_sy6 | log | 1 | 0 |
| gr_prsity5 | gr_prsity5 | log | 705 | 0 |
| pp_vka2 | pp_vka2 | log | 32 | 0 |
| gr_sy3 | gr_sy3 | log | 705 | 0 |
| cn_hk7 | cn_hk7 | log | 1 | 0 |
| pp_rech0 | pp_rech0 | log | 32 | 0 |
| cn_strt8 | cn_strt8 | log | 1 | 0 |
| gr_ss4 | gr_ss4 | log | 705 | 0 |
| pp_ss1 | pp_ss1 | log | 32 | 0 |
| pp_sy2 | pp_sy2 | log | 32 | 0 |
| cn_sy8 | cn_sy8 | log | 1 | 0 |
| pp_hk2 | pp_hk2 | log | 32 | 0 |
| cn_prsity8 | cn_prsity8 | log | 1 | 0 |
| pp_rech1 | pp_rech1 | log | 32 | 0 |
| gr_strt5 | gr_strt5 | log | 705 | 0 |
| pp_prsity0 | pp_prsity0 | log | 32 | 0 |
| pp_strt2 | pp_strt2 | log | 32 | 0 |
| pp_sy0 | pp_sy0 | log | 32 | 0 |

| | upper bound | lower bound | standard deviation |
|-------------|---------------------|---------------|----------------------|
| gr_vka4 | 1 | -1 | 0.5 |
| cn_vka7 | 1 | -1 | 0.5 |
| welflux | 0.176091 to 0.30103 | -0.30103 to 0 | 0.0752575 to 0.11928 |
| gr_rech2 | 0.0413927 | -0.0457575 | 0.0217875 |
| pp_vka0 | 1 | -1 | 0.5 |
| gr_sy5 | 0.243038 | -0.60206 | 0.211275 |
| gr_ss3 | 1 | -1 | 0.5 |
| gr_vka5 | 1 | -1 | 0.5 |
| cn_vka8 | 1 | -1 | 0.5 |
| cn_hk6 | 1 | -1 | 0.5 |
| cn_rech4 | 0.0791812 | -0.09691 | 0.0440228 |
| drncond_k00 | 1 | -1 | 0.5 |
| pp_hk1 | 1 | -1 | 0.5 |
| gr_sy4 | 0.243038 | -0.60206 | 0.211275 |
| flow | 0.09691 | -0.124939 | 0.0554622 |
| cn_ss6 | 1 | -1 | 0.5 |
| strk | 2 | -2 | 1 |
| cn_strt6 | 0.0211893 | -0.0222764 | 0.0108664 |
| pp_prsity2 | 0 | -1 | 0.25 |
| pp_sy1 | 0.243038 | -0.60206 | 0.211275 |
| gr_strt4 | 0.0211893 | -0.0222764 | 0.0108664 |
| gr_hk5 | 1 | -1 | 0.5 |
| gr_hk4 | 1 | -1 | 0.5 |
| cn_sy7 | 0.243038 | -0.60206 | 0.211275 |
| cn_ss7 | 1 | -1 | 0.5 |
| cn_ss8 | 1 | -1 | 0.5 |
| pp_strt0 | 0.0211893 | -0.0222764 | 0.0108664 |
| gr_prsity4 | 0 | -1 | 0.25 |
| gr_rech3 | 0.0413927 | -0.0457575 | 0.0217875 |
| gr_hk3 | 1 | -1 | 0.5 |
| ... | ... | ... | ... |
| cn_hk8 | 1 | -1 | 0.5 |
| cn_prsity6 | 0 | -1 | 0.25 |
| pp_ss0 | 1 | -1 | 0.5 |
| cn_prsity7 | 0 | -1 | 0.25 |
| cn_rech5 | -0.09691 | -1 | 0.225772 |
| welflux_k02 | 1 | -1 | 0.5 |
| gr_vka3 | 1 | -1 | 0.5 |
| pp_prsity1 | 0 | -1 | 0.25 |
| gr_strt3 | 0.0211893 | -0.0222764 | 0.0108664 |
| pp_hk0 | 1 | -1 | 0.5 |
| gr_ss5 | 1 | -1 | 0.5 |
| pp_ss2 | 1 | -1 | 0.5 |
| cn_sy6 | 0.243038 | -0.60206 | 0.211275 |
| gr_prsity5 | 0 | -1 | 0.25 |
| pp_vka2 | 1 | -1 | 0.5 |

| | | | |
|------------|-----------|------------|-----------|
| gr_sy3 | 0.243038 | -0.60206 | 0.211275 |
| cn_hk7 | 1 | -1 | 0.5 |
| pp_rech0 | 0.0413927 | -0.0457575 | 0.0217875 |
| cn_strt8 | 0.0211893 | -0.0222764 | 0.0108664 |
| gr_ss4 | 1 | -1 | 0.5 |
| pp_ss1 | 1 | -1 | 0.5 |
| pp_sy2 | 0.243038 | -0.60206 | 0.211275 |
| cn_sy8 | 0.243038 | -0.60206 | 0.211275 |
| pp_hk2 | 1 | -1 | 0.5 |
| cn_prsity8 | 0 | -1 | 0.25 |
| pp_rech1 | 0.0413927 | -0.0457575 | 0.0217875 |
| gr_strt5 | 0.0211893 | -0.0222764 | 0.0108664 |
| pp_prsity0 | 0 | -1 | 0.25 |
| pp_strt2 | 0.0211893 | -0.0222764 | 0.0108664 |
| pp_sy0 | 0.243038 | -0.60206 | 0.211275 |

[65 rows x 7 columns]

1.0.1 reduce the number of adjustable parameters

This is the painful part: we cant use 10K+ pars because we cant wait around for that many runs and then the linear algebra of factoring a 10k+ by 10K+ matrix is also difficult. So that means we need to fix a lot a parameters #frownyface

```
In [4]: par = pst.parameter_data
```

```
In [5]: # grid-scale pars
gr_pars = par.loc[par.pargp.apply(lambda x: "gr" in x), "parnme"]
par.loc[gr_pars, "partrans"] = "fixed"
pst.npar_adj
```

```
Out [5]: 719
```

```
In [6]: # these are the sfr conductance parameters - Ive left all 40 adjustable
# but if you uncomment this, it will tie them into 1 parameter effectively
# strk_pars = par.loc[par.pargp=="strk", "parnme"]
# p1 = strk_pars.iloc[0]
# par.loc[strk_pars.iloc[1:], "partrans"] = "tied"
# par.loc[strk_pars.iloc[1:], "partied"] = p1
pst.npar_adj
```

```
Out [6]: 719
```

```
In [7]: par.loc[par.pargp.apply(lambda x: "pp" in x), "pargp"].unique()
```

```
Out [7]: array(['pp_hk0', 'pp_hk1', 'pp_hk2', 'pp_prsity0', 'pp_prsity1',
                'pp_prsity2', 'pp_rech0', 'pp_rech1', 'pp_ss0', 'pp_ss1', 'pp_ss2',
                'pp_strt0', 'pp_strt1', 'pp_strt2', 'pp_sy0', 'pp_sy1', 'pp_sy2',
                'pp_vka0', 'pp_vka1', 'pp_vka2'], dtype=object)
```

Fix the storage pilot points - we still have layer-scale storage pars adjustable

```
In [8]: #s_pars = par.loc[par.pargp.apply(lambda x: "pp" in x and ("ss" in x or "sy" in x)), "p"]  
        #par.loc[s_pars, "partrans"] = "fixed"  
        pst.npar_adj
```

```
Out[8]: 719
```

```
In [9]: adj_par = par.loc[par.partrans=="log", :]  
        adj_par.pargp.value_counts().sort_values()
```

```
Out[9]: cn_sy8          1  
        cn_rech4        1  
        cn_hk6          1  
        cn_ss6          1  
        cn_sy7          1  
        flow            1  
        cn_vka8         1  
        cn_vka7         1  
        cn_strt7        1  
        cn_strt6        1  
        cn_hk8          1  
        cn_prsity8      1  
        cn_rech5        1  
        cn_ss7          1  
        cn_strt8        1  
        cn_hk7          1  
        cn_vka6         1  
        cn_sy6          1  
        cn_prsity6      1  
        cn_ss8          1  
        cn_prsity7      1  
        welflux         2  
        welflux_k02     6  
        drncond_k00     10  
        pp_ss0          32  
        pp_sy0          32  
        pp_sy2          32  
        pp_vka0         32  
        pp_hk1          32  
        pp_prsity1      32  
        pp_prsity2      32  
        pp_rech1        32  
        pp_ss1          32  
        pp_sy1          32  
        pp_strt0        32  
        pp_vka1         32  
        pp_strt1        32  
        pp_ss2          32
```

```

pp_rech0      32
pp_prsity0    32
pp_hk0        32
pp_hk2        32
pp_vka2       32
pp_strt2      32
strk          40
Name: pargp, dtype: int64

```

fix the future recharge pilot points, vka in layers 1 and 3 and the initial condition pilot points (we still have layer-scale pars for each of these types)

```

In [10]: fi_grps = ["pp_rech1", "pp_vka0", "pp_vka2", "pp_strt0", "pp_strt1", "pp_strt2"]
          par.loc[par.pargp.apply(lambda x: x in fi_grps), "partrans"] = "fixed"
          pst.npar_adj

```

```

Out[10]: 527

```

Ok, thats better...so lets run PESTPP-GLM. We will use a single "base parameter" jacobian matrix as the basis for 6 super parameter iterations. Then we will draw 100 realizations from the FOSM posterior parameter covariance matrix and run those 100 realizations to get the psoterior forecast PDFs

```

In [11]: pst.control_data.noptmax = 3
          pst.pestpp_options["n_iter_base"] = -1
          pst.pestpp_options["n_iter_super"] = 3
          pst.pestpp_options["num_reals"] = 50 # this is how many ies uses
          pst.pestpp_options["parcov"] = "prior_cov.jcb"
          pst.write(os.path.join(t_d, "freyberg_pp.pst"))

```

```

In [12]: #pyemu.os_utils.start_slaves(t_d, "pestpp-glm", "freyberg_pp.pst", num_slaves=20, slave_r
          #                                     master_dir=m_d)

```

```

In [13]: df = df=pd.read_csv(os.path.join(m_d, "freyberg_pp.post.obsen.csv"), index_col=0)
          oe = pyemu.ObservationEnsemble.from_dataframe(pst=pst, df=df)

```

```

In [14]: ax = oe.phi_vector.hist()#bins=np.linspace(0, 100, 20))
          oe.phi_vector.sort_values().iloc[:20]

```

```


Out[14]: real_name
21      105.630787
33      250.056863
15      341.328184
45      341.959245
5       361.923751
23      361.985973
36      379.002255
19      433.680020
35      435.836633

```

```

8      566.688279
48     575.856958
37     591.763851
7      682.465841
17     734.368859
39     803.951847
11     913.109707
10     936.800544
34     957.916272
1      1119.577753
49     1168.722748
dtype: float64

```



Here we see the distribution of phi values across the 100 posterior realizations. Should we accept all of these??? The theoretical phi we should accept is number of nonzero obs (14).

To get a “posterior” ensemble, we need to throw out the realizations with large phi - lets just take the 20 best:

```
In [15]: oe_pt = oe.loc[oe.phi_vector.sort_values().index[:20],:] #just take the 20 lowest phi
```

We can also load and plot the FOSM forecast results along side of the ensemble results:

```
In [16]: f_df = pd.read_csv(os.path.join(m_d, "freyberg_pp.pred.usum.csv"), index_col=0)
         f_df.index = f_df.index.map(str.lower)
         f_df
```

```
Out[16]:
```



| | prior_mean | prior_stdev | prior_lower_bound | \ |
|------|------------|-------------|-------------------|---|
| name | | | | |



| | | | |
|--------------------|-----------|-----------|------------|
| fa_hw_19791230 | -977.2390 | 295.32800 | -1567.8900 |
| fa_hw_19801229 | -351.2160 | 409.77000 | -1170.7600 |
| fa_tw_19791230 | -453.0330 | 409.35100 | -1271.7400 |
| fa_tw_19801229 | 108.9600 | 506.73200 | -904.5040 |
| hds_00_013_002_000 | 39.6102 | 3.96314 | 31.6840 |
| hds_00_013_002_001 | 38.3838 | 4.05782 | 30.2681 |
| part_status | 2.0000 | 0.00000 | 2.0000 |
| part_time | 907.7020 | 704.75100 | -501.8010 |



| | prior_upper_bound | post_mean | post_stdev \ |
|--------------------|-------------------|-----------|--------------|
| name | | | |
| fa_hw_19791230 | -386.5840 | -729.6690 | 253.767000 |
| fa_hw_19801229 | 468.3240 | -160.1280 | 344.320000 |
| fa_tw_19791230 | 365.6690 | -149.8650 | 180.227000 |
| fa_tw_19801229 | 1122.4200 | 361.9180 | 272.345000 |
| hds_00_013_002_000 | 47.5365 | 36.9915 | 0.327097 |
| hds_00_013_002_001 | 46.4994 | 36.0961 | 0.710410 |
| part_status | 2.0000 | 1.0000 | 0.000000 |
| part_time | 2317.2000 | 4015.0000 | 604.279000 |



| | post_lower_bound | post_upper_bound |
|--------------------|------------------|------------------|
| name | | |
| fa_hw_19791230 | -1237.2000 | -222.1350 |
| fa_hw_19801229 | -848.7680 | 528.5130 |
| fa_tw_19791230 | -510.3190 | 210.5890 |
| fa_tw_19801229 | -182.7730 | 906.6090 |
| hds_00_013_002_000 | 36.3373 | 37.6457 |
| hds_00_013_002_001 | 34.6753 | 37.5170 |
| part_status | 1.0000 | 1.0000 |
| part_time | 2806.4400 | 5223.5600 |

```
In [17]: obs = pst.observation_data
fnames = pst.pestpp_options["forecasts"].split(",")
for forecast in fnames:
    ax = plt.subplot(111)
    oe_pt.loc[:,forecast].hist(ax=ax,color="b",alpha=0.5,normed=True)
    ax.plot([obs.loc[forecast,"obsval"],obs.loc[forecast,"obsval"]],ax.get_ylim(),"r")
    axt = plt.twinx()
    x,y = pyemu.plot_utils.gaussian_distribution(f_df.loc[forecast,"prior_mean"],f_df.loc[forecast,"prior_stdev"])
    axt.fill_between(x,0,y,facecolor="0.5",alpha=0.25)
    x,y = pyemu.plot_utils.gaussian_distribution(f_df.loc[forecast,"post_mean"],f_df.loc[forecast,"post_stdev"])
    axt.fill_between(x,0,y,facecolor="b",alpha=0.25)
    axt.set_ylim(0,axt.get_ylim()[1])
    axt.set_yticks([])
    ax.set_title(forecast)
plt.show()
```







1.0.2 Setup of Tikhonov regularization

Now lets setup and use some formal regularization to bring the final phi up to around 14. We will use first-order regularization based on the covariance matrix we build earlier:

```
In [18]: cov = pyemu.Cov.from_binary(os.path.join(t_d,"prior_cov.jcb"))
```

```
new binary format detected...
```

```
In [19]: pyemu.helpers.first_order_pearson_tikhonov(pst,cov)
```

```
getting CC matrix  
processing
```

```
In [20]: pst.prior_information.head()
```

```
Out[20]:
```

| | equation | obgnme | \ |
|-------|---|----------|---|
| pilbl | | | |
| pcc_1 | 1.0 * log(dc0000390005) - 1.0 * log(dc0000390006) = 0.0 | regul_cc | |
| pcc_2 | 1.0 * log(dc0000390005) - 1.0 * log(dc0000390007) = 0.0 | regul_cc | |
| pcc_3 | 1.0 * log(dc0000390005) - 1.0 * log(dc0000390008) = 0.0 | regul_cc | |
| pcc_4 | 1.0 * log(dc0000390005) - 1.0 * log(dc0000390009) = 0.0 | regul_cc | |
| pcc_5 | 1.0 * log(dc0000390005) - 1.0 * log(dc0000390010) = 0.0 | regul_cc | |

| | pilbl | weight |
|-------|-------|----------|
| pilbl | | |
| pcc_1 | pcc_1 | 0.904837 |
| pcc_2 | pcc_2 | 0.818731 |
| pcc_3 | pcc_3 | 0.740818 |
| pcc_4 | pcc_4 | 0.670320 |
| pcc_5 | pcc_5 | 0.606531 |

```
In [21]: shutil.copy2(os.path.join(m_d,"freyberg_pp.jcb"),os.path.join(t_d,"restart_pp.jcb"))
```

```
Out[21]: 'template/restart_pp.jcb'
```


```
In [22]: pst.pestpp_options["base_jacobian"] = "restart_pp.jcb"  
pst.reg_data.phimlim = pst.nnz_obs  
pst.reg_data.phimaccept = pst.reg_data.phimlim * 1.1  
pst.write(os.path.join(t_d,"freyberg_pp.pst"))
```

```
In [23]: pyemu.os_utils.start_slaves(t_d,"pestpp-glm","freyberg_pp.pst",num_slaves=20,slave_ro  
master_dir=m_d)
```

```
In [24]: df = df=pd.read_csv(os.path.join(m_d,"freyberg_pp.post.obsen.csv"),index_col=0)  
oe = pyemu.ObservationEnsemble.from_dataframe(pst=pst,df=df)
```

```
In [25]: ax = oe.phi_vector.hist()#bins=np.linspace(0,100,20))  
oe.phi_vector.sort_values().iloc[:20]
```

```
Out[25]: real_name
0      59.324204
38     82.194416
5      97.574472
17    117.532065
8     165.624313
23    168.838806
36    212.341022
49    218.655514
9     251.151530
37    264.760800
39    265.784626
44    274.808989
45    284.772193
48    294.292926
47    341.094426
7     352.179906
19    366.867432
26    399.599443
11    403.500953
35    429.953038
dtype: float64
```



Same as before, to get a “posterior” ensemble, we need to throw out the realizations with large ϕ - lets just take the 20 best:

```
In [26]: oe_pt = oe.loc[oe.phi_vector.sort_values().index[:20],:]
```

```
In [27]: f_df = pd.read_csv(os.path.join(m_d,"freyberg_pp.pred.usum.csv"),index_col=0)
f_df.index = f_df.index.map(str.lower)
f_df
```

```
Out [27]:
```

| | prior_mean | prior_stdev | prior_lower_bound \ |
|--------------------|------------|-------------|---------------------|
| name | | | |
| fa_hw_19791230 | -977.2390 | 295.32800 | -1567.8900 |
| fa_hw_19801229 | -351.2160 | 409.77000 | -1170.7600 |
| fa_tw_19791230 | -453.0330 | 409.35100 | -1271.7400 |
| fa_tw_19801229 | 108.9600 | 506.73200 | -904.5040 |
| hds_00_013_002_000 | 39.6102 | 3.96314 | 31.6840 |
| hds_00_013_002_001 | 38.3838 | 4.05782 | 30.2681 |
| part_status | 2.0000 | 0.00000 | 2.0000 |
| part_time | 907.7020 | 704.75100 | -501.8010 |

| | prior_upper_bound | post_mean | post_stdev \ |
|--------------------|-------------------|-----------|--------------|
| name | | | |
| fa_hw_19791230 | -386.5840 | -682.3380 | 252.978000 |
| fa_hw_19801229 | 468.3240 | -227.9100 | 342.479000 |
| fa_tw_19791230 | 365.6690 | -138.5270 | 179.850000 |
| fa_tw_19801229 | 1122.4200 | 264.8790 | 272.352000 |
| hds_00_013_002_000 | 47.5365 | 38.8495 | 0.409407 |
| hds_00_013_002_001 | 46.4994 | 37.7571 | 0.749881 |
| part_status | 2.0000 | 2.0000 | 0.000000 |
| part_time | 2317.2000 | 1059.5000 | 604.211000 |

| | post_lower_bound | post_upper_bound |
|--------------------|------------------|------------------|
| name | | |
| fa_hw_19791230 | -1188.2900 | -176.3830 |
| fa_hw_19801229 | -912.8690 | 457.0490 |
| fa_tw_19791230 | -498.2270 | 221.1720 |
| fa_tw_19801229 | -279.8240 | 809.5820 |
| hds_00_013_002_000 | 38.0307 | 39.6683 |
| hds_00_013_002_001 | 36.2573 | 39.2568 |
| part_status | 2.0000 | 2.0000 |
| part_time | -148.9210 | 2267.9200 |

```
In [28]: obs = pst.observation_data
fnames = pst.pestpp_options["forecasts"].split(",")
for forecast in fnames:
    ax = plt.subplot(111)
    oe_pt.loc[:,forecast].hist(ax=ax,color="b",alpha=0.5,normed=True)
    ax.plot([obs.loc[forecast,"obsval"],obs.loc[forecast,"obsval"]],ax.get_ylim(),"r")
    axt = plt.twinx()
    x,y = pyemu.plot_utils.gaussian_distribution(f_df.loc[forecast,"prior_mean"],f_df.
    axt.fill_between(x,0,y,facecolor="0.5",alpha=0.25)
    x,y = pyemu.plot_utils.gaussian_distribution(f_df.loc[forecast,"post_mean"],f_df.
    axt.fill_between(x,0,y,facecolor="b",alpha=0.25)
```

```

axt.set_ylim(0,axt.get_ylim()[1])
axt.set_yticks([])
ax.set_title(forecast)
plt.show()

```

