

pestpp-glm

June 5, 2019

1 PESTPP-GLM

In this notebook, we will run PESTPP-GLM in standard parameter estimation mode and regularization mode. In both cases, we will use the baked-in bayes-linear posterior monte carlo analysis to get posterior forecast PDFs. We will use the prior monte carlo outputs as the prior forecast PDF.

```
In [1]: %matplotlib inline
import os
import shutil
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib as mpl
plt.rcParams['font.size']=12
import flopy
import pyemu
```

flopy is installed in /Users/jeremyw/Dev/gw1876/activities_2day_mfm/notebooks/flopy

1.1 SUPER IMPORTANT: SET HOW MANY PARALLEL WORKERS TO USE

```
In [2]: num_workers = 20
```

```
In [3]: t_d = "template"
m_d = "master_glm"
```

```
In [4]: pst = pyemu.Pst(os.path.join(t_d,"freyberg.pst"))
pst.write_par_summary_table(filename="none")
```

```
Out[4]:
```

	type	transform	count	initial value	\
cn_strt8	cn_strt8	log	1	0	
pp_prsity0	pp_prsity0	log	32	0	
cn_prsity7	cn_prsity7	log	1	0	
pp_strt2	pp_strt2	log	32	0	
pp_strt0	pp_strt0	log	32	0	
cn_hk7	cn_hk7	log	1	0	
pp_vka1	pp_vka1	log	32	0	

gr_vka5	gr_vka5	log	705	0
pp_hk2	pp_hk2	log	32	0
cn_prsity6	cn_prsity6	log	1	0
pp_rech1	pp_rech1	log	32	0
pp_prsity2	pp_prsity2	log	32	0
pp_sy2	pp_sy2	log	32	0
gr_rech2	gr_rech2	log	705	0
gr_rech3	gr_rech3	log	705	0
pp_sy1	pp_sy1	log	32	0
cn_rech5	cn_rech5	log	1	-0.39794
cn_strt7	cn_strt7	log	1	0
gr_strt4	gr_strt4	log	705	0
welflux_k02	welflux_k02	log	6	0
pp_sy0	pp_sy0	log	32	0
gr_strt3	gr_strt3	log	705	0
welflux	welflux	log	2	0 to 0.176091
cn_vka6	cn_vka6	log	1	0
cn_strt6	cn_strt6	log	1	0
gr_ss5	gr_ss5	log	705	0
cn_ss7	cn_ss7	log	1	0
gr_hk5	gr_hk5	log	705	0
cn_ss8	cn_ss8	log	1	0
gr_vka3	gr_vka3	log	705	0
...
pp_hk1	pp_hk1	log	32	0
pp_hk0	pp_hk0	log	32	0
gr_prsity5	gr_prsity5	log	705	0
pp_ss1	pp_ss1	log	32	0
cn_sy8	cn_sy8	log	1	0
gr_ss4	gr_ss4	log	705	0
gr_vka4	gr_vka4	log	705	0
pp_strt1	pp_strt1	log	32	0
cn_vka7	cn_vka7	log	1	0
pp_ss2	pp_ss2	log	32	0
drncond_k00	drncond_k00	log	10	0
gr_sy4	gr_sy4	log	705	0
cn_vka8	cn_vka8	log	1	0
gr_prsity4	gr_prsity4	log	705	0
cn_hk6	cn_hk6	log	1	0
flow	flow	log	1	0
cn_rech4	cn_rech4	log	1	0
pp_rech0	pp_rech0	log	32	0
gr_sy5	gr_sy5	log	705	0
cn_prsity8	cn_prsity8	log	1	0
pp_vka0	pp_vka0	log	32	0
pp_vka2	pp_vka2	log	32	0
gr_hk3	gr_hk3	log	705	0
gr_strt5	gr_strt5	log	705	0

gr_hk4	gr_hk4	log	705	0
cn_sy6	cn_sy6	log	1	0
pp_ss0	pp_ss0	log	32	0
cn_sy7	cn_sy7	log	1	0
gr_prsity3	gr_prsity3	log	705	0
cn_hk8	cn_hk8	log	1	0

	upper bound	lower bound	standard deviation
cn_strt8	0.0211893	-0.0222764	0.0108664
pp_prsity0	0.176091	-0.30103	0.11928
cn_prsity7	0.176091	-0.30103	0.11928
pp_strt2	0.0211893	-0.0222764	0.0108664
pp_strt0	0.0211893	-0.0222764	0.0108664
cn_hk7	1	-1	0.5
pp_vka1	1	-1	0.5
gr_vka5	1	-1	0.5
pp_hk2	1	-1	0.5
cn_prsity6	0.176091	-0.30103	0.11928
pp_rech1	0.0413927	-0.0457575	0.0217875
pp_prsity2	0.176091	-0.30103	0.11928
pp_sy2	0.243038	-0.60206	0.211275
gr_rech2	0.0413927	-0.0457575	0.0217875
gr_rech3	0.0413927	-0.0457575	0.0217875
pp_sy1	0.243038	-0.60206	0.211275
cn_rech5	-0.09691	-1	0.225772
cn_strt7	0.0211893	-0.0222764	0.0108664
gr_strt4	0.0211893	-0.0222764	0.0108664
welflux_k02	1	-1	0.5
pp_sy0	0.243038	-0.60206	0.211275
gr_strt3	0.0211893	-0.0222764	0.0108664
welflux	0.176091 to 0.30103	-0.30103 to 0	0.0752575 to 0.11928
cn_vka6	1	-1	0.5
cn_strt6	0.0211893	-0.0222764	0.0108664
gr_ss5	1	-1	0.5
cn_ss7	1	-1	0.5
gr_hk5	1	-1	0.5
cn_ss8	1	-1	0.5
gr_vka3	1	-1	0.5
...
pp_hk1	1	-1	0.5
pp_hk0	1	-1	0.5
gr_prsity5	0.176091	-0.30103	0.11928
pp_ss1	1	-1	0.5
cn_sy8	0.243038	-0.60206	0.211275
gr_ss4	1	-1	0.5
gr_vka4	1	-1	0.5
pp_strt1	0.0211893	-0.0222764	0.0108664
cn_vka7	1	-1	0.5

pp_ss2	1	-1	0.5
drncond_k00	1	-1	0.5
gr_sy4	0.243038	-0.60206	0.211275
cn_vka8	1	-1	0.5
gr_prsity4	0.176091	-0.30103	0.11928
cn_hk6	1	-1	0.5
flow	0.09691	-0.124939	0.0554622
cn_rech4	0.0791812	-0.09691	0.0440228
pp_rech0	0.0413927	-0.0457575	0.0217875
gr_sy5	0.243038	-0.60206	0.211275
cn_prsity8	0.176091	-0.30103	0.11928
pp_vka0	1	-1	0.5
pp_vka2	1	-1	0.5
gr_hk3	1	-1	0.5
gr_strt5	0.0211893	-0.0222764	0.0108664
gr_hk4	1	-1	0.5
cn_sy6	0.243038	-0.60206	0.211275
pp_ss0	1	-1	0.5
cn_sy7	0.243038	-0.60206	0.211275
gr_prsity3	0.176091	-0.30103	0.11928
cn_hk8	1	-1	0.5

[65 rows x 7 columns]

1.1.1 reduce the number of adjustable parameters

This is the painful part: we cant use 10K+ pars because we cant wait around for that many runs and then the linear algebra of factoring a 10k+ by 10K+ matrix is also difficult. So that means we need to fix a lot a parameters #frownyface

```
In [5]: par = pst.parameter_data
```

```
In [6]: # grid-scale pars
```

```
gr_pars = par.loc[par.pargp.apply(lambda x: "gr" in x), "parnme"]
par.loc[gr_pars, "partrans"] = "fixed"
pst.npar_adj
```

```
Out[6]: 719
```

```
In [7]: # these are the sfr conductance parameters - Ive left all 40 adjustable
# but if you uncomment this, it will tie them into 1 parameter effectively
# strk_pars = par.loc[par.pargp=="strk", "parnme"]
# p1 = strk_pars.iloc[0]
# par.loc[strk_pars.iloc[1:], "partrans"] = "tied"
# par.loc[strk_pars.iloc[1:], "partied"] = p1
pst.npar_adj
```

```
Out[7]: 719
```

```
In [8]: par.loc[par.pargp.apply(lambda x: "pp" in x), "pargp"].unique()
```

```
Out[8]: array(['pp_hk0', 'pp_hk1', 'pp_hk2', 'pp_prsity0', 'pp_prsity1',
               'pp_prsity2', 'pp_rech0', 'pp_rech1', 'pp_ss0', 'pp_ss1', 'pp_ss2',
               'pp_strt0', 'pp_strt1', 'pp_strt2', 'pp_sy0', 'pp_sy1', 'pp_sy2',
               'pp_vka0', 'pp_vka1', 'pp_vka2'], dtype=object)
```

Fix the storage pilot points - we still have layer-scale storage pars adjustable

```
In [9]: #s_pars = par.loc[par.pargp.apply(lambda x: "pp" in x and ("ss" in x or "sy" in x)), "p
        #par.loc[s_pars, "partrans"] = "fixed"
        pst.npar_adj
```

```
Out[9]: 719
```

```
In [10]: adj_par = par.loc[par.partrans=="log", :]
         adj_par.pargp.value_counts().sort_values()
```

```
Out[10]: cn_hk8          1
         cn_hk6          1
         cn_prsity7      1
         cn_prsity8      1
         cn_hk7          1
         cn_rech4        1
         flow            1
         cn_prsity6      1
         cn_vka8          1
         cn_rech5        1
         cn_vka7          1
         cn_ss8          1
         cn_sy6          1
         cn_sy7          1
         cn_vka6          1
         cn_strt6        1
         cn_ss7          1
         cn_strt7        1
         cn_strt8        1
         cn_ss6          1
         cn_sy8          1
         welflux         2
         welflux_k02     6
         drncond_k00     10
         pp_ss2          32
         pp_ss0          32
         pp_hk0          32
         pp_rech0        32
         pp_strt2        32
         pp_strt0        32
         pp_vka1         32
```

```

pp_hk2          32
pp_rech1        32
pp_prsity2      32
pp_sy2          32
pp_sy0          32
pp_prsity1      32
pp_hk1          32
pp_prsity0      32
pp_ss1          32
pp_vka2         32
pp_strt1        32
pp_vka0         32
pp_sy1          32
strk            40
Name: pargp, dtype: int64

```

fix the future recharge pilot points, vka in layers 1 and 3 and the initial condition pilot points (we still have layer-scale pars for each of these types)

```

In [11]: fi_grps = ["pp_rech1", "pp_vka0", "pp_vka2", "pp_strt0", "pp_strt1", "pp_strt2"]
          par.loc[par.pargp.apply(lambda x: x in fi_grps), "partrans"] = "fixed"
          pst.npar_adj

```

```

Out[11]: 527

```

Ok, thats better...so lets run PESTPP-GLM. We will use a single "base parameter" jacobian matrix as the basis for 6 super parameter iterations. Then we will draw 100 realizations from the FOSM posterior parameter covariance matrix and run those 100 realizations to get the psoterior forecast PDFs

```

In [12]: pst.control_data.noptmax = 3
          pst.pestpp_options["n_iter_base"] = -1
          pst.pestpp_options["n_iter_super"] = 3
          pst.pestpp_options["num_reals"] = 50 # this is how many ies uses
          pst.pestpp_options["parcov"] = "prior_cov.jcb"
          pst.write(os.path.join(t_d, "freyberg_pp.pst"))

```

```

noptmax:3, npar_adj:527, nnz_obs:14

```

```

In [13]: pyemu.os_utils.start_slaves(t_d, "pestpp-glm", "freyberg_pp.pst", num_slaves=num_workers
          master_dir=m_d)

```

```

In [14]: df = df=pd.read_csv(os.path.join(m_d, "freyberg_pp.post.obsen.csv"), index_col=0)
          oe = pyemu.ObservationEnsemble.from_dataframe(pst=pst, df=df)

```

```


In [15]: ax = oe.phi_vector.hist()#bins=np.linspace(0,100,20))
          oe.phi_vector.sort_values().iloc[:20]

```

```

Out[15]: real_name
16      93.563008
25     170.672590
11     215.054772
33     229.920615
6      246.861326
24     251.063969
45     272.209134
23     279.243029
37     322.455007
26     332.400050
21     343.341730
48     354.597653
34     401.095146
46     408.511642
7      420.700508
39     441.698089
42     442.051026
12     489.628147
3      489.754867
10     545.568931
dtype: float64

```



Here we see the distribution of phi values across the 100 posterior realizations. Should we accept all of these??? The theoretical phi we should accept is number of nonzero obs (14).

To get a “posterior” ensemble, we need to throw out the realizations with large phi - lets just take the 20 best:

```
In [16]: oe_pt = oe.loc[oe.phi_vector.sort_values().index[:20],:] #just take the 20 lowest phi
```

We can also load and plot the FOSM forecast results along side of the ensemble results:

```
In [17]: f_df = pd.read_csv(os.path.join(m_d,"freyberg_pp.pred.usum.csv"),index_col=0)
f_df.index = f_df.index.map(str.lower)
f_df
```

```
Out [17]:
```

	prior_mean	prior_stdev	prior_lower_bound \
name			
fa_hw_19791230	-977.2390	295.32800	-1567.8900
fa_hw_19801229	-351.2160	409.77000	-1170.7600
fa_tw_19791230	-453.0330	409.35100	-1271.7400
fa_tw_19801229	108.9600	506.73200	-904.5040
hds_00_013_002_000	39.6102	3.96314	31.6840
hds_00_013_002_001	38.3838	4.05782	30.2681
part_status	2.0000	0.00000	2.0000
part_time	907.7020	570.98600	-234.2690

	prior_upper_bound	post_mean	post_stdev \
name			
fa_hw_19791230	-386.5840	-675.0430	247.837000
fa_hw_19801229	468.3240	37.6125	335.549000
fa_tw_19791230	365.6690	-299.4690	170.577000
fa_tw_19801229	1122.4200	333.7460	263.672000
hds_00_013_002_000	47.5365	37.6359	0.273990
hds_00_013_002_001	46.4994	36.4927	0.687833
part_status	2.0000	2.0000	0.000000
part_time	2049.6700	703.9470	436.945000

	post_lower_bound	post_upper_bound
name		
fa_hw_19791230	-1170.7200	-179.3700
fa_hw_19801229	-633.4850	708.7100
fa_tw_19791230	-640.6230	41.6858
fa_tw_19801229	-193.5970	861.0890
hds_00_013_002_000	37.0879	38.1838
hds_00_013_002_001	35.1171	37.8684
part_status	2.0000	2.0000
part_time	-169.9420	1577.8400


```
In [18]: obs = pst.observation_data
fnames = pst.pestpp_options["forecasts"].split(",")
for forecast in fnames:
    ax = plt.subplot(111)
    oe_pt.loc[:,forecast].hist(ax=ax,color="b",alpha=0.5,normed=True)
    ax.plot([obs.loc[forecast,"obsval"],obs.loc[forecast,"obsval"]],ax.get_ylim(),"r")
    axt = plt.twinx()
    x,y = pyemu.plot_utils.gaussian_distribution(f_df.loc[forecast,"prior_mean"],f_df
```







```


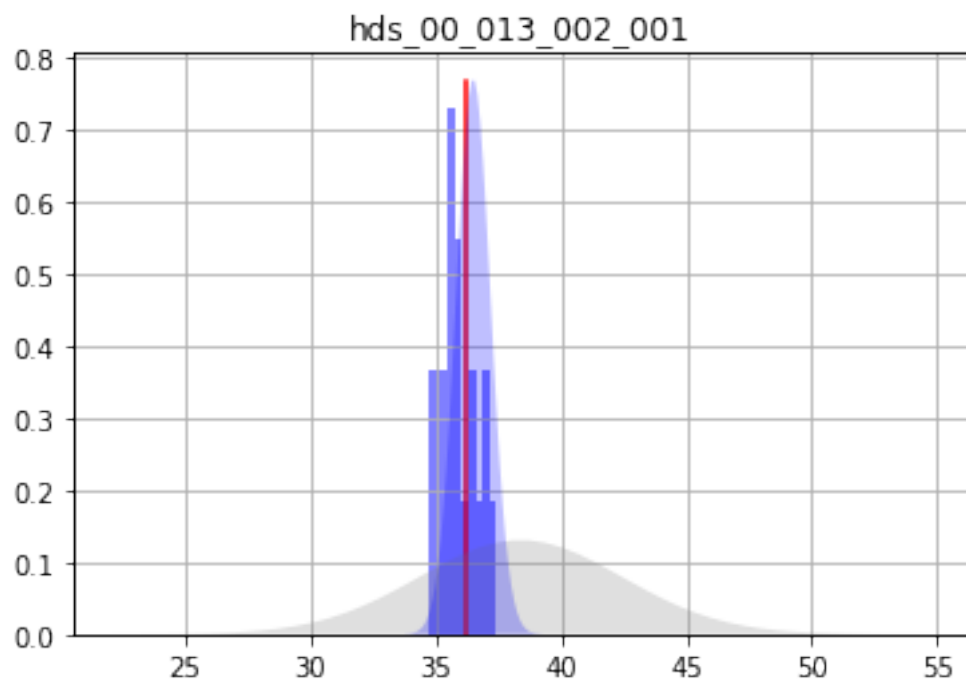
axt.fill_between(x,0,y,facecolor="0.5",alpha=0.25)
x,y = pyemu.plot_utils.gaussian_distribution(f_df.loc[forecast,"post_mean"],f_df.
axt.fill_between(x,0,y,facecolor="b",alpha=0.25)
axt.set_ylim(0,axt.get_ylim()[1])
axt.set_yticks([])
ax.set_title(forecast)
plt.show()

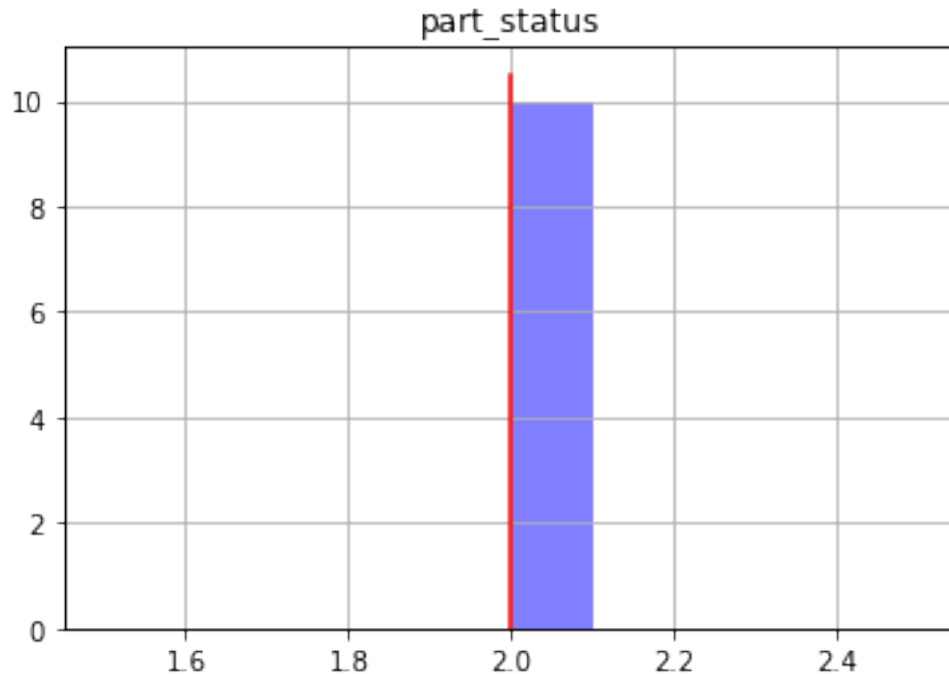
```











1.1.2 Setup of Tikhonov regularization

Now let's setup and use some formal regularization to bring the final phi up to around 14. We will use first-order regularization based on the covariance matrix we build earlier:

```
In [19]: cov = pyemu.Cov.from_binary(os.path.join(t_d,"prior_cov.jcb"))
```

```
new binary format detected...
```

```
In [20]: pyemu.helpers.first_order_pearson_tikhonov(pst,cov)
```

```
getting CC matrix
processing
```

```
In [21]: pst.prior_information.head()
```

```
Out[21]:
```

	equation	obgnme	\
pilbl			
pcc_1	1.0 * log(dc0000390005) - 1.0 * log(dc0000390006) = 0.0	regul_cc	
pcc_2	1.0 * log(dc0000390005) - 1.0 * log(dc0000390007) = 0.0	regul_cc	
pcc_3	1.0 * log(dc0000390005) - 1.0 * log(dc0000390008) = 0.0	regul_cc	
pcc_4	1.0 * log(dc0000390005) - 1.0 * log(dc0000390009) = 0.0	regul_cc	
pcc_5	1.0 * log(dc0000390005) - 1.0 * log(dc0000390010) = 0.0	regul_cc	

	pilbl	weight
pilbl		
pcc_1	pcc_1	0.904837
pcc_2	pcc_2	0.818731
pcc_3	pcc_3	0.740818
pcc_4	pcc_4	0.670320
pcc_5	pcc_5	0.606531

```
In [22]: shutil.copy2(os.path.join(m_d, "freyberg_pp.jcb"), os.path.join(t_d, "restart_pp.jcb"))
```

```
Out[22]: 'template/restart_pp.jcb'
```

```
In [23]: pst.pestpp_options["base_jacobian"] = "restart_pp.jcb"
pst.reg_data.phimlim = pst.nnz_obs
pst.reg_data.phimaccept = pst.reg_data.phimlim * 1.1
pst.write(os.path.join(t_d, "freyberg_pp.pst"))
```

```
noptmax:3, npar_adj:527, nnz_obs:14
```

```
In [24]: pyemu.os_utils.start_slaves(t_d, "pestpp-glm", "freyberg_pp.pst", num_slaves=num_workers,
                                     master_dir=m_d)
```

```
In [25]: df = df=pd.read_csv(os.path.join(m_d, "freyberg_pp.post.obsen.csv"), index_col=0)
oe = pyemu.ObservationEnsemble.from_dataframe(pst=pst, df=df)
```


```
In [26]: ax = oe.phi_vector.hist()#bins=np.linspace(0,100,20))
oe.phi_vector.sort_values().iloc[:20]
```

```
Out[26]: real_name
46      41.263489
16      63.951582
35     121.513255
7      137.446945
11     208.807130
45     210.737423
12     261.316576
9      281.400104
47     321.419440
19     330.908864
24     331.863379
25     333.766324
39     352.884354
30     361.820742
26     383.330645
21     390.853366
15     408.456242
44     413.387639
```

```

1      498.360492
33     514.880477
dtype: float64

```



Same as before, to get a “posterior” ensemble, we need to throw out the realizations with large phi - lets just take the 20 best:

```
In [27]: oe_pt = oe.loc[oe.phi_vector.sort_values().index[:20],:]
```

```
In [28]: f_df = pd.read_csv(os.path.join(m_d,"freyberg_pp.pred.usum.csv"),index_col=0)
f_df.index = f_df.index.map(str.lower)
f_df
```

```
Out [28]:
```

	prior_mean	prior_stdev	prior_lower_bound \
name			
fa_hw_19791230	-977.2390	295.32800	-1567.8900
fa_hw_19801229	-351.2160	409.77000	-1170.7600
fa_tw_19791230	-453.0330	409.35100	-1271.7400
fa_tw_19801229	108.9600	506.73200	-904.5040
hds_00_013_002_000	39.6102	3.96314	31.6840
hds_00_013_002_001	38.3838	4.05782	30.2681
part_status	2.0000	0.00000	2.0000
part_time	907.7020	570.98600	-234.2690

	prior_upper_bound	post_mean	post_stdev \
name			
fa_hw_19791230	-386.5840	-640.3370	250.802000

fa_hw_19801229	468.3240	-94.8227	339.971000
fa_tw_19791230	365.6690	-290.4730	170.616000
fa_tw_19801229	1122.4200	184.3570	263.999000
hds_00_013_002_000	47.5365	37.7893	0.274188
hds_00_013_002_001	46.4994	36.9350	0.687976
part_status	2.0000	2.0000	0.000000
part_time	2049.6700	966.2250	437.215000

	post_lower_bound	post_upper_bound
name		
fa_hw_19791230	-1141.9400	-138.7340
fa_hw_19801229	-774.7640	585.1190
fa_tw_19791230	-631.7040	50.7581
fa_tw_19801229	-343.6410	712.3540
hds_00_013_002_000	37.2409	38.3377
hds_00_013_002_001	35.5591	38.3110
part_status	2.0000	2.0000
part_time	91.7946	1840.6600

```
In [29]: obs = pst.observation_data
fnames = pst.pestpp_options["forecasts"].split(",")
for forecast in fnames:
    ax = plt.subplot(111)
    oe_pt.loc[:,forecast].hist(ax=ax,color="b",alpha=0.5,normed=True)
    ax.plot([obs.loc[forecast,"obsval"],obs.loc[forecast,"obsval"]],ax.get_ylim(),"r")
    axt = plt.twinx()
    x,y = pyemu.plot_utils.gaussian_distribution(f_df.loc[forecast,"prior_mean"],f_df.loc[forecast,"prior_std"])
    axt.fill_between(x,0,y,facecolor="0.5",alpha=0.25)
    x,y = pyemu.plot_utils.gaussian_distribution(f_df.loc[forecast,"post_mean"],f_df.loc[forecast,"post_std"])
    axt.fill_between(x,0,y,facecolor="b",alpha=0.25)
    axt.set_ylim(0,axt.get_ylim()[1])
    axt.set_yticks([])
    ax.set_title(forecast)
plt.show()
```