

# pestpp-glm

May 2, 2019

## 1 PESTPP-GLM

In this notebook, we will run PESTPP-GLM in standard parameter estimation mode and regularization mode. In both cases, we will use the baked-in bayes-linear posterior monte carlo analysis to get posterior forecast PDFs. We will use the prior monte carlo outputs as the prior forecast PDF.

```
In [1]: import os
import shutil
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import flopy
import pyemu
```

flopy is installed in /Users/jeremyw/Dev/gw1876/activities\_2day\_mfm/notebooks/flopy

```
In [2]: t_d = "template"
m_d = "master_glm"
```

```
In [3]: pst = pyemu.Pst(os.path.join(t_d, "freyberg.pst"))
pst.write_par_summary_table(filename="none")
```

```
Out[3]:
```

	type	transform	count	initial value \
pp_strt2	pp_strt2	log	32	0
gr_rech2	gr_rech2	log	705	0
cn_strt8	cn_strt8	log	1	0
gr_vka4	gr_vka4	log	705	0
gr_ss3	gr_ss3	log	705	0
pp_hk0	pp_hk0	log	32	0
cn_hk7	cn_hk7	log	1	0
pp_rech0	pp_rech0	log	32	0
cn_hk6	cn_hk6	log	1	0
welflux	welflux	log	2	0 to 0.176091
gr_hk5	gr_hk5	log	705	0
cn_sy8	cn_sy8	log	1	0
gr_strt3	gr_strt3	log	705	0

pp_ss0	pp_ss0	log	32	0
gr_ss5	gr_ss5	log	705	0
gr_vka5	gr_vka5	log	705	0
pp_rech1	pp_rech1	log	32	0
gr_sy4	gr_sy4	log	705	0
pp_strt1	pp_strt1	log	32	0
pp_vka1	pp_vka1	log	32	0
pp_prsity1	pp_prsity1	log	32	0
gr_sy3	gr_sy3	log	705	0
cn_prsity8	cn_prsity8	log	1	0
pp_hk1	pp_hk1	log	32	0
cn_hk8	cn_hk8	log	1	0
gr_strt5	gr_strt5	log	705	0
pp_sy0	pp_sy0	log	32	0
cn_strt7	cn_strt7	log	1	0
gr_ss4	gr_ss4	log	705	0
cn_ss7	cn_ss7	log	1	0
...	...	...	...	...
pp_sy1	pp_sy1	log	32	0
cn_vka6	cn_vka6	log	1	0
gr_hk4	gr_hk4	log	705	0
cn_prsity7	cn_prsity7	log	1	0
flow	flow	log	1	0
gr_prsity4	gr_prsity4	log	705	0
welflux_k02	welflux_k02	log	6	0
cn_vka8	cn_vka8	log	1	0
cn_ss6	cn_ss6	log	1	0
cn_rech4	cn_rech4	log	1	0
pp_ss2	pp_ss2	log	32	0
gr_rech3	gr_rech3	log	705	0
pp_strt0	pp_strt0	log	32	0
pp_prsity0	pp_prsity0	log	32	0
cn_prsity6	cn_prsity6	log	1	0
gr_strt4	gr_strt4	log	705	0
cn_rech5	cn_rech5	log	1	-0.39794
cn_strt6	cn_strt6	log	1	0
cn_ss8	cn_ss8	log	1	0
gr_hk3	gr_hk3	log	705	0
cn_sy6	cn_sy6	log	1	0
gr_prsity3	gr_prsity3	log	705	0
pp_vka0	pp_vka0	log	32	0
gr_sy5	gr_sy5	log	705	0
pp_hk2	pp_hk2	log	32	0
pp_vka2	pp_vka2	log	32	0
pp_sy2	pp_sy2	log	32	0
pp_prsity2	pp_prsity2	log	32	0
strk	strk	log	40	0
gr_vka3	gr_vka3	log	705	0

	upper bound	lower bound	standard deviation
pp_strt2	0.0211893	-0.0222764	0.0108664
gr_rech2	0.0413927	-0.0457575	0.0217875
cn_strt8	0.0211893	-0.0222764	0.0108664
gr_vka4	1	-1	0.5
gr_ss3	1	-1	0.5
pp_hk0	1	-1	0.5
cn_hk7	1	-1	0.5
pp_rech0	0.0413927	-0.0457575	0.0217875
cn_hk6	1	-1	0.5
welflux	0.176091 to 0.30103	-0.30103 to 0	0.0752575 to 0.11928
gr_hk5	1	-1	0.5
cn_sy8	0.243038	-0.60206	0.211275
gr_strt3	0.0211893	-0.0222764	0.0108664
pp_ss0	1	-1	0.5
gr_ss5	1	-1	0.5
gr_vka5	1	-1	0.5
pp_rech1	0.0413927	-0.0457575	0.0217875
gr_sy4	0.243038	-0.60206	0.211275
pp_strt1	0.0211893	-0.0222764	0.0108664
pp_vka1	1	-1	0.5
pp_prsity1	0	-1	0.25
gr_sy3	0.243038	-0.60206	0.211275
cn_prsity8	0	-1	0.25
pp_hk1	1	-1	0.5
cn_hk8	1	-1	0.5
gr_strt5	0.0211893	-0.0222764	0.0108664
pp_sy0	0.243038	-0.60206	0.211275
cn_strt7	0.0211893	-0.0222764	0.0108664
gr_ss4	1	-1	0.5
cn_ss7	1	-1	0.5
...	...	...	...
pp_sy1	0.243038	-0.60206	0.211275
cn_vka6	1	-1	0.5
gr_hk4	1	-1	0.5
cn_prsity7	0	-1	0.25
flow	0.09691	-0.124939	0.0554622
gr_prsity4	0	-1	0.25
welflux_k02	1	-1	0.5
cn_vka8	1	-1	0.5
cn_ss6	1	-1	0.5
cn_rech4	0.0791812	-0.09691	0.0440228
pp_ss2	1	-1	0.5
gr_rech3	0.0413927	-0.0457575	0.0217875
pp_strt0	0.0211893	-0.0222764	0.0108664
pp_prsity0	0	-1	0.25
cn_prsity6	0	-1	0.25

gr_strt4	0.0211893	-0.0222764	0.0108664
cn_rech5	-0.09691	-1	0.225772
cn_strt6	0.0211893	-0.0222764	0.0108664
cn_ss8	1	-1	0.5
gr_hk3	1	-1	0.5
cn_sy6	0.243038	-0.60206	0.211275
gr_prsity3	0	-1	0.25
pp_vka0	1	-1	0.5
gr_sy5	0.243038	-0.60206	0.211275
pp_hk2	1	-1	0.5
pp_vka2	1	-1	0.5
pp_sy2	0.243038	-0.60206	0.211275
pp_prsity2	0	-1	0.25
strk	2	-2	1
gr_vka3	1	-1	0.5

[65 rows x 7 columns]

### 1.0.1 reduce the number of adjustable parameters

This is the painful part: we cant use 10K+ pars because we cant wait around for that many runs and then the linear algebra of factoring a 10k+ by 10K+ matrix is also difficult. So that means we need to fix a lot a parameters #frownyface

```
In [4]: par = pst.parameter_data
```

```
In [5]: # grid-scale pars
gr_pars = par.loc[par.pargp.apply(lambda x: "gr" in x), "parnme"]
par.loc[gr_pars, "partrans"] = "fixed"
pst.npar_adj
```

```
Out[5]: 719
```

```
In [6]: # these are the sfr conductance parameters - Ive left all 40 adjustable
# but if you uncomment this, it will tie them into 1 parameter effectively
# strk_pars = par.loc[par.pargp=="strk", "parnme"]
# p1 = strk_pars.iloc[0]
# par.loc[strk_pars.iloc[1:], "partrans"] = "tied"
# par.loc[strk_pars.iloc[1:], "partied"] = p1
pst.npar_adj
```

```
Out[6]: 719
```

```
In [7]: par.loc[par.pargp.apply(lambda x: "pp" in x), "pargp"].unique()
```

```
Out[7]: array(['pp_hk0', 'pp_hk1', 'pp_hk2', 'pp_prsity0', 'pp_prsity1',
               'pp_prsity2', 'pp_rech0', 'pp_rech1', 'pp_ss0', 'pp_ss1', 'pp_ss2',
               'pp_strt0', 'pp_strt1', 'pp_strt2', 'pp_sy0', 'pp_sy1', 'pp_sy2',
               'pp_vka0', 'pp_vka1', 'pp_vka2'], dtype=object)
```

Fix the storage pilot points - we still have layer-scale storage pars adjustable

```
In [8]: #s_pars = par.loc[par.pargp.apply(lambda x: "pp" in x and ("ss" in x or "sy" in x)), "p"]  
        #par.loc[s_pars, "partrans"] = "fixed"  
        pst.npar_adj
```

```
Out[8]: 719
```

```
In [9]: adj_par = par.loc[par.partrans=="log", :]  
        adj_par.pargp.value_counts().sort_values()
```

```
Out[9]: cn_rech5          1  
        cn_hk7            1  
        cn_sy7            1  
        cn_hk6            1  
        cn_sy8            1  
        cn_prsity7        1  
        cn_strt8          1  
        cn_rech4          1  
        cn_ss8            1  
        cn_vka6           1  
        cn_prsity8        1  
        cn_hk8            1  
        cn_strt7          1  
        cn_strt6          1  
        cn_vka7           1  
        cn_vka8           1  
        cn_ss6            1  
        cn_ss7            1  
        cn_prsity6        1  
        cn_sy6            1  
        flow              1  
        welflux           2  
        welflux_k02       6  
        drncond_k00      10  
        pp_prsity0       32  
        pp_vka0          32  
        pp_strt2         32  
        pp_rech0         32  
        pp_ss0           32  
        pp_rech1         32  
        pp_vka1          32  
        pp_prsity1       32  
        pp_hk1           32  
        pp_sy0           32  
        pp_ss1           32  
        pp_prsity2       32  
        pp_vka2          32  
        pp_sy1           32
```

```

pp_strt0      32
pp_strt1      32
pp_ss2        32
pp_hk0        32
pp_sy2        32
pp_hk2        32
strk          40
Name: pargp, dtype: int64

```

fix the future recharge pilot points, vka in layers 1 and 3 and the initial condition pilot points (we still have layer-scale pars for each of these types)

```

In [10]: fi_grps = ["pp_rech1", "pp_vka0", "pp_vka2", "pp_strt0", "pp_strt1", "pp_strt2"]
          par.loc[par.pargp.apply(lambda x: x in fi_grps), "partrans"] = "fixed"
          pst.npar_adj

```

```

Out[10]: 527

```

Ok, thats better...so lets run PESTPP-GLM. We will use a single "base parameter" jacobian matrix as the basis for 6 super parameter iterations. Then we will draw 100 realizations from the FOSM posterior parameter covariance matrix and run those 100 realizations to get the psoterior forecast PDFs

```

In [11]: pst.control_data.noptmax = 3
          pst.pestpp_options["n_iter_base"] = -1
          pst.pestpp_options["n_iter_super"] = 3
          pst.pestpp_options["num_reals"] = 50 # this is how many ies uses
          pst.pestpp_options["parcov"] = "prior_cov.jcb"
          pst.write(os.path.join(t_d, "freyberg_pp.pst"))

In [12]: #pyemu.os_utils.start_slaves(t_d, "pestpp-glm", "freyberg_pp.pst", num_slaves=20, slave_r
          #                                     master_dir=m_d)

In [13]: df = df=pd.read_csv(os.path.join(m_d, "freyberg_pp.post.obsen.csv"), index_col=0)
          oe = pyemu.ObservationEnsemble.from_dataframe(pst=pst, df=df)

In [14]: ax = oe.phi_vector.hist()#bins=np.linspace(0, 100, 20))
          oe.phi_vector.sort_values().iloc[:20]

```

```


Out[14]: real_name
1      16.417398
23     40.379792
21    136.184310
10    188.214172
26    229.596887
8     280.037659
16    302.087583
41    303.432559
12    332.762416

```

```

37    342.523602
30    346.262108
17    412.799856
0     444.770367
15    476.092295
5     479.962961
47    499.062079
46    531.175817
31    531.603828
33    561.768950
25    562.670514
dtype: float64

```



Here we see the distribution of phi values across the 100 posterior realizations. Should we accept all of these??? The theoretical phi we should accept is number of nonzero obs (14).

To get a “posterior” ensemble, we need to throw out the realizations with large phi - lets just take the 20 best:

```
In [15]: oe_pt = oe.loc[oe.phi_vector.sort_values().index[:20],:] #just take the 20 lowest phi
```

We can also load and plot the FOSM forecast results along side of the ensemble results:

```
In [16]: f_df = pd.read_csv(os.path.join(m_d,"freyberg_pp.pred.usum.csv"),index_col=0)
         f_df.index = f_df.index.map(str.lower)
         f_df
```

```
Out[16]:
```

	prior_mean	prior_stdev	prior_lower_bound	\
name				



fa_hw_19791230	-977.2390	295.32800	-1567.8900
fa_hw_19801229	-351.2160	409.77000	-1170.7600
fa_tw_19791230	-453.0330	409.35100	-1271.7400
fa_tw_19801229	108.9600	506.73200	-904.5040
hds_00_013_002_000	39.6102	3.96314	31.6840
hds_00_013_002_001	38.3838	4.05782	30.2681
part_status	2.0000	0.00000	2.0000
part_time	907.7020	704.75100	-501.8010



	prior_upper_bound	post_mean	post_stdev \
name			
fa_hw_19791230	-386.5840	-1226.1600	253.662000
fa_hw_19801229	468.3240	-438.0650	342.835000
fa_tw_19791230	365.6690	-685.1120	185.061000
fa_tw_19801229	1122.4200	25.8356	275.087000
hds_00_013_002_000	47.5365	40.1357	0.528593
hds_00_013_002_001	46.4994	38.6689	0.821579
part_status	2.0000	2.0000	0.000000
part_time	2317.2000	864.5700	609.653000



	post_lower_bound	post_upper_bound
name		
fa_hw_19791230	-1733.4900	-718.8380
fa_hw_19801229	-1123.7400	247.6060
fa_tw_19791230	-1055.2300	-314.9900
fa_tw_19801229	-524.3380	576.0100
hds_00_013_002_000	39.0785	41.1929
hds_00_013_002_001	37.0257	40.3120
part_status	2.0000	2.0000
part_time	-354.7370	2083.8800



```
In [17]: obs = pst.observation_data
fnames = pst.pestpp_options["forecasts"].split(",")
for forecast in fnames:
    ax = plt.subplot(111)
    oe_pt.loc[:,forecast].hist(ax=ax,color="b",alpha=0.5,normed=True)
    ax.plot([obs.loc[forecast,"obsval"],obs.loc[forecast,"obsval"]],ax.get_ylim(),"r")
    axt = plt.twinx()
    x,y = pyemu.plot_utils.gaussian_distribution(f_df.loc[forecast,"prior_mean"],f_df.loc[forecast,"prior_stdev"])
    axt.fill_between(x,0,y,facecolor="0.5",alpha=0.25)
    x,y = pyemu.plot_utils.gaussian_distribution(f_df.loc[forecast,"post_mean"],f_df.loc[forecast,"post_stdev"])
    axt.fill_between(x,0,y,facecolor="b",alpha=0.25)
    axt.set_ylim(0,axt.get_ylim()[1])
    axt.set_yticks([])
    ax.set_title(forecast)
plt.show()
```











## 1.0.2 Setup of Tikhonov regularization

Now lets setup and use some formal regularization to bring the final phi up to around 14. We will use first-order regularization based on the covariance matrix we build earlier:

```
In [18]: cov = pyemu.Cov.from_binary(os.path.join(t_d,"prior_cov.jcb"))
```

```
new binary format detected...
```

```
In [19]: pyemu.helpers.first_order_pearson_tikhonov(pst,cov)
```

```
getting CC matrix  
processing
```

```
In [20]: pst.prior_information.head()
```

```
Out[20]:
```

	equation	obgnme	\
pilbl			
pcc_1	1.0 * log(dc0000390005) - 1.0 * log(dc0000390006) = 0.0	regul_cc	
pcc_2	1.0 * log(dc0000390005) - 1.0 * log(dc0000390007) = 0.0	regul_cc	
pcc_3	1.0 * log(dc0000390005) - 1.0 * log(dc0000390008) = 0.0	regul_cc	
pcc_4	1.0 * log(dc0000390005) - 1.0 * log(dc0000390009) = 0.0	regul_cc	
pcc_5	1.0 * log(dc0000390005) - 1.0 * log(dc0000390010) = 0.0	regul_cc	

	pilbl	weight
pilbl		
pcc_1	pcc_1	0.904837
pcc_2	pcc_2	0.818731
pcc_3	pcc_3	0.740818
pcc_4	pcc_4	0.670320
pcc_5	pcc_5	0.606531

```
In [21]: shutil.copy2(os.path.join(m_d,"freyberg_pp.jcb"),os.path.join(t_d,"restart_pp.jcb"))
```

```
Out[21]: 'template/restart_pp.jcb'
```

```
In [22]: pst.pestpp_options["base_jacobian"] = "restart_pp.jcb"  
pst.reg_data.phimlim = pst.nnz_obs  
pst.reg_data.phimaccept = pst.reg_data.phimlim * 1.1  
pst.write(os.path.join(t_d,"freyberg_pp.pst"))
```

```
In [23]: pyemu.os_utils.start_slaves(t_d,"pestpp-glm","freyberg_pp.pst",num_slaves=20,slave_roo  
master_dir=m_d)
```


```
In [24]: df = df=pd.read_csv(os.path.join(m_d,"freyberg_pp.post.obsen.csv"),index_col=0)  
oe = pyemu.ObservationEnsemble.from_dataframe(pst=pst,df=df)
```

```
In [25]: ax = oe.phi_vector.hist()#bins=np.linspace(0,100,20))  
oe.phi_vector.sort_values().iloc[:20]
```

```

Out[25]: real_name
36      15.065424
5       32.066449
15      37.431013
23      48.812383
8       51.256638
37      58.762065
33      58.770261
45      66.702450
48      72.439946
7      119.140950
35     135.031556
39     142.732374
21     164.359717
19     172.289419
10     189.037788
34     215.624717
17     268.322434
49     302.529537
1      380.312298
6      382.312478
dtype: float64

```



Same as before, to get a “posterior” ensemble, we need to throw out the realizations with large  $\phi$  - lets just take the 20 best:

```

In [26]: oe_pt = oe.loc[oe.phi_vector.sort_values().index[:20],:]

```

```
In [27]: f_df = pd.read_csv(os.path.join(m_d,"freyberg_pp.pred.usum.csv"),index_col=0)
f_df.index = f_df.index.map(str.lower)
f_df
```

```
Out [27]:
```

	prior_mean	prior_stdev	prior_lower_bound \
name			
fa_hw_19791230	-977.2390	295.32800	-1567.8900
fa_hw_19801229	-351.2160	409.77000	-1170.7600
fa_tw_19791230	-453.0330	409.35100	-1271.7400
fa_tw_19801229	108.9600	506.73200	-904.5040
hds_00_013_002_000	39.6102	3.96314	31.6840
hds_00_013_002_001	38.3838	4.05782	30.2681
part_status	2.0000	0.00000	2.0000
part_time	907.7020	704.75100	-501.8010

	prior_upper_bound	post_mean	post_stdev \
name			
fa_hw_19791230	-386.5840	-729.6690	253.767000
fa_hw_19801229	468.3240	-160.1280	344.320000
fa_tw_19791230	365.6690	-149.8650	180.227000
fa_tw_19801229	1122.4200	361.9180	272.345000
hds_00_013_002_000	47.5365	36.9915	0.327097
hds_00_013_002_001	46.4994	36.0961	0.710410
part_status	2.0000	1.0000	0.000000
part_time	2317.2000	4015.0000	604.279000

	post_lower_bound	post_upper_bound
name		
fa_hw_19791230	-1237.2000	-222.1350
fa_hw_19801229	-848.7680	528.5130
fa_tw_19791230	-510.3190	210.5890
fa_tw_19801229	-182.7730	906.6090
hds_00_013_002_000	36.3373	37.6457
hds_00_013_002_001	34.6753	37.5170
part_status	1.0000	1.0000
part_time	2806.4400	5223.5600

```
In [28]: obs = pst.observation_data
fnames = pst.pestpp_options["forecasts"].split(",")
for forecast in fnames:
    ax = plt.subplot(111)
    oe_pt.loc[:,forecast].hist(ax=ax,color="b",alpha=0.5,normed=True)
    ax.plot([obs.loc[forecast,"obsval"],obs.loc[forecast,"obsval"]],ax.get_ylim(),"r")
    axt = plt.twinx()
    x,y = pyemu.plot_utils.gaussian_distribution(f_df.loc[forecast,"prior_mean"],f_df.
    axt.fill_between(x,0,y,facecolor="0.5",alpha=0.25)
    x,y = pyemu.plot_utils.gaussian_distribution(f_df.loc[forecast,"post_mean"],f_df.
    axt.fill_between(x,0,y,facecolor="b",alpha=0.25)
```

```
axt.set_ylim(0,axt.get_ylim()[1])  
axt.set_yticks([])  
ax.set_title(forecast)  
plt.show()
```

