

# prior\_montecarlo

May 11, 2019

## 1 Run and process the prior monte carlo and pick a “truth” realization

A great advantage of exploring a synthetic model is that we can enforce a “truth” and then evaluate how our various attempts to estimate it perform. One way to do this is to run a monte carlo ensemble of multiple parameter realizations and then choose one of them to represent the “truth”. That will be accomplished in this notebook.

```
In [1]: import os
import shutil
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib as mpl
plt.rcParams['font.size']=12
import flopy
import pyemu
```

flopy is installed in /Users/jeremyw/Dev/gw1876/activities\_2day\_mfm/notebooks/flopy

### 1.0.1 set the t\_d or “template directory” variable to point at the template folder and read in the PEST control file

```
In [2]: t_d = "template"
pst = pyemu.Pst(os.path.join(t_d, "freyberg.pst"))
```

### 1.0.2 Decide what pars are uncertain in the truth

We need to decide what our truth looks like - should the pilot points or the grid-scale pars be the source of spatial variability? or both?

```
In [3]: par = pst.parameter_data
# grid pars
#should_fix = par.loc[par.pargp.apply(lambda x: "gr" in x), "parname"]
# pp pars
#should_fix = par.loc[par.pargp.apply(lambda x: "pp" in x), "parname"]
#pst.npar - should_fix.shape[0]
```

```
In [4]: pe = pyemu.ParameterEnsemble.from_binary(pst=pst,filename=os.path.join(t_d,"prior.jcb")
        #pe.loc[:,should_fix] = 1.0
        pe.to_csv(os.path.join(t_d,"sweep_in.csv"))
```

new binary format detected...

### 1.0.3 run the prior ensemble in parallel locally

This takes advantage of the program pestpp-swp which runs a parameter sweep through a set of parameters. By default, pestpp-swp reads in the ensemble from a file called sweep\_in.csv which in this case we made just above.

```
In [5]: m_d = "master_prior_sweep"
        pyemu.os_utils.start_slaves(t_d,"pestpp-swp","freyberg.pst",num_slaves=20,slave_root="
```

### 1.0.4 Load the output ensemble and plot a few things

```
In [6]: obs_df = pd.read_csv(os.path.join(m_d,"sweep_out.csv"),index_col=0)
        print('number of realization in the ensemble before dropping: ' + str(obs_df.shape[0]))
```

number of realization in the ensemble before dropping: 200

drop any failed runs

```
In [7]: obs_df = obs_df.loc[obs_df.failed_flag==0,:]
        print('number of realization in the ensemble **after** dropping: ' + str(obs_df.shape[
```

number of realization in the ensemble \*\*after\*\* dropping: 200


### 1.0.5 confirm which quantities were identified as forecasts



```
In [8]: fnames = pst.pestpp_options["forecasts"].split(',')
        fnames
```



```
Out[8]: ['fa_hw_19791230',
        'fa_hw_19801229',
        'fa_tw_19791230',
        'fa_tw_19801229',
        'hds_00_013_002_000',
        'hds_00_013_002_001',
        'part_time',
        'part_status']
```



### 1.0.6 now we can plot the distributions of each forecast

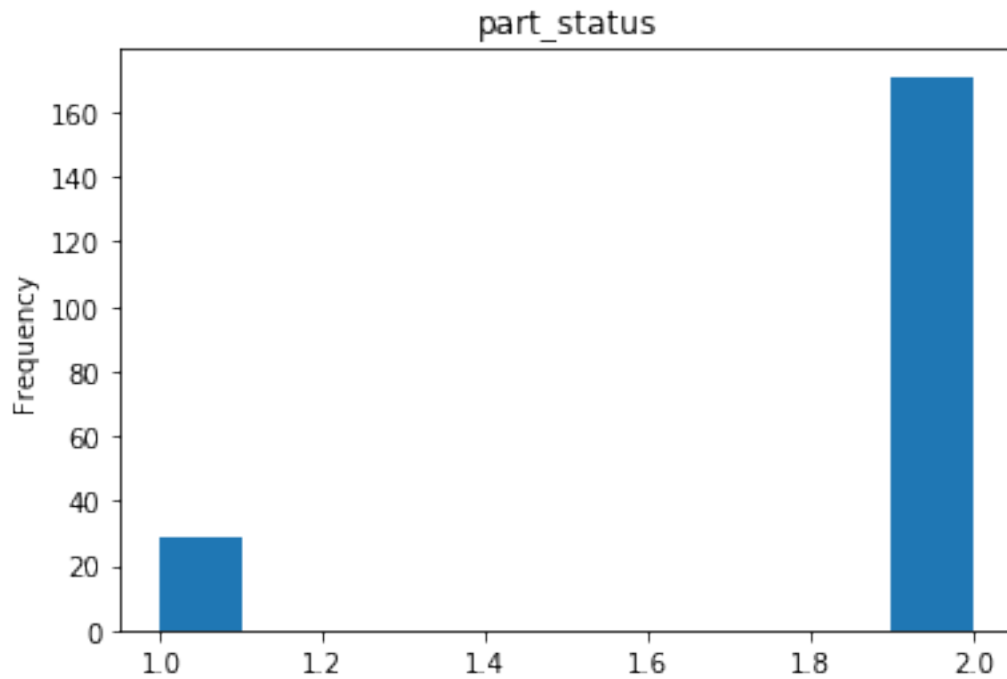
```
In [9]: for forecast in fnames:
        plt.figure()
        ax = obs_df.loc[:,forecast].plot(kind="hist")
        ax.set_title(forecast)
        plt.show()
```














We see that under scenario conditions, many more realizations for the flow to the aquifer in the headwaters are positive (as expected). Lets difference these two:

```
In [10]: sfnames = [f for f in fnames if "1980" in f or "_001" in f]
          hfnames = [f for f in fnames if "1979" in f or "_000" in f]
          diff = obs_df.loc[:,hfnames].values - obs_df.loc[:,sfnames].values
          diff = pd.DataFrame(diff,columns=sfnames)
          diff.hist(figsize=(10,10))
          plt.show()
```



We now see that the most extreme scenario yields a large decrease in flow from the aquifer to the headwaters (the most negative value)

### 1.0.7 setting the “truth”

We just need to replace the observed values (obsval) in the control file with the outputs for one of the realizations on obs\_df. In this way, we now have the nonzero values for history matching, but also the truth values for comparing how we are doing with other unobserved quantities. I’m going to pick a realization that yields an “average” variability of the observed gw levels:

```
In [11]: # choose the realization with a low historic gw to sw headwater flux
#hist_swgw = obs_df.loc[:, "fa_hw_19791230"].sort_values()
hist_swgw = obs_df.loc[:, "part_time"].sort_values()
idx = hist_swgw.index[20]
```



```
idx  
hist_swgw
```

```
Out[11]: run_id  
149      198.1387  
41       265.3273  
165      276.1939  
120      293.1571  
37       313.4481  
96       325.5403  
0        342.4888  
173      358.4053  
77       361.9629  
64       372.9315  
73       374.5064  
145      380.7252  
16       394.1340  
40       395.1985  
35       405.8629  
128      424.2145  
78       427.1186  
51       447.5162  
21       451.1969  
117      451.4228  
194      458.4655  
131      463.7558  
27       464.6402  
184      476.3412  
57       484.2876  
19       501.6304  
162      520.8668  
182      524.6132  
107      525.4871  
141      529.1481  
      ...  
192      3467.2040  
101      4015.0000  
17       4015.0000  
187      4015.0000  
191      4015.0000  
118      4015.0000  
20       4015.0000  
193      4015.0000  
6        4015.0000  
3        4015.0000  
12       4015.0000  
183      4015.0000  
138      4015.0000
```

```

100    4015.0000
157    4015.0000
61     4015.0000
47     4015.0000
166    4015.0000
46     4015.0000
153    4015.0000
24     4015.0000
75     4015.0000
82     4015.0000
88     4015.0000
175    4015.0000
176    4015.0000
159    4015.0000
29     4015.0000
76     4015.0000
130    4015.0000
Name: part_time, Length: 200, dtype: float64

```

```
In [12]: obs_df.loc[idx,pst.nnz_obs_names]
```

```

Out[12]: fo_39_19791230      11302.000000
hds_00_002_009_000         36.058125
hds_00_002_015_000         35.212799
hds_00_003_008_000         36.083511
hds_00_009_001_000         37.389400
hds_00_013_010_000         35.048706
hds_00_015_016_000         34.838463
hds_00_021_010_000         34.814125
hds_00_022_015_000         34.552216
hds_00_024_004_000         35.209511
hds_00_026_006_000         34.681568
hds_00_029_015_000         34.122902
hds_00_033_007_000         33.345665
hds_00_034_010_000         33.317856
Name: 194, dtype: float64

```

Lets see how our selected truth does with the sw/gw forecasts:

```
In [13]: obs_df.loc[idx,fnames]
```

```

Out[13]: fa_hw_19791230      -826.113600
fa_hw_19801229      -412.139600
fa_tw_19791230       341.545050
fa_tw_19801229       704.034190
hds_00_013_002_000     36.989079
hds_00_013_002_001     36.173130
part_time             458.465500
part_status           2.000000
Name: 194, dtype: float64

```

Assign some initial weights. Now, it is custom to add noise to the observed values... we will use the classic Gaussian noise... zero mean and standard deviation of 1 over the weight

```
In [14]: pst = pyemu.Pst(os.path.join(t_d,"freyberg.pst"))
obs = pst.observation_data
obs.loc[:, "obsval"] = obs_df.loc[idx, pst.obs_names]
obs.loc[obs.obgnme=="calhead", "weight"] = 10.0
obs.loc[obs.obgnme=="calflux", "weight"] = 1.0
```

here we just get a sample from a random normal distribution with mean=0 and std=1. The argument indicates how many samples we want - and we choose `pst.nnz_obs` which is the the number of nonzero-weighted observations in the PST file

```
In [15]: np.random.seed(seed=0)
snd = np.random.randn(pst.nnz_obs)
noise = snd * 1./obs.loc[pst.nnz_obs_names, "weight"]
pst.observation_data.loc[noise.index, "obsval"] += noise
noise
```

```
Out[15]: obsnme
fo_39_19791230      1.764052
hds_00_002_009_000  0.040016
hds_00_002_015_000  0.097874
hds_00_003_008_000  0.224089
hds_00_009_001_000  0.186756
hds_00_013_010_000 -0.097728
hds_00_015_016_000  0.095009
hds_00_021_010_000 -0.015136
hds_00_022_015_000 -0.010322
hds_00_024_004_000  0.041060
hds_00_026_006_000  0.014404
hds_00_029_015_000  0.145427
hds_00_033_007_000  0.076104
hds_00_034_010_000  0.012168
Name: weight, dtype: float64
```

Then we write this out to a new file and run `pestpp-ies` to see how the objective function looks

```
In [16]: pst.write(os.path.join(t_d,"freyberg.pst"))
pyemu.os_utils.run("pestpp-ies freyberg.pst", cwd=t_d)
```

```
noptmax:0, npar_adj:14819, nnz_obs:14
```

Now we can read in the results and make some figures showing residuals and the balance of the objective function


```
In [17]: pst = pyemu.Pst(os.path.join(t_d,"freyberg.pst"))
print(pst.phi)
plt.figure()
pst.plot(kind='phi_pie')
print('Here are the non-zero weighted observation names')

figs = pst.plot(kind="1to1")
plt.show()
pst.res.loc[pst.nnz_obs_names,:]
```

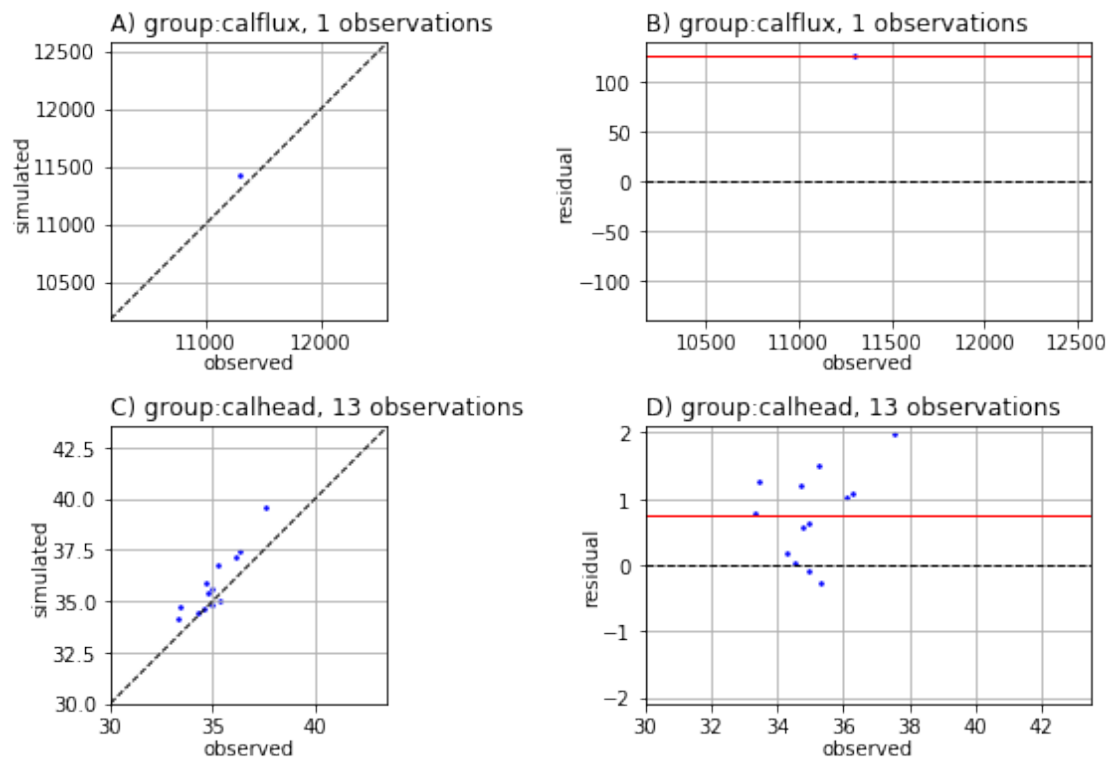
17221.065640613586

Here are the non-zero weighted observation names

<Figure size 432x288 with 0 Axes>



<Figure size 576x756 with 0 Axes>



Out[17]:

	name	group	measured	modelled	\
name					
fo_39_19791230	fo_39_19791230	calflux	11303.764052	11430.000000	
hds_00_002_009_000	hds_00_002_009_000	calhead	36.098140	37.107498	
hds_00_002_015_000	hds_00_002_015_000	calhead	35.310673	35.045185	
hds_00_003_008_000	hds_00_003_008_000	calhead	36.307601	37.397289	
hds_00_009_001_000	hds_00_009_001_000	calhead	37.576156	39.546417	
hds_00_013_010_000	hds_00_013_010_000	calhead	34.950978	35.571774	
hds_00_015_016_000	hds_00_015_016_000	calhead	34.933472	34.835716	
hds_00_021_010_000	hds_00_021_010_000	calhead	34.798989	35.386250	
hds_00_022_015_000	hds_00_022_015_000	calhead	34.541894	34.577492	
hds_00_024_004_000	hds_00_024_004_000	calhead	35.250571	36.760464	
hds_00_026_006_000	hds_00_026_006_000	calhead	34.695973	35.896149	
hds_00_029_015_000	hds_00_029_015_000	calhead	34.268329	34.453842	
hds_00_033_007_000	hds_00_033_007_000	calhead	33.421769	34.678810	
hds_00_034_010_000	hds_00_034_010_000	calhead	33.330023	34.118073	
	residual	weight			
name					
fo_39_19791230	-126.235948	1.0			
hds_00_002_009_000	-1.009358	10.0			
hds_00_002_015_000	0.265488	10.0			
hds_00_003_008_000	-1.089689	10.0			
hds_00_009_001_000	-1.970261	10.0			
hds_00_013_010_000	-0.620795	10.0			
hds_00_015_016_000	0.097755	10.0			
hds_00_021_010_000	-0.587260	10.0			
hds_00_022_015_000	-0.035598	10.0			
hds_00_024_004_000	-1.509893	10.0			
hds_00_026_006_000	-1.200176	10.0			
hds_00_029_015_000	-0.185513	10.0			
hds_00_033_007_000	-1.257041	10.0			
hds_00_034_010_000	-0.788049	10.0			


Publication ready figs - oh snap!

Depending on the truth you chose, we may have a problem - we set the weights for both the heads and the flux to reasonable values based on what we expect for measurement noise. But the contributions to total phi might be out of balance - if contribution of the flux measurement to total phi is too low, the history matching excersizes (coming soon!) will focus almost entirely on minimizing head residuals. So we need to balance the objective function. This is a subtle but very important step, especially since some of our forecasts deal with sw-gw exchange

```
In [18]: pc = pst.phi_components
         #target = {"calflux":0.3 * pc["calhead"]}
         target = {"calhead":500,"calflux":500}
```

```
pst.adjust_weights(obsgrp_dict=target)
pst.plot(kind='phi_pie')
```

Out[18]: <matplotlib.axes.\_subplots.AxesSubplot at 0x181d85e5c0>



Lets see what the new flux observation weight is:

```
In [19]: pst.observation_data.loc[pst.nnz_obs_names, "weight"]
```

```
Out[19]: obsnme
fo_39_19791230      0.177134
hds_00_002_009_000  6.236491
hds_00_002_015_000  6.236491
hds_00_003_008_000  6.236491
```

```

hds_00_009_001_000    6.236491
hds_00_013_010_000    6.236491
hds_00_015_016_000    6.236491
hds_00_021_010_000    6.236491
hds_00_022_015_000    6.236491
hds_00_024_004_000    6.236491
hds_00_026_006_000    6.236491
hds_00_029_015_000    6.236491
hds_00_033_007_000    6.236491
hds_00_034_010_000    6.236491
Name: weight, dtype: float64

```

Now, for some super trickery: since we changed the weight, we need to generate the observation noise using these new weights for the error model (so meta!)

```

In [20]: obs = pst.observation_data
         np.random.seed(seed=0)
         snd = np.random.randn(pst.nnz_obs)
         noise = snd * 1./obs.loc[pst.nnz_obs_names,"weight"]
         obs.loc[:, "obsval"] = obs_df.loc[idx, pst.obs_names]
         pst.observation_data.loc[noise.index, "obsval"] += noise
         noise

```

```

Out [20]: obsnme
fo_39_19791230        9.958857
hds_00_002_009_000    0.064164
hds_00_002_015_000    0.156937
hds_00_003_008_000    0.359320
hds_00_009_001_000    0.299457
hds_00_013_010_000   -0.156703
hds_00_015_016_000    0.152343
hds_00_021_010_000   -0.024270
hds_00_022_015_000   -0.016551
hds_00_024_004_000    0.065838
hds_00_026_006_000    0.023097
hds_00_029_015_000    0.233188
hds_00_033_007_000    0.122030
hds_00_034_010_000    0.019510
Name: weight, dtype: float64

```

```

In [21]: pst.write(os.path.join(t_d, "freyberg.pst"))
         pyemu.os_utils.run("pestpp-ies freyberg.pst", cwd=t_d)
         pst = pyemu.Pst(os.path.join(t_d, "freyberg.pst"))
         print(pst.phi)
         pst.plot(kind='phi_pie')
         plt.show()


```

```

noptmax:0, npar_adj:14819, nnz_obs:14
903.6102175392298

```





Whew! confused yet? Ok, let's leave all this confusion behind...its mostly academic, just to make sure we are using weights that are in harmony with the noise we added to the truth...Just to make sure we have everything working right, we should be able to load the truth parameters, run the model once and have a phi equivalent to the noise vector:

```
In [22]: par_df = pd.read_csv(os.path.join(m_d,"sweep_in.csv"),index_col=0)
          pst.parameter_data.loc[:, "parval1"] = par_df.loc[idx,pst.par_names]
          pst.write(os.path.join(m_d,"test.pst"))
```

```
noptmax:0, npar_adj:14819, nnz_obs:14
```

we will run this with noptmax=0 to preform a single run. Pro-tip: you can use any of the pestpp-### binaries/executables to run noptmax=0

```
In [23]: pyemu.os_utils.run("pestpp-ies.exe test.pst", cwd=m_d)
        pst = pyemu.Pst(os.path.join(m_d, "test.pst"))
        print(pst.phi)
        pst.res.loc[pst.nnz_obs_names,:]
```

17.52884697222985

```
Out [23]:
```

	name	group	measured	modelled \
	name			
	fo_39_19791230	fo_39_19791230	calflux	11311.958857 11302.000000
	hds_00_002_009_000	hds_00_002_009_000	calhead	36.122288 36.058125
	hds_00_002_015_000	hds_00_002_015_000	calhead	35.369736 35.212799
	hds_00_003_008_000	hds_00_003_008_000	calhead	36.442831 36.083511
	hds_00_009_001_000	hds_00_009_001_000	calhead	37.688857 37.389400
	hds_00_013_010_000	hds_00_013_010_000	calhead	34.892003 35.048706
	hds_00_015_016_000	hds_00_015_016_000	calhead	34.990806 34.838463
	hds_00_021_010_000	hds_00_021_010_000	calhead	34.789855 34.814125
	hds_00_022_015_000	hds_00_022_015_000	calhead	34.535665 34.552216
	hds_00_024_004_000	hds_00_024_004_000	calhead	35.275349 35.209511
	hds_00_026_006_000	hds_00_026_006_000	calhead	34.704665 34.681568
	hds_00_029_015_000	hds_00_029_015_000	calhead	34.356090 34.122902
	hds_00_033_007_000	hds_00_033_007_000	calhead	33.467695 33.345665
	hds_00_034_010_000	hds_00_034_010_000	calhead	33.337366 33.317856

	residual	weight
name		
fo_39_19791230	9.958857	0.177134
hds_00_002_009_000	0.064164	6.236491
hds_00_002_015_000	0.156937	6.236491
hds_00_003_008_000	0.359320	6.236491
hds_00_009_001_000	0.299457	6.236491
hds_00_013_010_000	-0.156703	6.236491
hds_00_015_016_000	0.152343	6.236491
hds_00_021_010_000	-0.024270	6.236491
hds_00_022_015_000	-0.016551	6.236491
hds_00_024_004_000	0.065838	6.236491
hds_00_026_006_000	0.023097	6.236491
hds_00_029_015_000	0.233188	6.236491
hds_00_033_007_000	0.122030	6.236491
hds_00_034_010_000	0.019510	6.236491

The residual should be exactly the noise values from above. Lets load the model (that was just run using the true pars) and check some things

```
In [24]: m = flopy.modflow.Modflow.load("freyberg.nam", model_ws=m_d)
```


```
In [25]: a = m.upw.hk[0].array
        #a = m.rch.rech[0].array
```

```

a = np.ma.masked_where(m.bas6.ibound[0].array==0,a)
print(a.min(),a.max())
c = plt.imshow(a)
plt.colorbar()
plt.show()

```


0.4185674 28.04684



```

In [26]: lst = flopy.utils.MfListBudget(os.path.join(m_d,"freyberg.list"))
df = lst.get_dataframes(diff=True)[0]
ax = df.plot(kind="bar",figsize=(10,10), grid=True)
a = ax.set_xticklabels(["historic","scenario"],rotation=90)
plt.show(ax)



```


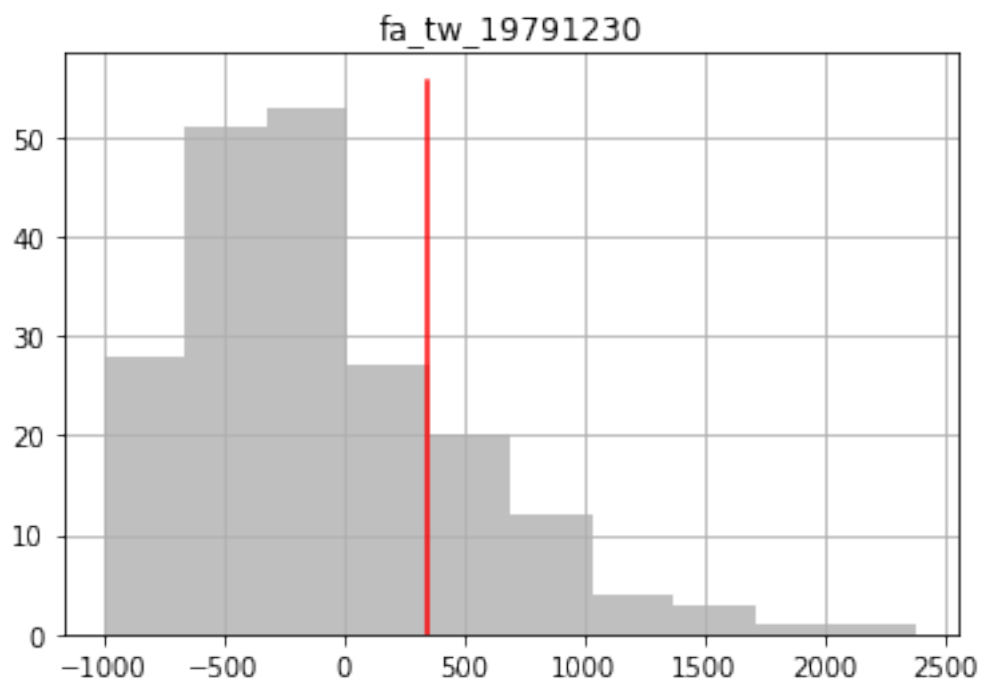




### 1.0.8 see how our existing observation ensemble compares to the truth

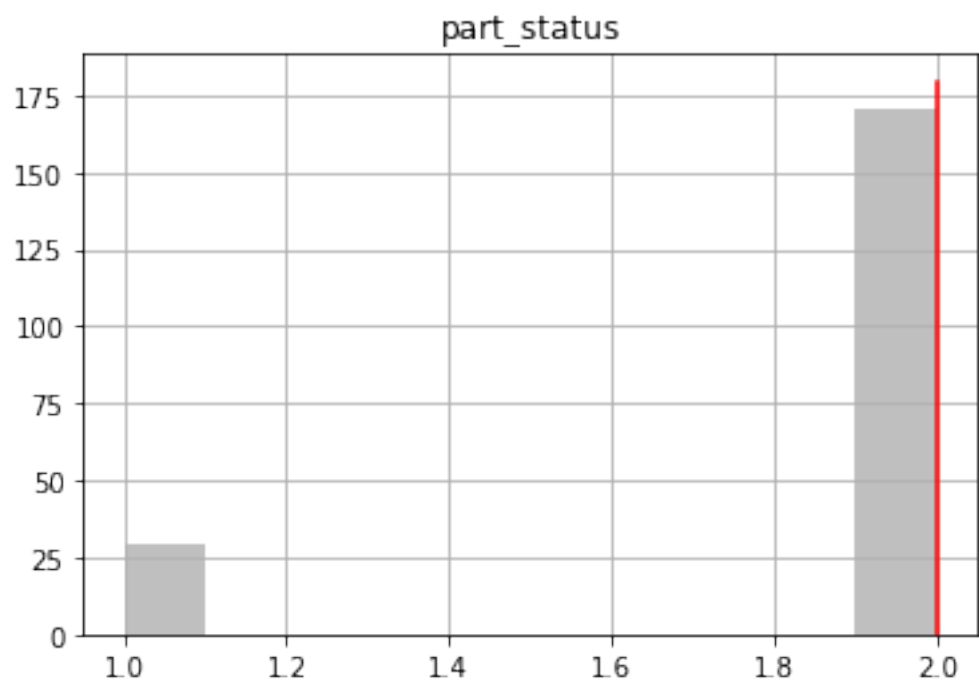

forecasts:

```
In [27]: obs = pst.observation_data
plt.figure()
for forecast in fnames:
    ax = plt.subplot(111)
    obs_df.loc[:,forecast].hist(ax=ax,color="0.5",alpha=0.5)
    ax.plot([obs.loc[forecast,"obsval"],obs.loc[forecast,"obsval"]],ax.get_ylim(),"r")
    ax.set_title(forecast)
plt.show()
```









observations:



```
In [28]: for oname in pst.nnz_obs_names:
          ax = plt.subplot(111)
          obs_df.loc[:, oname].hist(ax=ax, color="0.5", alpha=0.5)
          ax.plot([obs.loc[oname, "obsval"], obs.loc[oname, "obsval"]], ax.get_ylim(), "r")
          ax.set_title(oname)
          plt.show()
```

