

prior_montecarlo

May 8, 2019

1 Run and process the prior monte carlo and pick a “truth” realization

A great advantage of exploring a synthetic model is that we can enforce a “truth” and then evaluate how our various attempts to estimate it perform. One way to do this is to run a monte carlo ensemble of multiple parameter realizations and then choose one of them to represent the “truth”. That will be accomplished in this notebook.

```
In [1]: import os
import shutil
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib as mpl
plt.rcParams['font.size']=12
import flogy
import pyemu
```

flogy is installed in /Users/jeremyw/Dev/gw1876/activities_2day_mfm/notebooks/flogy

1.0.1 set the t_d or “template directory” variable to point at the template folder and read in the PEST control file

```
In [2]: t_d = "template"
pst = pyemu.Pst(os.path.join(t_d, "freyberg.pst"))
```

1.0.2 Decide what pars are uncertain in the truth

We need to decide what our truth looks like - should the pilot points or the grid-scale pars be the source of spatial variability? or both?

```
In [3]: par = pst.parameter_data
# grid pars
#should_fix = par.loc[par.pargp.apply(lambda x: "gr" in x), "parname"]
# pp pars
#should_fix = par.loc[par.pargp.apply(lambda x: "pp" in x), "parname"]
#pst.npar - should_fix.shape[0]
```

```
In [4]: pe = pyemu.ParameterEnsemble.from_binary(pst=pst,filename=os.path.join(t_d,"prior.jcb")
        #pe.loc[:,should_fix] = 1.0
        pe.to_csv(os.path.join(t_d,"sweep_in.csv"))
```

new binary format detected...

1.0.3 run the prior ensemble in parallel locally

This takes advantage of the program pestpp-swp which runs a parameter sweep through a set of parameters. By default, pestpp-swp reads in the ensemble from a file called sweep_in.csv which in this case we made just above.

```
In [5]: m_d = "master_prior_sweep"
        pyemu.os_utils.start_slaves(t_d,"pestpp-swp","freyberg.pst",num_slaves=20,slave_root="
```

1.0.4 Load the output ensemble and plot a few things

```
In [6]: obs_df = pd.read_csv(os.path.join(m_d,"sweep_out.csv"),index_col=0)
        print('number of realization in the ensemble before dropping: ' + str(obs_df.shape[0]))
```

number of realization in the ensemble before dropping: 200

drop any failed runs

```
In [7]: obs_df = obs_df.loc[obs_df.failed_flag==0,:]
        print('number of realization in the ensemble **after** dropping: ' + str(obs_df.shape[
```

number of realization in the ensemble **after** dropping: 200


1.0.5 confirm which quantities were identified as forecasts



```
In [8]: fnames = pst.pestpp_options["forecasts"].split(',')
        fnames
```



```
Out[8]: ['fa_hw_19791230',
        'fa_hw_19801229',
        'fa_tw_19791230',
        'fa_tw_19801229',
        'hds_00_013_002_000',
        'hds_00_013_002_001',
        'part_time',
        'part_status']
```



1.0.6 now we can plot the distributions of each forecast


```
In [9]: for forecast in fnames:
        plt.figure()
        ax = obs_df.loc[:,forecast].plot(kind="hist")
        ax.set_title(forecast)
        plt.show()
```














We see that under scenario conditions, many more realizations for the flow to the aquifer in the headwaters are positive (as expected). Lets difference these two:

```
In [10]: sfnames = [f for f in fnames if "1980" in f or "_001" in f]
          hfnames = [f for f in fnames if "1979" in f or "_000" in f]
          diff = obs_df.loc[:,hfnames].values - obs_df.loc[:,sfnames].values
          diff = pd.DataFrame(diff,columns=sfnames)
          diff.hist(figsize=(10,10))
          plt.show()
```



We now see that the most extreme scenario yields a large decrease in flow from the aquifer to the headwaters (the most negative value)

1.0.7 setting the “truth”

We just need to replace the observed values (obsval) in the control file with the outputs for one of the realizations on obs_df. In this way, we now have the nonzero values for history matching, but also the truth values for comparing how we are doing with other unobserved quantities. I’m going to pick a realization that yields an “average” variability of the observed gw levels:

```
In [11]: # choose the realization with a low historic gw to sw headwater flux
#hist_swgw = obs_df.loc[:, "fa_hw_19791230"].sort_values()
hist_swgw = obs_df.loc[:, "part_time"].sort_values()
idx = hist_swgw.index[20]
```



```
idx  
hist_swgw
```

```
Out[11]: run_id  
149      184.1517  
166      206.2137  
77       276.1604  
59       314.8956  
100      331.3931  
174      338.6796  
35       343.6858  
161      351.1038  
120      365.6083  
82       376.8029  
0        384.3066  
73       394.6332  
98       398.3461  
41       401.3468  
21       402.4046  
180      408.5494  
52       420.0638  
168      424.3355  
57       446.0459  
40       458.2126  
3        462.9558  
118      475.5129  
132      477.4851  
184      477.9743  
117      479.2090  
131      487.5205  
37       502.1088  
34       513.7058  
56       521.4803  
95       521.8308  
      ...  
4        2438.2690  
178      2617.6400  
115      2927.8600  
116      2943.3860  
32       3007.2270  
83       3019.2510  
158      3052.2840  
17       4015.0000  
5        4015.0000  
24       4015.0000  
43       4015.0000  
25       4015.0000  
147      4015.0000
```

```

12      4015.0000
165     4015.0000
69      4015.0000
96      4015.0000
157     4015.0000
159     4015.0000
51      4015.0000
175     4015.0000
173     4015.0000
55      4015.0000
78      4015.0000
76      4015.0000
137     4015.0000
64      4015.0000
65      4015.0000
108     4015.0000
169     4015.0000
Name: part_time, Length: 200, dtype: float64

```

```
In [12]: obs_df.loc[idx,pst.nnz_obs_names]
```

```

Out[12]: fo_39_19791230      9714.600000
hds_00_002_009_000      35.826954
hds_00_002_015_000      34.983761
hds_00_003_008_000      35.923908
hds_00_009_001_000      36.820004
hds_00_013_010_000      34.999874
hds_00_015_016_000      34.566364
hds_00_021_010_000      34.772247
hds_00_022_015_000      34.603825
hds_00_024_004_000      34.965450
hds_00_026_006_000      34.759605
hds_00_029_015_000      34.272018
hds_00_033_007_000      34.105034
hds_00_034_010_000      33.770840
Name: 3, dtype: float64

```

Lets see how our selected truth does with the sw/gw forecasts:

```
In [13]: obs_df.loc[idx,fnames]
```

```

Out[13]: fa_hw_19791230      -1123.553000
fa_hw_19801229      -197.152255
fa_tw_19791230       58.586980
fa_tw_19801229      631.809570
hds_00_013_002_000      36.341404
hds_00_013_002_001      35.233444
part_time      462.955800
part_status      2.000000
Name: 3, dtype: float64

```

Assign some initial weights. Now, it is custom to add noise to the observed values... we will use the classic Gaussian noise... zero mean and standard deviation of 1 over the weight

```
In [14]: pst = pyemu.Pst(os.path.join(t_d,"freyberg.pst"))
obs = pst.observation_data
obs.loc[:, "obsval"] = obs_df.loc[idx, pst.obs_names]
obs.loc[obs.obgnme=="calhead", "weight"] = 10.0
obs.loc[obs.obgnme=="calflux", "weight"] = 1.0
```

here we just get a sample from a random normal distribution with mean=0 and std=1. The argument indicates how many samples we want - and we choose `pst.nnz_obs` which is the the number of nonzero-weighted observations in the PST file

```
In [15]: np.random.seed(seed=0)
snd = np.random.randn(pst.nnz_obs)
noise = snd * 1./obs.loc[pst.nnz_obs_names, "weight"]
pst.observation_data.loc[noise.index, "obsval"] += noise
noise
```

```
Out[15]: obsnme
fo_39_19791230      1.764052
hds_00_002_009_000  0.040016
hds_00_002_015_000  0.097874
hds_00_003_008_000  0.224089
hds_00_009_001_000  0.186756
hds_00_013_010_000 -0.097728
hds_00_015_016_000  0.095009
hds_00_021_010_000 -0.015136
hds_00_022_015_000 -0.010322
hds_00_024_004_000  0.041060
hds_00_026_006_000  0.014404
hds_00_029_015_000  0.145427
hds_00_033_007_000  0.076104
hds_00_034_010_000  0.012168
Name: weight, dtype: float64
```

Then we write this out to a new file and run `pestpp-ies` to see how the objective function looks

```
In [16]: pst.write(os.path.join(t_d,"freyberg.pst"))
pyemu.os_utils.run("pestpp-ies freyberg.pst", cwd=t_d)
```

```
noptmax:0, npar_adj:14819, nnz_obs:14
```

Now we can read in the results and make some figures showing residuals and the balance of the objective function

```

In [17]: pst = pyemu.Pst(os.path.join(t_d,"freyberg.pst"))
         print(pst.phi)
         plt.figure()
         pst.plot(kind='phi_pie')
         print('Here are the non-zero weighted observation names')


         figs = pst.plot(kind="1to1")
         plt.show()
         pst.res.loc[pst.nnz_obs_names,:]

```

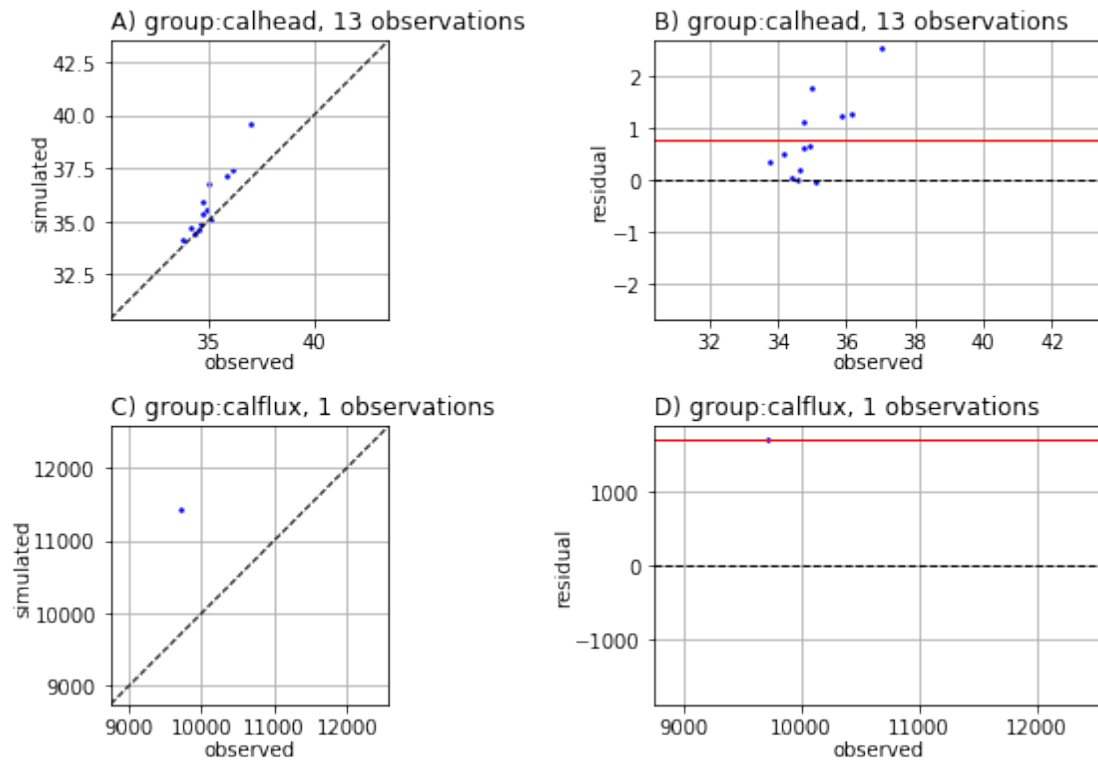
2938060.4131643856

Here are the non-zero weighted observation names

<Figure size 432x288 with 0 Axes>



<Figure size 576x756 with 0 Axes>



```

Out [17]:

```

	name	group	measured	modelled \
	name			
	fo_39_19791230	fo_39_19791230	calflux	9716.364052 11430.000000
	hds_00_002_009_000	hds_00_002_009_000	calhead	35.866970 37.107498
	hds_00_002_015_000	hds_00_002_015_000	calhead	35.081635 35.045185
	hds_00_003_008_000	hds_00_003_008_000	calhead	36.147998 37.397289
	hds_00_009_001_000	hds_00_009_001_000	calhead	37.006759 39.546417
	hds_00_013_010_000	hds_00_013_010_000	calhead	34.902146 35.571774
	hds_00_015_016_000	hds_00_015_016_000	calhead	34.661373 34.835716
	hds_00_021_010_000	hds_00_021_010_000	calhead	34.757112 35.386250
	hds_00_022_015_000	hds_00_022_015_000	calhead	34.593503 34.577492
	hds_00_024_004_000	hds_00_024_004_000	calhead	35.006510 36.760464
	hds_00_026_006_000	hds_00_026_006_000	calhead	34.774010 35.896149
	hds_00_029_015_000	hds_00_029_015_000	calhead	34.417446 34.453842
	hds_00_033_007_000	hds_00_033_007_000	calhead	34.181138 34.678810
	hds_00_034_010_000	hds_00_034_010_000	calhead	33.783007 34.118073

	residual	weight
name		
fo_39_19791230	-1713.635948	1.0
hds_00_002_009_000	-1.240529	10.0
hds_00_002_015_000	0.036450	10.0
hds_00_003_008_000	-1.249292	10.0
hds_00_009_001_000	-2.539658	10.0
hds_00_013_010_000	-0.669627	10.0
hds_00_015_016_000	-0.174343	10.0
hds_00_021_010_000	-0.629138	10.0
hds_00_022_015_000	0.016011	10.0
hds_00_024_004_000	-1.753954	10.0
hds_00_026_006_000	-1.122139	10.0
hds_00_029_015_000	-0.036396	10.0
hds_00_033_007_000	-0.497672	10.0
hds_00_034_010_000	-0.335065	10.0

Publication ready figs - oh snap!

Depending on the truth you chose, we may have a problem - we set the weights for both the heads and the flux to reasonable values based on what we expect for measurement noise. But the contributions to total phi might be out of balance - if contribution of the flux measurement to total phi is too low, the history matching excersizes (coming soon!) will focus almost entirely on minimizing head residuals. So we need to balance the objective function. This is a subtle but very important step, especially since some of our forecasts deal with sw-gw exchange

```


In [18]: pc = pst.phi_components
         #target = {"calflux":0.3 * pc["calhead"]}
         target = {"calhead":500,"calflux":500}
         pst.adjust_weights(obsgrp_dict=target)
         pst.plot(kind='phi_pie')

```

```

Out [18]: <matplotlib.axes._subplots.AxesSubplot at 0x18264b24e0>

```



Lets see what the new flux observation weight is:

```
In [19]: pst.observation_data.loc[pst.nnz_obs_names, "weight"]
```

```
Out [19]: obsnme  
fo_39_19791230      0.013049  
hds_00_002_009_000  5.750067  
hds_00_002_015_000  5.750067  
hds_00_003_008_000  5.750067  
hds_00_009_001_000  5.750067  
hds_00_013_010_000  5.750067  
hds_00_015_016_000  5.750067  
hds_00_021_010_000  5.750067  
hds_00_022_015_000  5.750067
```

```

hds_00_024_004_000    5.750067
hds_00_026_006_000    5.750067
hds_00_029_015_000    5.750067
hds_00_033_007_000    5.750067
hds_00_034_010_000    5.750067
Name: weight, dtype: float64

```

Now, for some super trickery: since we changed the weight, we need to generate the observation noise using these new weights for the error model (so meta!)

```

In [20]: obs = pst.observation_data
         np.random.seed(seed=0)
         snd = np.random.randn(pst.nnz_obs)
         noise = snd * 1./obs.loc[pst.nnz_obs_names, "weight"]
         obs.loc[:, "obsval"] = obs_df.loc[idx, pst.obs_names]
         pst.observation_data.loc[noise.index, "obsval"] += noise
         noise

```

```

Out[20]: obsnme
fo_39_19791230        135.190144
hds_00_002_009_000     0.069592
hds_00_002_015_000     0.170213
hds_00_003_008_000     0.389716
hds_00_009_001_000     0.324789
hds_00_013_010_000    -0.169959
hds_00_015_016_000     0.165231
hds_00_021_010_000    -0.026323
hds_00_022_015_000    -0.017951
hds_00_024_004_000     0.071408
hds_00_026_006_000     0.025051
hds_00_029_015_000     0.252914
hds_00_033_007_000     0.132353
hds_00_034_010_000     0.021161
Name: weight, dtype: float64

```

```


In [21]: pst.write(os.path.join(t_d, "freyberg.pst"))
         pyemu.os_utils.run("pestpp-ies freyberg.pst", cwd=t_d)
         pst = pyemu.Pst(os.path.join(t_d, "freyberg.pst"))
         print(pst.phi)
         pst.plot(kind='phi_pie')
         plt.show()

```

```

noptmax:0, npar_adj:14819, nnz_obs:14
884.8855030468067

```

Whew! confused yet? Ok, let's leave all this confusion behind...its mostly academic, just to make sure we are using weights that are in harmony with the noise we added to the truth...Just to make sure we have everything working right, we should be able to load the truth parameters, run the model once and have a phi equivalent to the noise vector:

```
In [22]: par_df = pd.read_csv(os.path.join(m_d,"sweep_in.csv"),index_col=0)
         pst.parameter_data.loc[:, "parval1"] = par_df.loc[idx,pst.par_names]
         pst.write(os.path.join(m_d,"test.pst"))
```

```
noptmax:0, npar_adj:14819, nnz_obs:14
```

we will run this with noptmax=0 to preform a single run. Pro-tip: you can use any of the pestpp-### binaries/executables to run noptmax=0

```
In [23]: pyemu.os_utils.run("pestpp-ies.exe test.pst", cwd=m_d)
        pst = pyemu.Pst(os.path.join(m_d, "test.pst"))
        print(pst.phi)
        pst.res.loc[pst.nnz_obs_names,:]
```

17.528847226654747

```
Out [23]:
```

	name	group	measured	modelled \
	name			
	fo_39_19791230	fo_39_19791230	calflux	9849.790144 9714.600000
	hds_00_002_009_000	hds_00_002_009_000	calhead	35.896546 35.826954
	hds_00_002_015_000	hds_00_002_015_000	calhead	35.153974 34.983761
	hds_00_003_008_000	hds_00_003_008_000	calhead	36.313624 35.923908
	hds_00_009_001_000	hds_00_009_001_000	calhead	37.144792 36.820004
	hds_00_013_010_000	hds_00_013_010_000	calhead	34.829915 34.999874
	hds_00_015_016_000	hds_00_015_016_000	calhead	34.731595 34.566364
	hds_00_021_010_000	hds_00_021_010_000	calhead	34.745925 34.772247
	hds_00_022_015_000	hds_00_022_015_000	calhead	34.585874 34.603825
	hds_00_024_004_000	hds_00_024_004_000	calhead	35.036858 34.965450
	hds_00_026_006_000	hds_00_026_006_000	calhead	34.784656 34.759605
	hds_00_029_015_000	hds_00_029_015_000	calhead	34.524933 34.272018
	hds_00_033_007_000	hds_00_033_007_000	calhead	34.237387 34.105034
	hds_00_034_010_000	hds_00_034_010_000	calhead	33.792000 33.770840

	residual	weight
name		
fo_39_19791230	135.190144	0.013049
hds_00_002_009_000	0.069592	5.750067
hds_00_002_015_000	0.170213	5.750067
hds_00_003_008_000	0.389716	5.750067
hds_00_009_001_000	0.324789	5.750067
hds_00_013_010_000	-0.169959	5.750067
hds_00_015_016_000	0.165231	5.750067
hds_00_021_010_000	-0.026323	5.750067
hds_00_022_015_000	-0.017951	5.750067
hds_00_024_004_000	0.071408	5.750067
hds_00_026_006_000	0.025051	5.750067
hds_00_029_015_000	0.252914	5.750067
hds_00_033_007_000	0.132353	5.750067
hds_00_034_010_000	0.021161	5.750067

The residual should be exactly the noise values from above. Lets load the model (that was just run using the true pars) and check some things

```
In [24]: m = flopy.modflow.Modflow.load("freyberg.nam", model_ws=m_d)
```

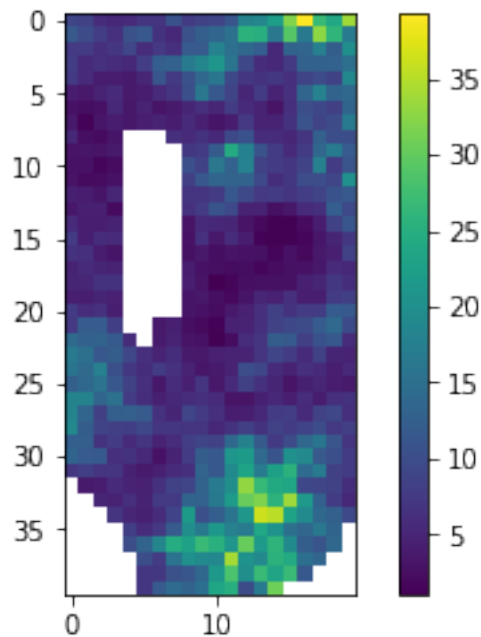
```
In [25]: a = m.upw.hk[0].array
        #a = m.rch.rech[0].array
```

```

a = np.ma.masked_where(m.bas6.ibound[0].array==0,a)
print(a.min(),a.max())
c = plt.imshow(a)
plt.colorbar()
plt.show()

```


1.017649 39.28789



```

In [26]: lst = flopy.utils.MfListBudget(os.path.join(m_d,"freyberg.list"))
df = lst.get_dataframes(diff=True)[0]
ax = df.plot(kind="bar",figsize=(10,10), grid=True)
a = ax.set_xticklabels(["historic","scenario"],rotation=90)
plt.show(ax)



```







1.0.8 see how our existing observation ensemble compares to the truth



forecasts:

```
In [27]: obs = pst.observation_data
plt.figure()
for forecast in fnames:
    ax = plt.subplot(111)
    obs_df.loc[:,forecast].hist(ax=ax,color="0.5",alpha=0.5)
    ax.plot([obs.loc[forecast,"obsval"],obs.loc[forecast,"obsval"]],ax.get_ylim(),"r")
    ax.set_title(forecast)
plt.show()
```









observations:


```
In [28]: for oname in pst.nnz_obs_names:
          ax = plt.subplot(111)
          obs_df.loc[:, oname].hist(ax=ax, color="0.5", alpha=0.5)
          ax.plot([obs.loc[oname, "obsval"], obs.loc[oname, "obsval"]], ax.get_ylim(), "r")
          ax.set_title(oname)
          plt.show()
```

