

prior_montecarlo

May 8, 2019

1 Run and process the prior monte carlo and pick a “truth” realization

A great advantage of exploring a synthetic model is that we can enforce a “truth” and then evaluate how our various attempts to estimate it perform. One way to do this is to run a monte carlo ensemble of multiple parameter realizations and then choose one of them to represent the “truth”. That will be accomplished in this notebook.

```
In [1]: import os
import shutil
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib as mpl
plt.rcParams['font.size']=12
import flopy
import pyemu
```

flopy is installed in /Users/jeremyw/Dev/gw1876/activities_2day_mfm/notebooks/flopy

1.0.1 set the t_d or “template directory” variable to point at the template folder and read in the PEST control file

```
In [2]: t_d = "template"
pst = pyemu.Pst(os.path.join(t_d, "freyberg.pst"))
```

1.0.2 Decide what pars are uncertain in the truth

We need to decide what our truth looks like - should the pilot points or the grid-scale pars be the source of spatial variability? or both?

```
In [3]: par = pst.parameter_data
# grid pars
#should_fix = par.loc[par.pargp.apply(lambda x: "gr" in x), "parname"]
# pp pars
#should_fix = par.loc[par.pargp.apply(lambda x: "pp" in x), "parname"]
#pst.npar - should_fix.shape[0]
```

```
In [4]: pe = pyemu.ParameterEnsemble.from_binary(pst=pst,filename=os.path.join(t_d,"prior.jcb")
        #pe.loc[:,should_fix] = 1.0
        pe.to_csv(os.path.join(t_d,"sweep_in.csv"))
```

new binary format detected...

1.0.3 run the prior ensemble in parallel locally

This takes advantage of the program pestpp-swp which runs a parameter sweep through a set of parameters. By default, pestpp-swp reads in the ensemble from a file called sweep_in.csv which in this case we made just above.

```
In [5]: m_d = "master_prior_sweep"
        pyemu.os_utils.start_slaves(t_d,"pestpp-swp","freyberg.pst",num_slaves=20,slave_root="
```

1.0.4 Load the output ensemble and plot a few things

```
In [6]: obs_df = pd.read_csv(os.path.join(m_d,"sweep_out.csv"),index_col=0)
        print('number of realization in the ensemble before dropping: ' + str(obs_df.shape[0]))
```

number of realization in the ensemble before dropping: 200

drop any failed runs

```
In [7]: obs_df = obs_df.loc[obs_df.failed_flag==0,:]
        print('number of realization in the ensemble **after** dropping: ' + str(obs_df.shape[
```

number of realization in the ensemble **after** dropping: 200


1.0.5 confirm which quantities were identified as forecasts



```
In [8]: fnames = pst.pestpp_options["forecasts"].split(',')
        fnames
```



```
Out[8]: ['fa_hw_19791230',
        'fa_hw_19801229',
        'fa_tw_19791230',
        'fa_tw_19801229',
        'hds_00_013_002_000',
        'hds_00_013_002_001',
        'part_time',
        'part_status']
```



1.0.6 now we can plot the distributions of each forecast


```
In [9]: for forecast in fnames:
        plt.figure()
        ax = obs_df.loc[:,forecast].plot(kind="hist")
        ax.set_title(forecast)
        plt.show()
```














We see that under scenario conditions, many more realizations for the flow to the aquifer in the headwaters are positive (as expected). Lets difference these two:

```
In [10]: sfnames = [f for f in fnames if "1980" in f or "_001" in f]
          hfnames = [f for f in fnames if "1979" in f or "_000" in f]
          diff = obs_df.loc[:,hfnames].values - obs_df.loc[:,sfnames].values
          diff = pd.DataFrame(diff,columns=sfnames)
          diff.hist(figsize=(10,10))
          plt.show()
```



We now see that the most extreme scenario yields a large decrease in flow from the aquifer to the headwaters (the most negative value)

1.0.7 setting the “truth”

We just need to replace the observed values (obsval) in the control file with the outputs for one of the realizations on obs_df. In this way, we now have the nonzero values for history matching, but also the truth values for comparing how we are doing with other unobserved quantities. I’m going to pick a realization that yields an “average” variability of the observed gw levels:

```
In [11]: # choose the realization with a low historic gw to sw headwater flux
hist_swgw = obs_df.loc[:, "fa_hw_19791230"].sort_values()
idx = hist_swgw.index[100]
idx
hist_swgw
```



```
Out[11]: run_id
72      -1439.898100
114     -1407.885900
135     -1400.619100
57      -1371.581000
100     -1330.061800
21      -1322.180700
145     -1310.682900
40      -1299.322400
93      -1297.668900
45      -1289.105600
137     -1282.997600
131     -1273.692600
33      -1270.382600
159     -1263.388500
69      -1234.232800
49      -1229.869700
23      -1222.724680
168     -1212.237700
119     -1202.302000
36      -1195.689330
151     -1183.319000
155     -1182.481800
43      -1176.382700
136     -1165.006500
133     -1161.195400
7       -1149.238900
17      -1141.957200
25      -1141.522700
105     -1138.499000
31      -1137.319600
...
63      -467.193200
75      -459.782621
89      -443.632510
59      -440.333500
27      -403.127710
38      -395.609500
61      -363.008265
161     -357.548200
116     -265.722100
103     -264.996300
76      -255.625800
180     -252.541616
52      -251.051210
173     -237.581829
92      -222.417300
108     -221.172100
```

```

190    -191.298431
30     -154.094510
111    -146.408040
71     -55.583140
98     -36.865800
56      49.053710
12     131.248200
51     251.162480
0       278.989000
165    387.888840
47     513.326680
113    781.076500
129    849.195100
120   1830.209000
Name: fa_hw_19791230, Length: 200, dtype: float64

```

```
In [12]: obs_df.loc[idx,pst.nnz_obs_names]
```

```

Out[12]: fo_39_19791230      12061.000000
hds_00_002_009_000        36.519390
hds_00_002_015_000        34.630413
hds_00_003_008_000        36.641518
hds_00_009_001_000        37.685490
hds_00_013_010_000        34.970249
hds_00_015_016_000        34.697254
hds_00_021_010_000        34.946480
hds_00_022_015_000        34.708725
hds_00_024_004_000        35.704624
hds_00_026_006_000        35.169151
hds_00_029_015_000        34.468609
hds_00_033_007_000        34.716385
hds_00_034_010_000        34.254543
Name: 8, dtype: float64

```

Lets see how our selected truth does with the sw/gw forecasts:

```
In [13]: obs_df.loc[idx,fnames]
```

```

Out[13]: fa_hw_19791230      -870.595760
fa_hw_19801229      -221.448077
fa_tw_19791230      -666.416960
fa_tw_19801229     -102.134800
hds_00_013_002_000      37.737610
hds_00_013_002_001      36.841492
part_time           1597.562000
part_status          2.000000
Name: 8, dtype: float64

```

Assign some initial weights. Now, it is custom to add noise to the observed values... we will use the classic Gaussian noise... zero mean and standard deviation of 1 over the weight

```
In [14]: pst = pyemu.Pst(os.path.join(t_d, "freyberg.pst"))
        obs = pst.observation_data
        obs.loc[:, "obsval"] = obs_df.loc[idx, pst.obs_names]
        obs.loc[obs.obgnme=="calhead", "weight"] = 10.0
        obs.loc[obs.obgnme=="calflux", "weight"] = 0.005
```

here we just get a sample from a random normal distribution with mean=0 and std=1. The argument indicates how many samples we want - and we choose `pst.nnz_obs` which is the number of nonzero-weighted observations in the PST file

```
In [15]: np.random.seed(seed=0)
        snd = np.random.randn(pst.nnz_obs)
        noise = snd * 1./obs.loc[pst.nnz_obs_names, "weight"]
        pst.observation_data.loc[noise.index, "obsval"] += noise
        noise
```

```
Out[15]: obsnme
fo_39_19791230      352.810469
hds_00_002_009_000    0.040016
hds_00_002_015_000    0.097874
hds_00_003_008_000    0.224089
hds_00_009_001_000    0.186756
hds_00_013_010_000   -0.097728
hds_00_015_016_000    0.095009
hds_00_021_010_000   -0.015136
hds_00_022_015_000   -0.010322
hds_00_024_004_000    0.041060
hds_00_026_006_000    0.014404
hds_00_029_015_000    0.145427
hds_00_033_007_000    0.076104
hds_00_034_010_000    0.012168
Name: weight, dtype: float64
```

Then we write this out to a new file and run `pestpp-ies` to see how the objective function looks

```
In [16]: pst.write(os.path.join(t_d, "freyberg.pst"))
        pyemu.os_utils.run("pestpp-ies freyberg.pst", cwd=t_d)

noptmax:0, npar_adj:14819, nnz_obs:14
```

Now we can read in the results and make some figures showing residuals and the balance of the objective function

```
In [17]: pst = pyemu.Pst(os.path.join(t_d, "freyberg.pst"))
        print(pst.phi)
        plt.figure()
        pst.plot(kind='phi_pie')
```


```
print('Here are the non-zero weighted observation names')

figs = pst.plot(kind="1to1")
plt.show()
pst.res.loc[pst.nnz_obs_names,:]
```





603.8982892954039

Here are the non-zero weighted observation names

<Figure size 432x288 with 0 Axes>



<Figure size 576x756 with 0 Axes>



```
Out[17]:
```

	name	group	measured	modelled \
	name			
	fo_39_19791230	fo_39_19791230 calflux	12413.810469	11430.000000

hds_00_002_009_000	hds_00_002_009_000	calhead	36.559406	37.107498
hds_00_002_015_000	hds_00_002_015_000	calhead	34.728287	35.045185
hds_00_003_008_000	hds_00_003_008_000	calhead	36.865607	37.397289
hds_00_009_001_000	hds_00_009_001_000	calhead	37.872245	39.546417
hds_00_013_010_000	hds_00_013_010_000	calhead	34.872521	35.571774
hds_00_015_016_000	hds_00_015_016_000	calhead	34.792263	34.835716
hds_00_021_010_000	hds_00_021_010_000	calhead	34.931344	35.386250
hds_00_022_015_000	hds_00_022_015_000	calhead	34.698403	34.577492
hds_00_024_004_000	hds_00_024_004_000	calhead	35.745684	36.760464
hds_00_026_006_000	hds_00_026_006_000	calhead	35.183556	35.896149
hds_00_029_015_000	hds_00_029_015_000	calhead	34.614036	34.453842
hds_00_033_007_000	hds_00_033_007_000	calhead	34.792489	34.678810
hds_00_034_010_000	hds_00_034_010_000	calhead	34.266711	34.118073

	residual	weight
name		
fo_39_19791230	983.810469	0.005
hds_00_002_009_000	-0.548092	10.000
hds_00_002_015_000	-0.316898	10.000
hds_00_003_008_000	-0.531682	10.000
hds_00_009_001_000	-1.674172	10.000
hds_00_013_010_000	-0.699252	10.000
hds_00_015_016_000	-0.043453	10.000
hds_00_021_010_000	-0.454905	10.000
hds_00_022_015_000	0.120911	10.000
hds_00_024_004_000	-1.014780	10.000
hds_00_026_006_000	-0.712593	10.000
hds_00_029_015_000	0.160194	10.000
hds_00_033_007_000	0.113679	10.000
hds_00_034_010_000	0.148638	10.000

Publication ready figs - oh snap!

Depending on the truth you chose, we may have a problem - we set the weights for both the heads and the flux to reasonable values based on what we expect for measurement noise. But the contributions to total phi might be out of balance - if contribution of the flux measurement to total phi is too low, the history matching excersizes (coming soon!) will focus almost entirely on minimizing head residuals. So we need to balance the objective function. This is a subtle but very important step, especially since some of our forecasts deal with sw-gw exchange

```
In [18]: #pc = pst.phi_components
#target = {"calflux":0.3 * pc["calhead"]}
#pst.adjust_weights(obsgrp_dict=target)
#pst.plot(kind='phi_pie')
```

Just to make sure we have everything working right, we should be able to load the truth parameters, run the model once and have a phi equivalent to the noise vector:

```
In [19]: par_df = pd.read_csv(os.path.join(m_d,"sweep_in.csv"),index_col=0)
pst.parameter_data.loc[:,"parval1"] = par_df.loc[idx,pst.par_names]
pst.write(os.path.join(m_d,"test.pst"))
```

```
noptmax:0, npar_adj:14819, nnz_obs:14
```

we will run this with noptmax=0 to preform a single run. Pro-tip: you can use any of the pestpp-### binaries/executables to run noptmax=0

```
In [20]: pyemu.os_utils.run("pestpp-ies.exe test.pst", cwd=m_d)
         pst = pyemu.Pst(os.path.join(m_d, "test.pst"))
         print(pst.phi)
         pst.res.loc[pst.nnz_obs_names, :]
```

```
17.52884723870795
```

```
Out [20]:
```

	name	group	measured	modelled \
name				
fo_39_19791230	fo_39_19791230	calflux	12413.810469	12061.000000
hds_00_002_009_000	hds_00_002_009_000	calhead	36.559406	36.519390
hds_00_002_015_000	hds_00_002_015_000	calhead	34.728287	34.630413
hds_00_003_008_000	hds_00_003_008_000	calhead	36.865607	36.641518
hds_00_009_001_000	hds_00_009_001_000	calhead	37.872245	37.685490
hds_00_013_010_000	hds_00_013_010_000	calhead	34.872521	34.970249
hds_00_015_016_000	hds_00_015_016_000	calhead	34.792263	34.697254
hds_00_021_010_000	hds_00_021_010_000	calhead	34.931344	34.946480
hds_00_022_015_000	hds_00_022_015_000	calhead	34.698403	34.708725
hds_00_024_004_000	hds_00_024_004_000	calhead	35.745684	35.704624
hds_00_026_006_000	hds_00_026_006_000	calhead	35.183556	35.169151
hds_00_029_015_000	hds_00_029_015_000	calhead	34.614036	34.468609
hds_00_033_007_000	hds_00_033_007_000	calhead	34.792489	34.716385
hds_00_034_010_000	hds_00_034_010_000	calhead	34.266711	34.254543


	residual	weight
name		
fo_39_19791230	352.810469	0.005
hds_00_002_009_000	0.040016	10.000
hds_00_002_015_000	0.097874	10.000
hds_00_003_008_000	0.224089	10.000
hds_00_009_001_000	0.186756	10.000
hds_00_013_010_000	-0.097728	10.000
hds_00_015_016_000	0.095009	10.000
hds_00_021_010_000	-0.015136	10.000
hds_00_022_015_000	-0.010322	10.000
hds_00_024_004_000	0.041060	10.000
hds_00_026_006_000	0.014404	10.000
hds_00_029_015_000	0.145427	10.000
hds_00_033_007_000	0.076104	10.000
hds_00_034_010_000	0.012168	10.000

The residual should be exactly the noise values from above. Lets load the model (that was just run using the true pars) and check some things


```
In [21]: m = flopy.modflow.Modflow.load("freyberg.nam",model_ws=m_d)
```

```
In [22]: a = m.upw.vka[1].array
        #a = m.rch.rech[0].array
        a = np.ma.masked_where(m.bas6.ibound[0].array==0,a)
        print(a.min(),a.max())
        c = plt.imshow(a)
        plt.colorbar()
        plt.show()
```

0.03417815 2.730791





```
In [23]: lst = flopy.utils.MfListBudget(os.path.join(m_d,"freyberg.list"))
        df = lst.get_dataframes(diff=True)[0]
        ax = df.plot(kind="bar",figsize=(6,6), grid=True)
        a = ax.set_xticklabels(["historic","scenario"],rotation=90)
        plt.show(ax)
```








1.0.8 see how our existing observation ensemble compares to the truth



forecasts:

```
In [24]: obs = pst.observation_data
plt.figure()
for forecast in fnames:
    ax = plt.subplot(111)
    obs_df.loc[:,forecast].hist(ax=ax,color="0.5",alpha=0.5)
    ax.plot([obs.loc[forecast,"obsval"],obs.loc[forecast,"obsval"]],ax.get_ylim(),"r")
    ax.set_title(forecast)
plt.show()
```









observations:

```

In [25]: for oname in pst.nnz_obs_names:
          ax = plt.subplot(111)
          obs_df.loc[:, oname].hist(ax=ax, color="0.5", alpha=0.5)
          ax.plot([obs.loc[oname, "obsval"], obs.loc[oname, "obsval"]], ax.get_ylim(), "r")
          ax.set_title(oname)
          plt.show()

```

