

prior_montecarlo

May 10, 2019

1 Run and process the prior monte carlo and pick a “truth” realization

A great advantage of exploring a synthetic model is that we can enforce a “truth” and then evaluate how our various attempts to estimate it perform. One way to do this is to run a monte carlo ensemble of multiple parameter realizations and then choose one of them to represent the “truth”. That will be accomplished in this notebook.

```
In [1]: import os
import shutil
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib as mpl
plt.rcParams['font.size']=12
import flopy
import pyemu
```

flopy is installed in /Users/jeremyw/Dev/gw1876/activities_2day_mfm/notebooks/flopy

1.0.1 set the t_d or “template directory” variable to point at the template folder and read in the PEST control file

```
In [2]: t_d = "template"
pst = pyemu.Pst(os.path.join(t_d, "freyberg.pst"))
```

1.0.2 Decide what pars are uncertain in the truth

We need to decide what our truth looks like - should the pilot points or the grid-scale pars be the source of spatial variability? or both?

```
In [3]: par = pst.parameter_data
# grid pars
#should_fix = par.loc[par.pargp.apply(lambda x: "gr" in x), "parname"]
# pp pars
#should_fix = par.loc[par.pargp.apply(lambda x: "pp" in x), "parname"]
#pst.npar - should_fix.shape[0]
```

```
In [4]: pe = pyemu.ParameterEnsemble.from_binary(pst=pst,filename=os.path.join(t_d,"prior.jcb")
        #pe.loc[:,should_fix] = 1.0
        pe.to_csv(os.path.join(t_d,"sweep_in.csv"))
```

new binary format detected...

1.0.3 run the prior ensemble in parallel locally

This takes advantage of the program pestpp-swp which runs a parameter sweep through a set of parameters. By default, pestpp-swp reads in the ensemble from a file called sweep_in.csv which in this case we made just above.

```
In [5]: m_d = "master_prior_sweep"
        pyemu.os_utils.start_slaves(t_d,"pestpp-swp","freyberg.pst",num_slaves=20,slave_root="
```

1.0.4 Load the output ensemble and plot a few things

```
In [6]: obs_df = pd.read_csv(os.path.join(m_d,"sweep_out.csv"),index_col=0)
        print('number of realization in the ensemble before dropping: ' + str(obs_df.shape[0]))
```

number of realization in the ensemble before dropping: 200

drop any failed runs

```
In [7]: obs_df = obs_df.loc[obs_df.failed_flag==0,:]
        print('number of realization in the ensemble **after** dropping: ' + str(obs_df.shape[
```

number of realization in the ensemble **after** dropping: 200


1.0.5 confirm which quantities were identified as forecasts



```
In [8]: fnames = pst.pestpp_options["forecasts"].split(',')
        fnames
```



```
Out[8]: ['fa_hw_19791230',
        'fa_hw_19801229',
        'fa_tw_19791230',
        'fa_tw_19801229',
        'hds_00_013_002_000',
        'hds_00_013_002_001',
        'part_time',
        'part_status']
```



1.0.6 now we can plot the distributions of each forecast


```
In [9]: for forecast in fnames:
        plt.figure()
        ax = obs_df.loc[:,forecast].plot(kind="hist")
        ax.set_title(forecast)
        plt.show()
```














We see that under scenario conditions, many more realizations for the flow to the aquifer in the headwaters are positive (as expected). Lets difference these two:

```
In [10]: sfnames = [f for f in fnames if "1980" in f or "_001" in f]
          hfnames = [f for f in fnames if "1979" in f or "_000" in f]
          diff = obs_df.loc[:,hfnames].values - obs_df.loc[:,sfnames].values
          diff = pd.DataFrame(diff,columns=sfnames)
          diff.hist(figsize=(10,10))
          plt.show()
```



We now see that the most extreme scenario yields a large decrease in flow from the aquifer to the headwaters (the most negative value)

1.0.7 setting the “truth”

We just need to replace the observed values (obsval) in the control file with the outputs for one of the realizations on obs_df. In this way, we now have the nonzero values for history matching, but also the truth values for comparing how we are doing with other unobserved quantities. I’m going to pick a realization that yields an “average” variability of the observed gw levels:

```
In [11]: # choose the realization with a low historic gw to sw headwater flux
#hist_swgw = obs_df.loc[:, "fa_hw_19791230"].sort_values()
hist_swgw = obs_df.loc[:, "part_time"].sort_values()
idx = hist_swgw.index[20]
```



```
idx
hist_swgw
```

```
Out[11]: run_id
149      207.9040
165      237.0716
191      280.1461
97       286.6308
77       289.1803
16       304.0466
173      326.0275
186      332.5973
180      355.7679
145      357.3499
78       359.7271
107      378.5330
17       391.6625
156      399.4463
76       410.1112
0        415.2489
96       437.2630
46       439.3421
2        442.4801
35       442.4812
64       452.2543
128      453.5696
141      464.6295
93       469.9462
21       480.9774
117      482.3959
6        489.8026
73       492.7455
11       501.1917
62       517.6446
...
8        4015.0000
118      4015.0000
190      4015.0000
69       4015.0000
193      4015.0000
194      4015.0000
12       4015.0000
47       4015.0000
75       4015.0000
27       4015.0000
82       4015.0000
104      4015.0000
153      4015.0000
```

```

138    4015.0000
136    4015.0000
157    4015.0000
159    4015.0000
41     4015.0000
125    4015.0000
74     4015.0000
84     4015.0000
166    4015.0000
29     4015.0000
169    4015.0000
171    4015.0000
172    4015.0000
144    4015.0000
175    4015.0000
131    4015.0000
42     4015.0000
Name: part_time, Length: 200, dtype: float64

```

```
In [12]: obs_df.loc[idx,pst.nnz_obs_names]
```

```

Out[12]: fo_39_19791230      10555.000000
hds_00_002_009_000         34.847549
hds_00_002_015_000         34.639633
hds_00_003_008_000         34.861164
hds_00_009_001_000         35.161556
hds_00_013_010_000         34.661835
hds_00_015_016_000         34.569942
hds_00_021_010_000         34.571648
hds_00_022_015_000         34.439877
hds_00_024_004_000         34.609158
hds_00_026_006_000         34.521420
hds_00_029_015_000         34.297997
hds_00_033_007_000         34.187607
hds_00_034_010_000         33.842358
Name: 64, dtype: float64

```

Lets see how our selected truth does with the sw/gw forecasts:

```
In [13]: obs_df.loc[idx,fnames]
```

```

Out[13]: fa_hw_19791230      -992.855500
fa_hw_19801229       -73.8777930
fa_tw_19791230      -256.953300
fa_tw_19801229      368.869266
hds_00_013_002_000      35.032547
hds_00_013_002_001      34.416519
part_time            452.254300
part_status           2.000000
Name: 64, dtype: float64

```

Assign some initial weights. Now, it is custom to add noise to the observed values... we will use the classic Gaussian noise... zero mean and standard deviation of 1 over the weight

```
In [14]: pst = pyemu.Pst(os.path.join(t_d,"freyberg.pst"))
obs = pst.observation_data
obs.loc[:, "obsval"] = obs_df.loc[idx, pst.obs_names]
obs.loc[obs.obgnme=="calhead", "weight"] = 10.0
obs.loc[obs.obgnme=="calflux", "weight"] = 1.0
```

here we just get a sample from a random normal distribution with mean=0 and std=1. The argument indicates how many samples we want - and we choose `pst.nnz_obs` which is the the number of nonzero-weighted observations in the PST file

```
In [15]: np.random.seed(seed=0)
snd = np.random.randn(pst.nnz_obs)
noise = snd * 1./obs.loc[pst.nnz_obs_names, "weight"]
pst.observation_data.loc[noise.index, "obsval"] += noise
noise
```

```
Out[15]: obsnme
fo_39_19791230      1.764052
hds_00_002_009_000  0.040016
hds_00_002_015_000  0.097874
hds_00_003_008_000  0.224089
hds_00_009_001_000  0.186756
hds_00_013_010_000 -0.097728
hds_00_015_016_000  0.095009
hds_00_021_010_000 -0.015136
hds_00_022_015_000 -0.010322
hds_00_024_004_000  0.041060
hds_00_026_006_000  0.014404
hds_00_029_015_000  0.145427
hds_00_033_007_000  0.076104
hds_00_034_010_000  0.012168
Name: weight, dtype: float64
```

Then we write this out to a new file and run `pestpp-ies` to see how the objective function looks

```
In [16]: pst.write(os.path.join(t_d,"freyberg.pst"))
pyemu.os_utils.run("pestpp-ies freyberg.pst", cwd=t_d)
```

```
noptmax:0, npar_adj:14819, nnz_obs:14
```

Now we can read in the results and make some figures showing residuals and the balance of the objective function

```

In [17]: pst = pyemu.Pst(os.path.join(t_d,"freyberg.pst"))
          print(pst.phi)
          plt.figure()
          pst.plot(kind='phi_pie')
          print('Here are the non-zero weighted observation names')


          figs = pst.plot(kind="1to1")
          plt.show()
          pst.res.loc[pst.nnz_obs_names,:]

```


766170.2973169514

Here are the non-zero weighted observation names

<Figure size 432x288 with 0 Axes>



<Figure size 576x756 with 0 Axes>



Out[17]:

name	group	measured	modelled \
------	-------	----------	------------

name				
fo_39_19791230	fo_39_19791230	calflux	10556.764052	11430.000000
hds_00_002_009_000	hds_00_002_009_000	calhead	34.887565	37.107498
hds_00_002_015_000	hds_00_002_015_000	calhead	34.737507	35.045185
hds_00_003_008_000	hds_00_003_008_000	calhead	35.085253	37.397289
hds_00_009_001_000	hds_00_009_001_000	calhead	35.348312	39.546417
hds_00_013_010_000	hds_00_013_010_000	calhead	34.564107	35.571774
hds_00_015_016_000	hds_00_015_016_000	calhead	34.664951	34.835716
hds_00_021_010_000	hds_00_021_010_000	calhead	34.556512	35.386250
hds_00_022_015_000	hds_00_022_015_000	calhead	34.429555	34.577492
hds_00_024_004_000	hds_00_024_004_000	calhead	34.650217	36.760464
hds_00_026_006_000	hds_00_026_006_000	calhead	34.535824	35.896149
hds_00_029_015_000	hds_00_029_015_000	calhead	34.443424	34.453842
hds_00_033_007_000	hds_00_033_007_000	calhead	34.263711	34.678810
hds_00_034_010_000	hds_00_034_010_000	calhead	33.854525	34.118073


	residual	weight
name		
fo_39_19791230	-873.235948	1.0
hds_00_002_009_000	-2.219933	10.0
hds_00_002_015_000	-0.307678	10.0
hds_00_003_008_000	-2.312036	10.0
hds_00_009_001_000	-4.198105	10.0
hds_00_013_010_000	-1.007667	10.0
hds_00_015_016_000	-0.170765	10.0
hds_00_021_010_000	-0.829738	10.0
hds_00_022_015_000	-0.147937	10.0
hds_00_024_004_000	-2.110246	10.0
hds_00_026_006_000	-1.360325	10.0
hds_00_029_015_000	-0.010418	10.0
hds_00_033_007_000	-0.415100	10.0
hds_00_034_010_000	-0.263547	10.0

Publication ready figs - oh snap!

Depending on the truth you chose, we may have a problem - we set the weights for both the heads and the flux to reasonable values based on what we expect for measurement noise. But the contributions to total phi might be out of balance - if contribution of the flux measurement to total phi is too low, the history matching excersizes (coming soon!) will focus almost entirely on minimizing head residuals. So we need to balance the objective function. This is a subtle but very important step, especially since some of our forecasts deal with sw-gw exchange

```
In [18]: pc = pst.phi_components
          #target = {"calflux":0.3 * pc["calhead"]}
          target = {"calhead":500,"calflux":500}
          pst.adjust_weights(obsgrp_dict=target)
          pst.plot(kind='phi_pie')
```

```
Out[18]: <matplotlib.axes._subplots.AxesSubplot at 0x10ec8f2b0>
```



Lets see what the new flux observation weight is:

```
In [19]: pst.observation_data.loc[pst.nnz_obs_names, "weight"]
```

```
Out [19]: obsnme  
fo_39_19791230      0.025607  
hds_00_002_009_000  3.711718  
hds_00_002_015_000  3.711718  
hds_00_003_008_000  3.711718  
hds_00_009_001_000  3.711718  
hds_00_013_010_000  3.711718  
hds_00_015_016_000  3.711718  
hds_00_021_010_000  3.711718  
hds_00_022_015_000  3.711718
```

```

hds_00_024_004_000    3.711718
hds_00_026_006_000    3.711718
hds_00_029_015_000    3.711718
hds_00_033_007_000    3.711718
hds_00_034_010_000    3.711718
Name: weight, dtype: float64

```

Now, for some super trickery: since we changed the weight, we need to generate the observation noise using these new weights for the error model (so meta!)

```

In [20]: obs = pst.observation_data
         np.random.seed(seed=0)
         snd = np.random.randn(pst.nnz_obs)
         noise = snd * 1./obs.loc[pst.nnz_obs_names,"weight"]
         obs.loc[:, "obsval"] = obs_df.loc[idx,pst.obs_names]
         pst.observation_data.loc[noise.index,"obsval"] += noise
         noise

```

```

Out[20]: obsnme
fo_39_19791230        68.890299
hds_00_002_009_000     0.107809
hds_00_002_015_000     0.263689
hds_00_003_008_000     0.603735
hds_00_009_001_000     0.503152
hds_00_013_010_000    -0.263295
hds_00_015_016_000     0.255970
hds_00_021_010_000    -0.040778
hds_00_022_015_000    -0.027809
hds_00_024_004_000     0.110622
hds_00_026_006_000     0.038808
hds_00_029_015_000     0.391806
hds_00_033_007_000     0.205037
hds_00_034_010_000     0.032781
Name: weight, dtype: float64

```

```


In [21]: pst.write(os.path.join(t_d,"freyberg.pst"))
         pyemu.os_utils.run("pestpp-ies freyberg.pst",cwd=t_d)
         pst = pyemu.Pst(os.path.join(t_d,"freyberg.pst"))
         print(pst.phi)
         pst.plot(kind='phi_pie')
         plt.show()

```

```

noptmax:0, npar_adj:14819, nnz_obs:14
863.2909159111605

```

Whew! confused yet? Ok, let's leave all this confusion behind...its mostly academic, just to make sure we are using weights that are in harmony with the noise we added to the truth...Just to make sure we have everything working right, we should be able to load the truth parameters, run the model once and have a phi equivalent to the noise vector:

```
In [22]: par_df = pd.read_csv(os.path.join(m_d, "sweep_in.csv"), index_col=0)
         pst.parameter_data.loc[:, "parval1"] = par_df.loc[idx, pst.par_names]
         pst.write(os.path.join(m_d, "test.pst"))
```

```
noptmax:0, npar_adj:14819, nnz_obs:14
```

we will run this with noptmax=0 to preform a single run. Pro-tip: you can use any of the pestpp-### binaries/executables to run noptmax=0

```
In [23]: pyemu.os_utils.run("pestpp-ies.exe test.pst", cwd=m_d)
        pst = pyemu.Pst(os.path.join(m_d, "test.pst"))
        print(pst.phi)
        pst.res.loc[pst.nnz_obs_names,:]
```

17.528847196782618

```
Out [23]:
```

	name	group	measured	modelled \
	name			
	fo_39_19791230	fo_39_19791230	calflux	10623.890299 10555.000000
	hds_00_002_009_000	hds_00_002_009_000	calhead	34.955359 34.847549
	hds_00_002_015_000	hds_00_002_015_000	calhead	34.903322 34.639633
	hds_00_003_008_000	hds_00_003_008_000	calhead	35.464899 34.861164
	hds_00_009_001_000	hds_00_009_001_000	calhead	35.664708 35.161556
	hds_00_013_010_000	hds_00_013_010_000	calhead	34.398539 34.661835
	hds_00_015_016_000	hds_00_015_016_000	calhead	34.825912 34.569942
	hds_00_021_010_000	hds_00_021_010_000	calhead	34.530869 34.571648
	hds_00_022_015_000	hds_00_022_015_000	calhead	34.412068 34.439877
	hds_00_024_004_000	hds_00_024_004_000	calhead	34.719780 34.609158
	hds_00_026_006_000	hds_00_026_006_000	calhead	34.560227 34.521420
	hds_00_029_015_000	hds_00_029_015_000	calhead	34.689803 34.297997
	hds_00_033_007_000	hds_00_033_007_000	calhead	34.392643 34.187607
	hds_00_034_010_000	hds_00_034_010_000	calhead	33.875139 33.842358

	residual	weight
name		
fo_39_19791230	68.890299	0.025607
hds_00_002_009_000	0.107809	3.711718
hds_00_002_015_000	0.263689	3.711718
hds_00_003_008_000	0.603735	3.711718
hds_00_009_001_000	0.503152	3.711718
hds_00_013_010_000	-0.263295	3.711718
hds_00_015_016_000	0.255970	3.711718
hds_00_021_010_000	-0.040778	3.711718
hds_00_022_015_000	-0.027809	3.711718
hds_00_024_004_000	0.110622	3.711718
hds_00_026_006_000	0.038808	3.711718
hds_00_029_015_000	0.391806	3.711718
hds_00_033_007_000	0.205037	3.711718
hds_00_034_010_000	0.032781	3.711718

The residual should be exactly the noise values from above. Lets load the model (that was just run using the true pars) and check some things

```
In [24]: m = flopy.modflow.Modflow.load("freyberg.nam", model_ws=m_d)
```


```
In [25]: a = m.upw.hk[0].array
        #a = m.rch.rech[0].array
```

```

a = np.ma.masked_where(m.bas6.ibound[0].array==0,a)
print(a.min(),a.max())
c = plt.imshow(a)
plt.colorbar()
plt.show()

```


0.1996812 12.40834



```

In [26]: lst = flopy.utils.MfListBudget(os.path.join(m_d,"freyberg.list"))
df = lst.get_dataframes(diff=True)[0]
ax = df.plot(kind="bar",figsize=(10,10), grid=True)
a = ax.set_xticklabels(["historic","scenario"],rotation=90)
plt.show(ax)



```







1.0.8 see how our existing observation ensemble compares to the truth



forecasts:

```
In [27]: obs = pst.observation_data
plt.figure()
for forecast in fnames:
    ax = plt.subplot(111)
    obs_df.loc[:,forecast].hist(ax=ax,color="0.5",alpha=0.5)
    ax.plot([obs.loc[forecast,"obsval"],obs.loc[forecast,"obsval"]],ax.get_ylim(),"r")
    ax.set_title(forecast)
plt.show()
```









observations:


```
In [28]: for oname in pst.nnz_obs_names:
          ax = plt.subplot(111)
          obs_df.loc[:, oname].hist(ax=ax, color="0.5", alpha=0.5)
          ax.plot([obs.loc[oname, "obsval"], obs.loc[oname, "obsval"]], ax.get_ylim(), "r")
          ax.set_title(oname)
          plt.show()
```

