

prior_montecarlo

May 7, 2019

1 Run and process the prior monte carlo and pick a “truth” realization

A great advantage of exploring a synthetic model is that we can enforce a “truth” and then evaluate how our various attempts to estimate it perform. One way to do this is to run a monte carlo ensemble of multiple parameter realizations and then choose one of them to represent the “truth”. That will be accomplished in this notebook.

```
In [1]: import os
import shutil
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib as mpl
plt.rcParams['font.size']=12
import flopy
import pyemu
```

flopy is installed in /Users/jeremyw/Dev/gw1876/activities_2day_mfm/notebooks/flopy

1.0.1 set the t_d or “template directory” variable to point at the template folder and read in the PEST control file

```
In [2]: t_d = "template"
pst = pyemu.Pst(os.path.join(t_d, "freyberg.pst"))
```

1.0.2 Decide what pars are uncertain in the truth

We need to decide what our truth looks like - should the pilot points or the grid-scale pars be the source of spatial variability? or both?

```
In [3]: par = pst.parameter_data
# grid pars
#should_fix = par.loc[par.pargp.apply(lambda x: "gr" in x), "parname"]
# pp pars
#should_fix = par.loc[par.pargp.apply(lambda x: "pp" in x), "parname"]
#pst.npar - should_fix.shape[0]
```

```
In [4]: pe = pyemu.ParameterEnsemble.from_binary(pst=pst,filename=os.path.join(t_d,"prior.jcb")
        #pe.loc[:,should_fix] = 1.0
        pe.to_csv(os.path.join(t_d,"sweep_in.csv"))
```

new binary format detected...

1.0.3 run the prior ensemble in parallel locally

This takes advantage of the program pestpp-swp which runs a parameter sweep through a set of parameters. By default, pestpp-swp reads in the ensemble from a file called sweep_in.csv which in this case we made just above.

```
In [5]: m_d = "master_prior_sweep"
        pyemu.os_utils.start_slaves(t_d,"pestpp-swp","freyberg.pst",num_slaves=20,slave_root="
```

1.0.4 Load the output ensemble and plot a few things

```
In [6]: obs_df = pd.read_csv(os.path.join(m_d,"sweep_out.csv"),index_col=0)
        print('number of realization in the ensemble before dropping: ' + str(obs_df.shape[0]))
```

number of realization in the ensemble before dropping: 200

drop any failed runs

```
In [7]: obs_df = obs_df.loc[obs_df.failed_flag==0,:]
        print('number of realization in the ensemble **after** dropping: ' + str(obs_df.shape[0]))
```

number of realization in the ensemble **after** dropping: 200


1.0.5 confirm which quantities were identified as forecasts



```
In [8]: fnames = pst.pestpp_options["forecasts"].split(',')
        fnames
```



```
Out[8]: ['fa_hw_19791230',
        'fa_hw_19801229',
        'fa_tw_19791230',
        'fa_tw_19801229',
        'hds_00_013_002_000',
        'hds_00_013_002_001',
        'part_time',
        'part_status']
```



1.0.6 now we can plot the distributions of each forecast

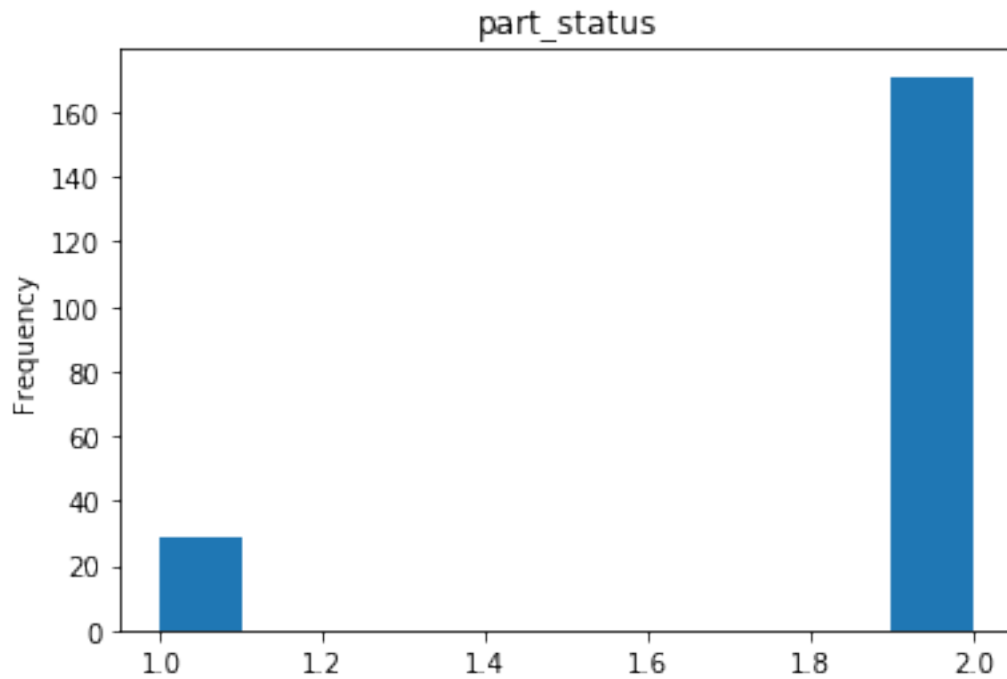
```
In [9]: for forecast in fnames:
        plt.figure()
        ax = obs_df.loc[:,forecast].plot(kind="hist")
        ax.set_title(forecast)
        plt.show()
```






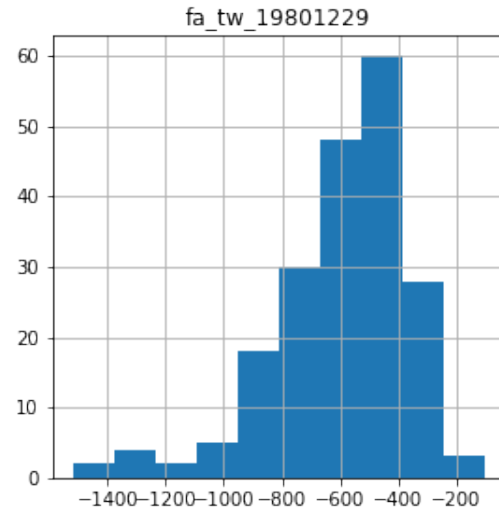







We see that under scenario conditions, many more realizations for the flow to the aquifer in the headwaters are positive (as expected). Lets difference these two:

```
In [10]: sfnames = [f for f in fnames if "1980" in f or "_001" in f]
          hfnames = [f for f in fnames if "1979" in f or "_000" in f]
          diff = obs_df.loc[:,hfnames].values - obs_df.loc[:,sfnames].values
          diff = pd.DataFrame(diff,columns=sfnames)
          diff.hist(figsize=(10,10))
          plt.show()
```



We now see that the most extreme scenario yields a large decrease in flow from the aquifer to the headwaters (the most negative value)

1.0.7 setting the “truth”

We just need to replace the observed values (obsval) in the control file with the outputs for one of the realizations on obs_df. In this way, we now have the nonzero values for history matching, but also the truth values for comparing how we are doing with other unobserved quantities. I’m going to pick a realization that yields an “average” variability of the observed gw levels:

```
In [11]: # choose the realization with a low historic gw to sw headwater flux
hist_sgw = obs_df.loc[:, "fa_hw_19791230"].sort_values()
idx = hist_sgw.index[100]
idx
hist_sgw
```



```
Out[11]: run_id
21      -1515.022600
5        -1478.211200
182     -1429.960200
20      -1370.134500
11      -1337.385700
116     -1308.621900
114     -1300.427400
38      -1292.823000
178     -1290.769200
61      -1268.275270
138     -1266.314400
49      -1250.714800
130     -1242.871000
170     -1242.660510
186     -1239.799900
26      -1239.114500
44      -1227.066600
65      -1221.035000
22      -1218.391100
48      -1216.842800
128     -1211.274800
113     -1209.880300
185     -1199.719800
2        -1184.331600
146     -1166.464300
46      -1165.863000
159     -1156.649900
134     -1155.403000
166     -1148.670800
69      -1148.229100
...
30      -307.996060
98      -302.095600
162     -294.096350
89      -249.515060
71      -208.660760
147     -205.904900
51      -199.145762
165     -166.591800
111     -165.552300
189     -163.077070
149     -119.544090
181     -109.619790
56      -83.603250
105     -74.408900
50      -57.105100
36      -25.374384
```

```

126      -13.427900
31       -12.679864
155        1.307900
40        11.568782
17        23.331010
60        76.068821
0         84.979200
12       119.832260
139      139.525690
163      160.345500
90       164.420400
41       186.825430
16       425.030270
27       484.226980
Name: fa_hw_19791230, Length: 200, dtype: float64

```

```
In [12]: obs_df.loc[idx,pst.nnz_obs_names]
```

```

Out[12]: fo_39_19791230      10597.000000
hds_00_002_009_000        37.343391
hds_00_002_015_000        34.760021
hds_00_003_008_000        37.572289
hds_00_009_001_000        39.558464
hds_00_013_010_000        35.365974
hds_00_015_016_000        34.653934
hds_00_021_010_000        34.711166
hds_00_022_015_000        34.419338
hds_00_024_004_000        35.141888
hds_00_026_006_000        34.682674
hds_00_029_015_000        34.207253
hds_00_033_007_000        34.075302
hds_00_034_010_000        33.636864
Name: 81, dtype: float64

```

Lets see how our selected truth does with the sw/gw forecasts:

```
In [13]: obs_df.loc[idx,fnames]
```

```

Out[13]: fa_hw_19791230      -853.092900
fa_hw_19801229      -210.374210
fa_tw_19791230      -184.237020
fa_tw_19801229       427.883580
hds_00_013_002_000       39.466167
hds_00_013_002_001       38.665234
part_time           637.733900
part_status          2.000000
Name: 81, dtype: float64

```

Assign some initial weights. Now, it is custom to add noise to the observed values... we will use the classic Gaussian noise... zero mean and standard deviation of 1 over the weight

```
In [14]: pst = pyemu.Pst(os.path.join(t_d,"freyberg.pst"))
        obs = pst.observation_data
        obs.loc[:, "obsval"] = obs_df.loc[idx, pst.obs_names]
        obs.loc[obs.obgnme=="calhead", "weight"] = 10.0
        obs.loc[obs.obgnme=="calflux", "weight"] = 0.005
```

here we just get a sample from a random normal distribution with mean=0 and std=1. The argument indicates how many samples we want - and we choose `pst.nnz_obs` which is the number of nonzero-weighted observations in the PST file

```
In [15]: np.random.seed(seed=0)
        snd = np.random.randn(pst.nnz_obs)
        noise = snd * 1./obs.loc[pst.nnz_obs_names, "weight"]
        pst.observation_data.loc[noise.index, "obsval"] += noise
        noise
```

```
Out[15]: obsnme
fo_39_19791230      352.810469
hds_00_002_009_000    0.040016
hds_00_002_015_000    0.097874
hds_00_003_008_000    0.224089
hds_00_009_001_000    0.186756
hds_00_013_010_000   -0.097728
hds_00_015_016_000    0.095009
hds_00_021_010_000   -0.015136
hds_00_022_015_000   -0.010322
hds_00_024_004_000    0.041060
hds_00_026_006_000    0.014404
hds_00_029_015_000    0.145427
hds_00_033_007_000    0.076104
hds_00_034_010_000    0.012168
Name: weight, dtype: float64
```

Then we write this out to a new file and run `pestpp-ies` to see how the objective function looks

```
In [16]: pst.write(os.path.join(t_d,"freyberg.pst"))
        pyemu.os_utils.run("pestpp-ies freyberg.pst", cwd=t_d)
```

Now we can read in the results and make some figures showing residuals and the balance of the objective function


```
In [17]: pst = pyemu.Pst(os.path.join(t_d,"freyberg.pst"))
        print(pst.phi)
        plt.figure()
        pst.plot(kind='phi_pie')
        print('Here are the non-zero weighted observation names')

        figs = pst.plot(kind="1to1")
        plt.show()
        pst.res.loc[pst.nnz_obs_names, :]
```





540.6801963885084

Here are the non-zero weighted observation names

<Figure size 432x288 with 0 Axes>



<Figure size 576x756 with 0 Axes>



Out[17]:

	name	group	measured	modelled \
	name			
	fo_39_19791230	fo_39_19791230 calflux	10949.810469	11430.000000

hds_00_002_009_000	hds_00_002_009_000	calhead	37.383407	37.107498
hds_00_002_015_000	hds_00_002_015_000	calhead	34.857895	35.045185
hds_00_003_008_000	hds_00_003_008_000	calhead	37.796378	37.397289
hds_00_009_001_000	hds_00_009_001_000	calhead	39.745220	39.546417
hds_00_013_010_000	hds_00_013_010_000	calhead	35.268247	35.571774
hds_00_015_016_000	hds_00_015_016_000	calhead	34.748943	34.835716
hds_00_021_010_000	hds_00_021_010_000	calhead	34.696031	35.386250
hds_00_022_015_000	hds_00_022_015_000	calhead	34.409016	34.577492
hds_00_024_004_000	hds_00_024_004_000	calhead	35.182948	36.760464
hds_00_026_006_000	hds_00_026_006_000	calhead	34.697079	35.896149
hds_00_029_015_000	hds_00_029_015_000	calhead	34.352680	34.453842
hds_00_033_007_000	hds_00_033_007_000	calhead	34.151406	34.678810
hds_00_034_010_000	hds_00_034_010_000	calhead	33.649031	34.118073

	residual	weight
name		
fo_39_19791230	-480.189531	0.005
hds_00_002_009_000	0.275909	10.000
hds_00_002_015_000	-0.187290	10.000
hds_00_003_008_000	0.399089	10.000
hds_00_009_001_000	0.198803	10.000
hds_00_013_010_000	-0.303527	10.000
hds_00_015_016_000	-0.086773	10.000
hds_00_021_010_000	-0.690219	10.000
hds_00_022_015_000	-0.168475	10.000
hds_00_024_004_000	-1.577516	10.000
hds_00_026_006_000	-1.199070	10.000
hds_00_029_015_000	-0.101162	10.000
hds_00_033_007_000	-0.527404	10.000
hds_00_034_010_000	-0.469041	10.000

Publication ready figs - oh snap!

Depending on the truth you chose, we may have a problem - we set the weights for both the heads and the flux to reasonable values based on what we expect for measurement noise. But the contributions to total phi might be out of balance - if contribution of the flux measurement to total phi is too low, the history matching excersizes (coming soon!) will focus almost entirely on minimizing head residuals. So we need to balance the objective function. This is a subtle but very important step, especially since some of our forecasts deal with sw-gw exchange

```
In [18]: #pc = pst.phi_components
#target = {"calflux":0.3 * pc["calhead"]}
#pst.adjust_weights(obsgrp_dict=target)
#pst.plot(kind='phi_pie')
```

Just to make sure we have everything working right, we should be able to load the truth parameters, run the model once and have a phi equivalent to the noise vector:

```
In [19]: par_df = pd.read_csv(os.path.join(m_d,"sweep_in.csv"),index_col=0)
pst.parameter_data.loc[:,"parval1"] = par_df.loc[idx,pst.par_names]
pst.write(os.path.join(m_d,"test.pst"))
```

we will run this with noptmax=0 to preform a single run. Pro-tip: you can use any of the pestpp-### binaries/executables to run noptmax=0

```
In [20]: pyemu.os_utils.run("pestpp-ies.exe test.pst", cwd=m_d)
         pst = pyemu.Pst(os.path.join(m_d, "test.pst"))
         print(pst.phi)
         pst.res.loc[pst.nnz_obs_names,:]
```

17.528847215722436

```
Out [20]:
```

	name	group	measured	modelled \
name				
fo_39_19791230	fo_39_19791230	calflux	10949.810469	10597.000000
hds_00_002_009_000	hds_00_002_009_000	calhead	37.383407	37.343391
hds_00_002_015_000	hds_00_002_015_000	calhead	34.857895	34.760021
hds_00_003_008_000	hds_00_003_008_000	calhead	37.796378	37.572289
hds_00_009_001_000	hds_00_009_001_000	calhead	39.745220	39.558464
hds_00_013_010_000	hds_00_013_010_000	calhead	35.268247	35.365974
hds_00_015_016_000	hds_00_015_016_000	calhead	34.748943	34.653934
hds_00_021_010_000	hds_00_021_010_000	calhead	34.696031	34.711166
hds_00_022_015_000	hds_00_022_015_000	calhead	34.409016	34.419338
hds_00_024_004_000	hds_00_024_004_000	calhead	35.182948	35.141888
hds_00_026_006_000	hds_00_026_006_000	calhead	34.697079	34.682674
hds_00_029_015_000	hds_00_029_015_000	calhead	34.352680	34.207253
hds_00_033_007_000	hds_00_033_007_000	calhead	34.151406	34.075302
hds_00_034_010_000	hds_00_034_010_000	calhead	33.649031	33.636864


	residual	weight
name		
fo_39_19791230	352.810469	0.005
hds_00_002_009_000	0.040016	10.000
hds_00_002_015_000	0.097874	10.000
hds_00_003_008_000	0.224089	10.000
hds_00_009_001_000	0.186756	10.000
hds_00_013_010_000	-0.097728	10.000
hds_00_015_016_000	0.095009	10.000
hds_00_021_010_000	-0.015136	10.000
hds_00_022_015_000	-0.010322	10.000
hds_00_024_004_000	0.041060	10.000
hds_00_026_006_000	0.014404	10.000
hds_00_029_015_000	0.145427	10.000
hds_00_033_007_000	0.076104	10.000
hds_00_034_010_000	0.012168	10.000

The residual should be exactly the noise values from above. Lets load the model (that was just run using the true pars) and check some things


```
In [21]: m = flopy.modflow.Modflow.load("freyberg.nam", model_ws=m_d)
```

```
In [22]: a = m.upw.vka[1].array
         #a = m.rch.rech[0].array
         a = np.ma.masked_where(m.bas6.ibound[0].array==0,a)
         print(a.min(),a.max())
         c = plt.imshow(a)
         plt.colorbar()
         plt.show()
```

0.00231499 2.546079





```
In [23]: lst = flopy.utils.MfListBudget(os.path.join(m_d,"freyberg.list"))
         df = lst.get_dataframes(diff=True)[0]
         ax = df.plot(kind="bar",figsize=(6,6), grid=True)
         a = ax.set_xticklabels(["historic","scenario"],rotation=90)
         plt.show(ax)
```








1.0.8 see how our existing observation ensemble compares to the truth



forecasts:

```
In [24]: obs = pst.observation_data
plt.figure()
for forecast in fnames:
    ax = plt.subplot(111)
    obs_df.loc[:,forecast].hist(ax=ax,color="0.5",alpha=0.5)
    ax.plot([obs.loc[forecast,"obsval"],obs.loc[forecast,"obsval"]],ax.get_ylim(),"r")
    ax.set_title(forecast)
plt.show()
```









observations:

```
In [25]: for oname in pst.nnz_obs_names:
          ax = plt.subplot(111)
          obs_df.loc[:, oname].hist(ax=ax, color="0.5", alpha=0.5)
          ax.plot([obs.loc[oname, "obsval"], obs.loc[oname, "obsval"]], ax.get_ylim(), "r")
          ax.set_title(oname)
          plt.show()
```

